

Overfitting Intuition

Project: Transparent Book Recommender

1. Your mission

Your learning target

I can compare training and validation evidence for model capacity.

I can explain overfitting without assuming every gap proves it.

I can use a declared selection rule instead of perfect training fit.

Remember before reading

Close your notes first. A rough first answer helps us choose the right starting point; it is not a score for your ability.

Without opening the old lesson, explain one key idea from Choose Informative Questions. Give an example and a mistake that the idea helps you avoid.


Without opening the old lesson, explain one key idea from Types of Classification Error. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Human Classifier Simulation. Give an example and a mistake that the idea helps you avoid.

After trying, check the relevant lesson or ask your teacher. Change the part of your explanation that the evidence shows was wrong.

Week 16: Overfitting Intuition

A rule that memorizes examples may fail on new cases.



THINKER MISSION

Contrast a simple rule with a seven-exception rule.

Thinker name and date:

Write or draw your reasoning and evidence.

Readiness for the connected explanation

Can a zero training error guarantee good validation performance?

This week's end goal


By the end, I can use development performance rather than training fit alone to choose model complexity.

2. Notice the evidence

PICTURE STORY

Notice → Predict → Explain

Begin with visible evidence. Do not explain more than the picture and information support.



MYSTERY

The learners must use overfitting intuition to solve a problem. What information do they have? What decision or calculation do they make? What would count as evidence that it worked?

Three observations (what you can point to):

Write or draw your reasoning and evidence.

My prediction and the clue supporting it:

Write or draw your reasoning and evidence.

A closer explanation

A more flexible model can represent complicated relationships and accidental details. Capacity can help when there is enough relevant evidence, but perfect training performance is not the goal. Compare training and validation using the same metric while checking leakage and representativeness.


Use fictional classroom data throughout this investigation. Keep the task, units, reference labels and data split visible as you reason. Earlier weeks supplied the model-building process; this week examines overfitting and model capacity.

3. See how it works

VISUAL MECHANISM

Input → Process → Output

Trace the information. At each stage, ask what changed and how you could test it.





PROCESS

A rule that memorizes examples may fail on new cases.



OUT

OUTPUT

Contrast a simple rule with a seven-exception rule.

Which stage could fail even when the other two are correct? Explain:

Write or draw your reasoning and evidence.

A closer explanation

On the same regression split, a constant has training MAE 5 and validation 6, a line has training 2 and validation 3, and a highly flexible curve has training 0 and validation 8. Under smallest validation MAE, choose the line. The curve has the lowest training loss but the worst validation result.

Before continuing, underline the decision rule and circle the evidence that supports the conclusion. Explain the example to a partner without replacing the counts or conditions with a stronger claim. Keep the distinction between an observed result and a proposed next step.

Connect the picture and the explanation


Point to the input, the deciding step and the result in this week's visual. Explain the same step in ordinary words. A picture helps when you can say what each part represents.


4. Compare the cases

CASE GALLERY


Four Tests, Four Kinds of Evidence

A system is not understood until it survives more than one comfortable example.





NORMAL CASE: predict and name evidence. BOUNDARY CASE: predict and name evidence.





MISSING INFORMATION: predict and name evidence. TRANSFER CASE: predict and name evidence.

Which case is most likely to reveal a hidden assumption? Why?

Write or draw your reasoning and evidence.

Change one thing in the pictured case

Change: Complex's validation error becomes 2 on a newly documented, fair development comparison.

Prediction: The stated rule now selects Complex.

Reason: The rule compares validation errors; it does not ban complex models.

5. Words that unlock the idea

VOCABULARY LAB

Words You Can Use to Reason

A definition matters when you can recognize the idea, use it, and contrast it with a near-miss.

1

OVERFITTING

matching training details
that do not generalize

2

GENERALIZE

work on new relevant cases

3

COMPLEXITY

number of conditions or exceptions

4

EXCEPTION

special case outside the main rule

Word · My own example · Not this / boundary

Word: overfitting

My own example: Adapting too closely to development examples so performance does not transfer well to new relevant cases.

Not this / boundary: _____

Word: generalize

My own example: Perform usefully on new relevant cases beyond those used to develop a system.

Not this / boundary: _____

Word: complexity

My own example: The amount of structure, work or resources a method requires, measured in a stated way.

Not this / boundary: _____

Word: exception

My own example: A condition that requires handling outside the usual path.

Not this / boundary: _____

Words used in the closer explanation

Word · Meaning

Word: Capacity

Meaning: Flexibility of a model to represent patterns

Word: Underfitting

Meaning: Failing to capture useful structure

Word: Generalization

Meaning: Performance on relevant unseen examples

6. Follow a worked example

WORKED EVIDENCE

Reason Step by Step

Predict first. Then check each step. A correct answer without visible reasoning is incomplete evidence.

WEEK 16 • WORKED EVIDENCE

- 1 Simple rule: $\text{difficulty} \leq \text{reader level} + 1$.
- 2 Memorized rule: seven title-specific exceptions.
- 3 Both fit the seen books; only new books reveal generalization.
- 4 Prefer the simplest rule that performs well on fair validation cases.

1. Simple rule: $\text{difficulty} \leq \text{reader level} + 1$.
2. Memorized rule: seven title-specific exceptions.
3. Both fit the seen books; only new books reveal generalization.
4. Prefer the simplest rule that performs well on fair validation cases.

Recalculate, retrace, or restate the reasoning in your own representation:

Write or draw your reasoning and evidence.

CHECK Did the evidence answer the exact question? Did the total, path, or source record stay consistent?

A closer explanation

Overfitting and model capacity at a glance

Part · Evidence or action · Interpretation

Part: Constant

Evidence or action: Training 5

Interpretation: Validation 6

Part: Line

Evidence or action: Training 2

Interpretation: Validation 3

Part: Flexible curve

Evidence or action: Training 0

Interpretation: Validation 8

The curve result is consistent with overfitting, but a gap alone does not prove its cause. A changed population, bad labels or a split error could also produce a gap. Inspect the procedure before claiming a diagnosis. The constant may miss useful structure, but these numbers alone do not identify the true relationship.

Does a training-validation gap alone prove overfitting? Explain the steps or evidence behind your response.

Learn from the worked case

Cover the result before checking it. Say what is given, choose the procedure, and follow one step at a time. Then uncover the result and explain your first mismatch. Ask why a step is allowed, not only what number it produces.

A pictorial worked case: Overfitting Intuition

A fictional project compares three book-length prediction models on the same development sets.

The facts we will use

Mean absolute errors in pages: Simple training 5, validation 6; Medium training 2, validation 3; Complex training 0, validation 7.

Select lowest validation error; untouched final tests are not used for selection.

Training and validation errors

	Training MAE pages	Validation MAE pages
Simple	5	6
Medium	2	3
Complex	0	7

Read a cell using both its row and column labels.

Overfitting Intuition

Start with the supplied facts

Follow each arrow. Ask which fact allows the next step.



1. Training comparison

Complex fits training perfectly with zero error.



2. Validation comparison

Medium has error 3, versus Simple 6 and Complex 7.



3. Choose for development

Select Medium under the declared lowest-validation-error rule.



4. Final evaluation

Freeze the selection, then evaluate new reserved cases.

CHANGE ONE THING

Complex's validation error becomes 2 on a newly documented, fair development comparison.

Explain the changed result using the matching lesson.

Read the picture one step at a time

Step · What happens and why

Step: 1. Training comparison

What happens and why: Complex fits training perfectly with zero error.

Step: 2. Validation comparison

What happens and why: Medium has error 3, versus Simple 6 and Complex 7.

Step: 3. Choose for development

What happens and why: Select Medium under the declared lowest-validation-error rule.

Step: 4. Final evaluation

What happens and why: Freeze the selection, then evaluate new reserved cases.

Match each pictured step to the supplied facts. This is a worked teaching case, not an independent assessment.

Why this result follows

A flexible model can fit quirks of its training set without predicting new examples well. Here Complex's training success occur together with the worst validation error. The numbers are invented comparisons, not a claim that complexity always hurts. Choosing Medium uses development evidence; its validation score must not be advertised as an untouched final-test result.

What this example does not establish

A single small validation set may be noisy; selection evidence does not prove future population performance.

7. Find and repair a mistake

ERROR DETECTIVE

Compare → Diagnose → Repair

Do not patch the answer. Find the earliest unsupported assumption, wrong step, or weak evidence.

!

COMPARE • DIAGNOSE

Weak idea: A rule that is perfect on training data must be the best rule. Strong move: locate evidence, change one cause, and retest.

Evidence record · Expected · Actual · Likely cause / next test

Evidence record: normal case

Expected: works

Actual: _____

Likely cause / next test: _____

Evidence record: boundary case

Expected: stated rule

Actual: _____

Likely cause / next test: _____

Evidence record: missing input

Expected: safe response

Actual: _____

Likely cause / next test: _____

Evidence record: fresh transfer

Expected: generalizes

Actual: _____

Likely cause / next test: _____

Repair one cause. Why is this repair better than changing several things?

Write or draw your reasoning and evidence.

A closer explanation

A simpler model is not always better. If a flexible candidate has better valid evidence under the declared rule, keep that evidence visible. Model capacity is one part of an experiment, alongside features, sample

size, label quality and the split.

What extra evidence would you inspect before concluding a model memorized noise?

8. Practise together

LOGIC STUDIO

Trace Before You Run

Block code, pseudocode, tables, and diagrams are different ways to expose the same reasoning.



TRACE THE LOGIC

Predict the result. Trace every condition or calculation. Compare expected with actual. Explain the first difference.

SET evidence_record PREDICT expected_result RUN one normal case and one boundary case
COMPARE actual_result with expected_result IF different: locate first cause; change one factor; RETEST
EXPLAIN result, limitation, and next safe step

Trace one input. Which line changes the state or chooses a path?

Write or draw your reasoning and evidence.

A closer explanation

Change one development factor at a time: simplify the candidate, improve data coverage or adjust a declared regularization setting. Refit on training and compare on validation. Do not remove difficult valid examples or repeatedly search the test for a favorable result.

Practise with less help: Interpret a gap

Use workbook W2 for the connected example “Overfitting and model capacity”. First fill the input and rule boxes. Try the middle steps before asking for a hint. After a hint, try the next step yourself.

Plan · Do · Check

Plan: What is given? Which rule or procedure applies?

Do: Record each intermediate step.

Check: Does the result follow? What remains unknown?

Explain your trace to a partner. Your partner points to one step to check. Swap roles on the next case. Each person writes their own explanation.

Check your reasoning

Use the worked case for Overfitting and model capacity. Proposed explanation: “A final answer is enough”. Decide whether this explanation is supported. Point to the step or supplied fact that decides; explain your repair if needed.

Answer on your own before discussion. If you are unsure, say which step you need help with.

Explain the picture

Does zero training error justify choosing Complex in the original record?

My explanation using a pictured step and the given facts

This is supported practice. Keep the separately named independent case for an attempt before feedback.

9. Run the investigation

EVIDENCE LAB

Predict → Test → Record → Explain

Keep expected and actual results separate. Surprises are useful data when recorded honestly.



PREDICT → TEST → RECORD → EXPLAIN

Compare simple and exception-heavy rules on training cards, then on sealed new cards. Count where performance changes and explain why.

Compare simple and exception-heavy rules on training cards, then on sealed new cards. Count where performance changes and explain why.

Test / input · Prediction · Actual · Evidence / calculation · Change or keep?

Test / input: normal

Prediction: _____

Actual: _____

Evidence / calculation: _____

Change or keep?: _____

Test / input: boundary

Prediction: _____

Actual: _____

Evidence / calculation: _____

Change or keep?: _____

Test / input: missing

Prediction: _____

Actual: _____

Evidence / calculation: _____

Change or keep?: _____

Test / input: fresh transfer

Prediction: _____

Actual: _____

Evidence / calculation: _____

Change or keep?: _____

What did the hardest test teach you that the normal test did not?

Write or draw your reasoning and evidence.

Plan → try → check

Before trying: what do I expect, and why? While trying: which step could first go wrong? Afterwards: what did I actually observe, and what should change? Keep “not run” if you only traced the procedure.

10. Build your understanding

MASTERY

From Knowing the Word to Owning the Idea

Move upward only when you can produce evidence without copying the worked example.



MASTERY LADDER

1 Name it • 2 Trace it • 3 Apply it • 4 Challenge it • 5 Explain limits

- 1 • NAME: Define the concept without repeating the textbook.
- 2 • TRACE: Show the information, state, branch, calculation, or evidence path.
- 3 • APPLY: Solve a different example independently.
- 4 • CHALLENGE: Find a boundary case, counterexample, or unfair outcome.
- 5 • EXPLAIN LIMITS: Say when the result should not be trusted or used alone.

My strongest evidence of independent understanding:

Write or draw your reasoning and evidence.

11. Connect your project

CUMULATIVE PROJECT

Transparent Book Recommender

Every week adds one inspectable piece. The system must show how it reached a result.



TRANSPARENT BOOK RECOMMENDER

Avoid title-by-title patches; improve features or thresholds that transfer.

THIS WEEK'S CHECKPOINT Avoid title-by-title patches; improve features or thresholds that transfer.

SYSTEM MAP Reader need → allowed inputs → transparent rule or score → ranked candidates → explanation → test record → human decision

What will you add, change, or verify in the project this week?

Write or draw your reasoning and evidence.

What should remain under human control?

Write or draw your reasoning and evidence.

RESPONSIBLE DESIGN Collect only necessary information. Show uncertainty and limits. Never confuse a ranking with a fact about a person.

Evidence for your project

Compare recommender revisions on fixed development cases, reserving new cases for the final

presentation.

12. Show what you know

INDEPENDENT MASTERY PROOF

Explain • Transfer • Test • Defend

Complete this page without copying the worked example. Your reasoning matters more than decoration.



INDEPENDENT PROOF

Contrast a simple rule with a seven-exception rule.

CHALLENGE Contrast a simple rule with a seven-exception rule.

1. My fresh example, trace, calculation, or design:

Write or draw your reasoning and evidence.

2. My evidence that it works (include a boundary or missing-data case):

Write or draw your reasoning and evidence.

3. One limitation, fairness concern, or situation needing human review:

Write or draw your reasoning and evidence.

SELF-CHECK ☐ I explained the mechanism ☐ I used evidence ☐ I tested a fresh case ☐ I named a limitation ☐ I can answer 'why?'

13. Study a complete case

Week 16 Worked case

Use the supplied details to practise the weekly concept before attempting the independent question. This case extends the illustrated lesson.

The situation

Rule A memorises six training book names. It scores 6/6 familiar and 2/6 new. Rule B uses relevant traits but is untested.

The question

What does the evidence support?

Worked explanation

A shows weak transfer. B might generalise better but requires testing; simplicity is not proof.

Explain it in your own words

Which detail supports the answer, and why does it matter?

Write or draw your reasoning and evidence.

Trace the supplied case visually

Perfect recall can hide weak transfer

Evidence or stage · Value or result

Evidence or stage: Familiar books

Value or result: 6/6

Evidence or stage: New books

Value or result: 2/6

Rule A memorises six training book names. It scores 6/6 familiar and 2/6 new. Rule B uses relevant traits but is untested.

What does the evidence support?

A shows weak transfer. B might generalise better but requires testing; simplicity is not proof.

14. Try a new case independently

Independent case: Apply independently

Models A training 4 validation 5, B training 1 validation 2, C training 0 validation 2. Lowest validation MAE then simpler; B is simpler than C. Choose and justify.

My independent explanation, trace and limits

Attempt this case before feedback. Record any help. Earlier examples below are further practice; a case

already practised with feedback cannot establish fresh transfer.

Week 16 Independent application

Close the worked explanation. Apply the idea to this different question without copying.

Rule A memorises six training book IDs. Rule B uses a relevant topic feature. Why might A fail on a new ID? What test is needed?

My answer, with a trace, calculation or supporting evidence

Write or draw your reasoning and evidence.

Create and test

Change one relevant detail. Predict the result before testing. Include a boundary or missing-information case when it helps reveal a limit.

My changed input and expected result

Write or draw your reasoning and evidence.

Connect to your project

Avoid title-by-title patches; improve features or thresholds that transfer.

The evidence I will keep

Write or draw your reasoning and evidence.

Your teacher has the independent answer and scoring criteria. Complete the matching workbook week to keep your practice evidence together.

Transfer practice from the original programme

Rule A memorises eight training IDs and gets 8/8 training, 2/8 validation. Rule B uses genre and gets 6/8 training, 6/8 validation. Which is preferable on this evidence, and what test remains?

Explain the mechanism, evidence and one limit.

Use feedback to improve

Attempt the independent case before hints. Record whether you worked alone, used a hint or followed a model. After feedback, fix the first unsupported step and explain why it changed. A later check uses a changed case, not a copied answer.