

Predictive vs Generative

Project: Transparent Book Recommender

1. Your mission

Your learning target

I can distinguish discriminative prediction from generation by inspecting the requested output.

I can calculate an empirical next-token distribution from counted continuations.

I can separate next-token fit from factual support and downstream usefulness.

Remember before reading

Close your notes first. A rough first answer helps us choose the right starting point; it is not a score for your ability.

Without opening the old lesson, explain one key idea from Image and Voice Pipelines. Give an example and a mistake that the idea helps you avoid.


Without opening the old lesson, explain one key idea from Feedback Loops. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Similarity by Selected Traits. Give an example and a mistake that the idea helps you avoid.

After trying, check the relevant lesson or ask your teacher. Change the part of your explanation that the evidence shows was wrong.

Week 27: Predictive vs Generative

Predictive systems select a target; generative systems construct new output.



THINKER MISSION

Match problems to classifier, predictor, or generator.

Thinker name and date:

Write or draw your reasoning and evidence.

Readiness for the connected explanation

What is the difference between a model prediction and evidence that a statement is true?

This week's end goal


By the end, I can distinguish a generated continuation from a prediction checked against the world.

2. Notice the evidence

PICTURE STORY

Notice → Predict → Explain

Begin with visible evidence. Do not explain more than the picture and information support.



MYSTERY

The learners must use predictive vs generative to solve a problem. What information do they have? What decision or calculation do they make? What would count as evidence that it worked?

Three observations (what you can point to):

Write or draw your reasoning and evidence.

My prediction and the clue supporting it:

Write or draw your reasoning and evidence.

A closer explanation

A classifier might map a message to the label request or report. A text generator instead produces a sequence, predicting a distribution over the next token given a prefix. Both make predictions, but the training targets and permitted outputs differ. One application can combine generation with a classifier, retrieval and human approval. Calling a whole product “generative” does not describe every component.


The objective defines what training attempts to fit. Next-token fit rewards matching observed continuations; it does not directly certify that an answer is true, well sourced or permitted. A fluent completion is still an output to evaluate against the task.

3. See how it works

VISUAL MECHANISM

Input → Process → Output

Trace the information. At each stage, ask what changed and how you could test it.



↓



PROCESS

Predictive systems select a target; generative systems construct new output.



OUT

OUTPUT

Match problems to classifier, predictor, or generator.

Which stage could fail even when the other two are correct? Explain:

Write or draw your reasoning and evidence.

A closer explanation

Prefix · Continuation · Count

Prefix: The marker is

Continuation: red

Count: 6

Prefix: The marker is

Continuation: blue

Count: 3

Prefix: The marker is

Continuation: green

Count: 1

There are ten observed continuations. The empirical probabilities are red $6/10 = 0.6$, blue $3/10 = 0.3$, green $1/10 = 0.1$. They are nonnegative and sum to 1. A simple frequentist classroom model uses these relative counts as its next-token distribution. These probabilities describe the supplied corpus counts, not the current colour of any actual marker.

Keep the evidence roles separate

Plan or input · Procedure · Permitted conclusion

Plan or input: Ten corpus continuations

Procedure: Count red 6, blue 3, green 1

Permitted conclusion: Distribution 0.6, 0.3, 0.1

Plan or input: Prefix: The marker is

Procedure: Select a next token

Permitted conclusion: An output, not an observation

Plan or input: Current marker question

Procedure: Check direct evidence

Permitted conclusion: Colour claim needs support

The continuation model and a factual observation have different evidence sources.

Connect the picture and the explanation


Point to the input, the deciding step and the result in this week's visual. Explain the same step in ordinary words. A picture helps when you can say what each part represents.


4. Compare the cases

CASE GALLERY

Four Tests, Four Kinds of Evidence

A system is not understood until it survives more than one comfortable example.






BOUNDARY

Near a decision edge

NORMAL CASE: predict and name evidence. BOUNDARY CASE: predict and name evidence.



MISSING

Needed information absent



MISSING INFORMATION: predict and name evidence. TRANSFER CASE: predict and name evidence.

Which case is most likely to reveal a hidden assumption? Why?

Write or draw your reasoning and evidence.

Change one thing in the pictured case

Change: Change the training table to full 2, empty 7, unknown 1.

Prediction: Greedy output becomes empty with probability 0.7.

Reason: The continuation distribution changed; the independently observed box has not changed.

5. Words that unlock the idea

VOCABULARY LAB

Words You Can Use to Reason

A definition matters when you can recognize the idea, use it, and contrast it with a near-miss.

1

PREDICTIVE

estimates a label or number

2

GENERATIVE

constructs new content

3

OUTPUT SPACE

set of possible outputs

4

FIT FOR PURPOSE

mechanism matches the
goal

Word · My own example · Not this / boundary

Word: predictive

My own example: Estimating a target from supplied inputs.

Not this / boundary: _____

Word: generative

My own example: Producing new content from learned patterns and supplied conditions.

Not this / boundary: _____

Word: output space

My own example: The set or range of outputs a system can produce for a task.

Not this / boundary: _____

Word: fit for purpose

My own example: Suitable for the stated users, task, conditions and consequences.

Not this / boundary: _____

Words used in the closer explanation

Word · Meaning

Word: Generative objective

Meaning: A training aim concerned with producing or modelling possible outputs

Word: Discriminative objective

Meaning: A training aim mapping inputs to a specified target such as a label

Word: Conditional distribution

Meaning: Probabilities for alternatives given stated information

Word: Autoregressive generation

Meaning: Producing a sequence one next item at a time using preceding items

6. Follow a worked example

WORKED EVIDENCE

Reason Step by Step

Predict first. Then check each step. A correct answer without visible reasoning is incomplete evidence.

WEEK 27 • WORKED EVIDENCE

- 1 'Is this book suitable?' → classification.
- 2 'How many days to finish?' → numeric prediction.
- 3 'Write a spoiler-free teaser' → generation.
- 4 A product may combine them, but each stage needs its own test.

1. 'Is this book suitable?' → classification.
2. 'How many days to finish?' → numeric prediction.
3. 'Write a spoiler-free teaser' → generation.
4. A product may combine them, but each stage needs its own test.

Recalculate, retrace, or restate the reasoning in your own representation:

Write or draw your reasoning and evidence.

CHECK Did the evidence answer the exact question? Did the total, path, or source record stay consistent?

A closer explanation

After producing red, the next step conditions on the expanded prefix "The marker is red". A second table might give an end token probability 0.8 and the token today probability 0.2. A generated sequence can stop at the end token or continue. Do not reuse the first-step table unchanged after adding a token unless the model explicitly ignores the new prefix.

A separate classifier can return red as a label from an image; that is a different input and target from continuing text. A downstream evaluation asks whether the system correctly reports the current marker. It must include ground truth or sources and appropriate failure handling, not just count plausible phrases.

Learn from the worked case

Cover the result before checking it. Say what is given, choose the procedure, and follow one step at a time. Then uncover the result and explain your first mismatch. Ask why a step is allowed, not only what number it produces.

A pictorial worked case: Predictive vs Generative

A fictional toy generator predicts the next word after “The box is”.

The facts we will use


Observed training continuations: full 6, empty 3, unknown 1.

Use counts divided by total 10. The current physical box is separately checked and found empty.

AI ACADEMY / AI-2/27

Toy next-word distribution

probability



Bar length shows the supplied value; the marked baseline is zero.

Predictive vs Generative

Start with the supplied facts

Follow each arrow. Ask which fact allows the next step.



1. Count table

Continuation counts are 6,3,1; total 10.



2. Probabilities

full=0.6; empty=0.3; unknown=0.1.



3. Greedy continuation

Largest probability chooses full for this toy prefix.



4. Check reality

Current observation says empty; frequent wording is not a present inventory reading.

CHANGE ONE THING

Change the training table to full 2, empty 7, unknown 1.

Explain the changed result using the matching lesson.

Read the picture one step at a time

Step · What happens and why

Step: 1. Count table

What happens and why: Continuation counts are 6,3,1; total 10.

Step: 2. Probabilities

What happens and why: full=0.6; empty=0.3; unknown=0.1.

Step: 3. Greedy continuation

What happens and why: Largest probability chooses full for this toy prefix.

Step: 4. Check reality

What happens and why: Current observation says empty; frequent wording is not a present inventory reading.

Match each pictured step to the supplied facts. This is a worked teaching case, not an independent assessment.

Why this result follows

The table predicts a continuation in a text exercise, not the contents of the current box. Generating the most frequent token is a model-defined decision. A factual inventory statement needs current evidence. The numerical table is intentionally small; real generative systems condition on richer context and do not simply consult these three counts.

What this example does not establish

A one-step toy distribution does not describe every predictive or generative model or validate factual accuracy.

Read the labelled picture: Does matching the physical box in the changed table prove the generator measures boxes?

Change one thing: Change the training table to full 2, empty 7, unknown 1. Predict the result before checking the explanation. What remains the same?

7. Find and repair a mistake

ERROR DETECTIVE

Compare → Diagnose → Repair

Do not patch the answer. Find the earliest unsupported assumption, wrong step, or weak evidence.

COMPARE • DIAGNOSE



Weak idea: Anything that produces an answer is generative AI. Strong move: locate evidence, change one cause, and retest.

Evidence record • Expected • Actual • Likely cause / next test

Evidence record: normal case

Expected: works

Actual: _____

Likely cause / next test: _____

Evidence record: boundary case

Expected: stated rule

Actual: _____

Likely cause / next test: _____

Evidence record: missing input

Expected: safe response

Actual: _____

Likely cause / next test: _____

Evidence record: fresh transfer

Expected: generalizes

Actual: _____

Likely cause / next test: _____

Repair one cause. Why is this repair better than changing several things?

Write or draw your reasoning and evidence.

A closer explanation

Keep three questions separate: did the model fit continuation data; did the output satisfy the requested format or task; and are factual claims supported? A system that drafts a fictional description has different criteria from one reporting a live measurement. Record the source and time when a claim refers to changing conditions. No AI account is needed to reason about these distinctions.

Exit: a team says “the highest-probability continuation proves the current fact”. Name the evidence they actually have and the evidence still needed.

8. Practise together

LOGIC STUDIO

Trace Before You Run

Block code, pseudocode, tables, and diagrams are different ways to expose the same reasoning.



TRACE THE LOGIC

Predict the result. Trace every condition or calculation. Compare expected with actual. Explain the first difference.

SET evidence_record PREDICT expected_result RUN one normal case and one boundary case
COMPARE actual_result with expected_result IF different: locate first cause; change one factor; RETEST
EXPLAIN result, limitation, and next safe step

Trace one input. Which line changes the state or chooses a path?

Write or draw your reasoning and evidence.

A closer explanation

Supported case: the prefix “The door is” is followed in a tiny corpus by open four times and closed once. The empirical probabilities are $\frac{4}{5}$ and $\frac{1}{5}$. A current sensor reports closed. Selecting the most frequent continuation open would be plausible under the corpus model but inconsistent with the sensor observation.

Guided: compute both probabilities and explain why selecting open does not establish the current door state.

Practise with less help: Compute the distribution

Use workbook W2 for the connected example “Generative and predictive objectives”. First fill the input and rule boxes. Try the middle steps before asking for a hint. After a hint, try the next step yourself.

Plan · Do · Check

Plan: What is given? Which rule or procedure applies?

Do: Record each intermediate step.

Check: Does the result follow? What remains unknown?

Explain your trace to a partner. Your partner points to one step to check. Swap roles on the next case. Each person writes their own explanation.

Check your reasoning

Use the worked case for Generative and predictive objectives. Proposed explanation: “Generative models do not predict”. Decide whether this explanation is supported. Point to the step or supplied fact that decides; explain your repair if needed.

Answer on your own before discussion. If you are unsure, say which step you need help with.

Explain the picture

Does matching the physical box in the changed table prove the generator measures boxes?

My explanation using a pictured step and the given facts

This is supported practice. Keep the separately named independent case for an attempt before feedback.

9. Run the investigation

EVIDENCE LAB

Predict → Test → Record → Explain

Keep expected and actual results separate. Surprises are useful data when recorded honestly.



PREDICT → TEST → RECORD → EXPLAIN

Sort nine task cards by mechanism, name the output space, and specify one evaluation method for each category.

Sort nine task cards by mechanism, name the output space, and specify one evaluation method for each category.

Test / input · Prediction · Actual · Evidence / calculation · Change or keep?

Test / input: normal

Prediction: _____

Actual: _____

Evidence / calculation: _____

Change or keep?: _____

Test / input: boundary

Prediction: _____

Actual: _____

Evidence / calculation: _____

Change or keep?: _____

Test / input: missing

Prediction: _____

Actual: _____

Evidence / calculation: _____

Change or keep?: _____

Test / input: fresh transfer

Prediction: _____

Actual: _____

Evidence / calculation: _____

Change or keep?: _____

What did the hardest test teach you that the normal test did not?

Write or draw your reasoning and evidence.

Plan → try → check

Before trying: what do I expect, and why? While trying: which step could first go wrong? Afterwards: what did I actually observe, and what should change? Keep “not run” if you only traced the procedure.

10. Build your understanding

MASTERY

From Knowing the Word to Owning the Idea

Move upward only when you can produce evidence without copying the worked example.



MASTERY LADDER
1 Name it • 2 Trace it • 3 Apply it • 4 Challenge it • 5 Explain limits

1 • NAME: Define the concept without repeating the textbook.

2 • TRACE: Show the information, state, branch, calculation, or evidence path.

3 • APPLY: Solve a different example independently.

4 • CHALLENGE: Find a boundary case, counterexample, or unfair outcome.

5 • EXPLAIN LIMITS: Say when the result should not be trusted or used alone.

My strongest evidence of independent understanding:

Write or draw your reasoning and evidence.

11. Connect your project

CUMULATIVE PROJECT

Transparent Book Recommender

Every week adds one inspectable piece. The system must show how it reached a result.



TRANSPARENT BOOK RECOMMENDER

Keep scoring predictive and explanations templated or carefully generated and verified.

THIS WEEK'S CHECKPOINT Keep scoring predictive and explanations templated or carefully generated and verified.

SYSTEM MAP Reader need → allowed inputs → transparent rule or score → ranked candidates → explanation → test record → human decision

What will you add, change, or verify in the project this week?

Write or draw your reasoning and evidence.

What should remain under human control?

Write or draw your reasoning and evidence.

RESPONSIBLE DESIGN Collect only necessary information. Show uncertainty and limits. Never confuse a ranking with a fact about a person.

Evidence for your project

Use approved book facts for recommender explanations; generated fluency is not a replacement for catalogue evidence.

12. Show what you know

INDEPENDENT MASTERY PROOF

Explain • Transfer • Test • Defend

Complete this page without copying the worked example. Your reasoning matters more than decoration.



INDEPENDENT PROOF

Match problems to classifier, predictor, or generator.

CHALLENGE Match problems to classifier, predictor, or generator.

1. My fresh example, trace, calculation, or design:

Write or draw your reasoning and evidence.

2. My evidence that it works (include a boundary or missing-data case):

Write or draw your reasoning and evidence.

3. One limitation, fairness concern, or situation needing human review:

Write or draw your reasoning and evidence.

SELF-CHECK ☐ I explained the mechanism ☐ I used evidence ☐ I tested a fresh case ☐ I named a limitation ☐ I can answer 'why?'

13. Study a complete case

Week 27 Worked case

Use the supplied details to practise the weekly concept before attempting the independent question. This case extends the illustrated lesson.

The situation

A labels genre; B writes a new blurb; C computes 24/6.

The question

Distinguish operations.

Worked explanation

A predicts a category, B generates text, C calculates 4. Products may combine them.

Explain it in your own words

Which detail supports the answer, and why does it matter?

Write or draw your reasoning and evidence.

Trace the supplied case visually

Three operations in one book helper

Evidence or stage · Value or result

Evidence or stage: Classify

Value or result: Select a genre label

Evidence or stage: Generate

Value or result: Write a new blurb

Evidence or stage: Calculate

Value or result: 24 divided by 6 equals 4

A labels genre; B writes a new blurb; C computes 24/6.

Distinguish operations.

A predicts a category, B generates text, C calculates 4. Products may combine them.

14. Try a new case independently

Independent case: Analyse a new corpus independently

For prefix The signal is, counts are high 2, low 5, unknown 3. Compute all probabilities and the most frequent continuation. A current measurement reports high. State what the most frequent continuation does and does not establish.

My independent explanation, trace and limits

Attempt this case before feedback. Record any help. Earlier examples below are further practice; a case already practised with feedback cannot establish fresh transfer.

Week 27 Independent application

Close the worked explanation. Apply the idea to this different question without copying.

One model labels a review positive or negative; another writes a new review. Which classifies and which generates? Must generated text be true?

My answer, with a trace, calculation or supporting evidence

Write or draw your reasoning and evidence.

Create and test

Change one relevant detail. Predict the result before testing. Include a boundary or missing-information case when it helps reveal a limit.

My changed input and expected result

Write or draw your reasoning and evidence.

Connect to your project

Keep scoring predictive and explanations templated or carefully generated and verified.

The evidence I will keep

Write or draw your reasoning and evidence.

Your teacher has the independent answer and scoring criteria. Complete the matching workbook week to keep your practice evidence together.

Transfer practice from the original programme

A tool retrieves a book description, predicts a genre, then writes a new blurb. Identify each operation and give one distinct failure to check at each stage.

Explain the mechanism, evidence and one limit.

Use feedback to improve

Attempt the independent case before hints. Record whether you worked alone, used a hint or followed a model. After feedback, fix the first unsupported step and explain why it changed. A later check uses a changed case, not a copied answer.

15. Your lesson route

Start with this question: A fictional toy generator predicts the next word after “The box is”.

By the end, I can · How I will show it

By the end, I can: Distinguish discriminative prediction from generation by inspecting the requested output

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Calculate an empirical next-token distribution from counted continuations

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Separate next-token fit from factual support and downstream usefulness

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

Remember → watch and explain → try with help → investigate → try independently → keep project evidence. Your teacher may spread this lesson across several classes.

Keep your first prediction. Compare it with what the evidence shows and explain any change.

16. Finish and return

Close your notes. Explain this goal in your own words: Distinguish discriminative prediction from generation by inspecting the requested output

My deciding step, evidence and one remaining question

If you are unsure, return to the worked picture and explain one arrow. If you are ready, change an input or assumption and justify your prediction.

Save your workbook evidence. At the next review, try a different case before opening your old answer.

17. Supported practice, repair and portfolio

Use this supported practice once, in the browser or on paper. Keep the reserved independent assessment for a separate first attempt before teacher feedback.

A fictional toy generator predicts the next word after “The box is”.

Practice 1: explain the deciding evidence

Which labelled step matches this detail? Current observation says empty; frequent wording is not a present inventory reading.

1. Probabilities
2. Count table
3. Greedy continuation
4. Check reality

My selected conclusion, deciding fact and explanation

Hint if needed: Read the steps in the pictorial case. Match the supplied detail to the stage that performs or records it.

After your attempt, compare with the supported explanation: Check reality Current observation says empty; frequent wording is not a present inventory reading. The table predicts a continuation in a text exercise, not the contents of the current box. Generating the most frequent token is a model-defined decision. A factual inventory statement needs current evidence. The numerical table is intentionally small; real generative systems condition on richer context and do not simply consult these three counts.

Practice 2: explain the deciding evidence

Changed-case practice: Change the training table to full 2, empty 7, unknown 1. Which conclusion is supported by the supplied facts?

1. Greedy output becomes empty with probability 0.7.
2. This one example guarantees correct results for every future input.
3. We can decide the result without examining the changed input or the procedure.

My selected conclusion, deciding fact and explanation

Hint if needed: Use the changed condition and the deciding step in the pictorial trace. The answer must say what follows from those facts.

After your attempt, compare with the supported explanation: Greedy output becomes empty with probability 0.7. The continuation distribution changed; the independently observed box has not changed.

Repair a missing building block

Review Image and Voice Pipelines (ai-2/26). Goal: Explain the mechanism and terminology of speech pipelines and multimodal evidence. Short practice: Which stage first fails, and why should the faulty transcript not be silently overwritten?. Return to this week and explain what you changed.

Review Feedback Loops (ai-2/24). Goal: Implement a deterministic recommender. Short practice: What extra field would reveal why comparing raw clicks is misleading?. Return to this week and explain what you changed.

Review Similarity by Selected Traits (ai-2/21). Goal: Represent items with numeric features. Short practice:

Explain why B ceases to be a closest option without any candidate data changing.. Return to this week and explain what you changed.

Keep your completed work

Use Download completed work in the hosted lesson to export your answers, reflections, practice feedback and code-run evidence. Open the HTML copy to read or print it. Drafts stay in this browser when storage is available; clear drafts on a shared device.