

Dataset Cards

Project: Responsible Plant Classifier

1. Your mission

Your learning target

I can inspect row meaning missing values and repeated records.

I can document dataset source schema and limitations.

I can separate observed quality flags from justified corrections.

Remember before reading

Close your notes first. A rough first answer helps us choose the right starting point; it is not a score for your ability.

Without opening the old lesson, explain one key idea from Leakage. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Label Quality. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Classification or Regression?. Give an example and a mistake that the idea helps you avoid.

After trying, check the relevant lesson or ask your teacher. Change the part of your explanation that the evidence shows was wrong.

A dataset card records where data came from, why it exists, its fields, collection process, known gaps, risks, and prohibited uses.



LEARNING TARGET Write five dataset-card fields for a leaf dataset and explain why '1000 images' alone is insufficient.

Builder name and experiment date:

Readiness for the connected explanation

Why is elapsed trip duration in minutes a regression target even when recorded as whole numbers?

This week's end goal

By the end, I can explain dataset cards using a traced example and a stated limitation.

2. Notice the evidence

PICTURE ENTRY

Notice → Predict → Explain

Begin with visible evidence; do not explain more than the picture and information support.



MYSTERY The team must use dataset cards to solve a model-building problem. What information do they have? What important evidence is still missing?

Name three visible details without interpreting them.

Predict the data, model, or evaluation decision being made.

Identify one hidden assumption that could change the result.

State what evidence would confirm or reject your prediction.

My prediction and supporting evidence:

A closer explanation

A dataset can look like a neat table and still contain missing outcomes, repeated records or impossible values. Before fitting a model, inspect what one row represents and document where the table came from. A dataset card is a short record of its purpose, contents, collection, changes and limits. It helps another person decide whether the data suits a proposed use.

Our source is a teacher-authored fictional export called Robot practice v1, created for this lesson. Each row is intended to describe one trip, with one unique trip ID and elapsed duration in minutes. Durations may be zero or positive; None means unknown. All trips are described as indoor practice trips. No real robot measurements or personal records were collected, and there is no evidence about outdoor performance.

Trip ID · Duration in minutes · What is visible

Trip ID: A

Duration in minutes: 8

What is visible: A recorded duration

Trip ID: B

Duration in minutes: None

What is visible: An unknown duration

Trip ID: B

Duration in minutes: None

What is visible: The same ID appears again

Trip ID: C

Duration in minutes: -2

What is visible: A negative duration

Trip ID: D

Duration in minutes: 12

What is visible: A recorded duration

Keep this raw version unchanged. At this stage a flag means investigate, not automatically delete or replace. A repeated ID might be a copied row or a problem with the supposed one-row-per-trip definition. We need evidence before deciding what happened.

3. See how it works

ML MECHANISM

Input → **Model** → **Evidence**

Trace the information, fitted behavior, and evaluation proof separately.



IN

INPUT

Complete a dataset card and identify two uses the data cannot support.



ML

MODEL

A dataset card records where data came from, why it exists, its fields, collection process, known gaps, risks, and prohibited uses.

OUT

EVIDENCE

Complete a dataset card and identify two uses the data cannot support.

QUESTION Complete a dataset card and identify two uses the data cannot support.

MODEL IDEA A dataset card records where data came from, why it exists, its fields, collection process, known gaps, risks, and prohibited uses.

PROOF NEEDED Use an unseen case, visible calculation, error analysis, and a stated limitation.

UNDER THE HOOD Source: student-collected classroom plant photos with permission. Purpose: teach a low-stakes healthy-versus-spotted classifier.

Draw a fresh data → model → prediction → evaluation trace:

A closer explanation

```
rows = [("A", 8), ("B", None), ("B", None),
        ("C", -2), ("D", 12)]

ids = set()
missing = 0
negative = 0

for trip_id, minutes in rows:
    ids.add(trip_id)

    if minutes is None:
        missing = missing + 1

    elif minutes < 0:
        negative = negative + 1
```

```
print(len(rows), len(ids), missing, negative)
```

A set stores each distinct ID once. The output is 5 4 2 1: five rows, four distinct IDs, two rows with unknown duration and one row with a negative duration. There is one extra occurrence beyond the four distinct IDs. This is not a count of all rows in repeated-ID groups: both B rows belong to that group.

Turn an observation into a documented action

Finding · What it establishes · Next action

Finding: B appears twice

What it establishes: Identity needs review

Next action: Check export history

Finding: Two None values

What it establishes: 2 of 5 rows lack duration

Next action: Keep unknown explicit

Finding: C has minus 2

What it establishes: Invalid elapsed duration

Next action: Check source and units

Flags can overlap: both rows with ID B also have missing duration. Do not add flag counts as if they were separate groups.

The code assumes every row is a pair and each duration is a number or None. It is a small audit, not full schema validation. It does not check text values, units, source truth or whether IDs refer to real unique events. A successful run does not certify data quality.

Connect the picture and the explanation

Point to the input, the deciding step and the result in this week's visual. Explain the same step in ordinary words. A picture helps when you can say what each part represents.

4. Compare the cases

EVALUATION DESIGN

Four Tests, Four Kinds of Evidence

A system is not understood until it survives more than one comfortable example.



NORMAL CASE

Expected data from the intended setting



BOUNDARY CASE

Value near a threshold or class boundary



MISSING CASE

Unknown field, label, or necessary evidence



SHIFTED CASE

New subgroup, device, condition, or time

NORMAL Expected data and expected behavior.

BOUNDARY Value close to a threshold, ambiguous label, or uncertain prediction.

MISSING Necessary field, label, source, or measurement is unavailable.

SHIFTED New lighting, species, device, subgroup, or time condition.

RELEASE RULE Passing one case never proves generalization. Predefine expected behavior for all four categories.

Which case most threatens today's claim, and why?

Change one thing in the pictured case

Change: The dataset adds 6 outdoor photos of the same plants.

Prediction: Photo count becomes 24, plant count remains 6; lighting coverage changes.


Reason: Additional views broaden conditions without creating new independent plant identities.

5. Words that unlock the idea

VOCABULARY

Words You Can Calculate and Defend

A definition matters when it lets you interpret a new dataset, model, or result.



2

PROVENANCE

where data originated and changed

3

LIMITATION

condition weakening valid use

4

PROHIBITED USE

use outside evidence or
permission

Term · Precise meaning · My ML example

Term: dataset card

Precise meaning: structured documentation for a dataset

My ML example: _____

Term: provenance

Precise meaning: where data originated and changed

My ML example: _____

Term: limitation

Precise meaning: condition weakening valid use

My ML example: _____

Term: prohibited use

Precise meaning: use outside evidence or permission

My ML example: _____

MISCONCEPTION TO REPAIR A file name and list of columns are enough documentation for responsible reuse.

Repair the claim using a definition, calculation, or test record:

Words used in the closer explanation

Word · Meaning

Word: Provenance

Meaning: Where data came from and how it was created

Word: Schema

Meaning: The stated meaning types units and rules for fields

Word: Missing value

Meaning: An unknown or unavailable value rather than a measured zero

Word: Dataset card

Meaning: A concise record of dataset contents creation intended use and limitations

6. Follow a worked example

WORKED EVIDENCE

Reason Step by Step

Predict first. Then check each step. A correct answer without visible evidence cannot be audited.

WEEK 9 • WORKED EVIDENCE

- 1 Source: student-collected classroom plant photos with permission.
- 2 Purpose: teach a low-stakes healthy-versus-spotted classifier.
- 3 Gap: few evening-light photos and no outdoor disease diagnosis evidence.
- 4 Prohibit medical, agricultural, or pesticide decisions; require teacher review.

STEP 1 Source: student-collected classroom plant photos with permission.

STEP 2 Purpose: teach a low-stakes healthy-versus-spotted classifier.

STEP 3 Gap: few evening-light photos and no outdoor disease diagnosis evidence.

STEP 4 Prohibit medical, agricultural, or pesticide decisions; require teacher review.

CHECK Verify the denominator, unit, split, class definition, and data version before interpreting a score.

Recalculate, retrace, or restate the evidence independently:

A closer explanation

For Robot practice v1, record the purpose as practising inspection, the source as teacher-authored fictional data, and the row meaning as an intended single indoor trip. Record the five-row count, the ID field, minutes as the duration unit and None as unknown. List the duplicate-ID and invalid-duration findings. State that no cleaning has occurred, no train/test split is assigned, and no model has been evaluated.

Add intended use, excluded uses, known gaps, version and a contact or owner. Here the lesson teacher maintains the practice copy; its allowed classroom purpose is inspection practice. Do not describe it as validated navigation data. For a real external dataset, also verify its actual permission or licence and collection context; an absent field in the documentation is an unanswered question, not permission to invent a value.

What fraction of the five raw rows has unknown duration? Can you subtract two missing rows, one negative row and one extra ID occurrence from five to count distinct usable trips? Explain using the B rows.

Learn from the worked case

Cover the result before checking it. Say what is given, choose the procedure, and follow one step at a time. Then uncover the result and explain your first mismatch. Ask why a step is allowed, not only what number it produces.

A pictorial worked case: Dataset Cards

Build a dataset card that reveals what the photo count hides.

The facts we will use

Teacher-created dataset: 18 photos of 6 fictional plants, 3 views each, one tablet, indoor fluorescent light. Label rule is visible spot/no visible spot/review.

All inputs and outcomes in this worked example are invented classroom data; no real model run is claimed.

Dataset Cards

Start with the supplied facts

Follow each arrow. Ask which fact allows the next step.



1. Describe source

Record teacher-created fictional source, one tablet and indoor fluorescent lighting.



2. Describe units

18 photo rows come from only 6 plant identities.



3. Record limits

Split by plant; state that outdoor use and disease diagnosis have not been evaluated.

CHANGE ONE THING

The dataset adds 6 outdoor photos of the same plants.

Explain the changed result using the matching lesson.

Read the picture one step at a time

Step · What happens and why

Step: 1. Describe source

What happens and why: Record teacher-created fictional source, one tablet and indoor fluorescent lighting.

Step: 2. Describe units

What happens and why: 18 photo rows come from only 6 plant identities.

Step: 3. Record limits

What happens and why: Split by plant; state that outdoor use and disease diagnosis have not been evaluated.

Match each pictured step to the supplied facts. This is a worked teaching case, not an independent assessment.

Why this result follows

Read the picture in order. Split by plant; state that outdoor use and disease diagnosis have not been evaluated. The arrows show which recorded input or intermediate result each decision uses. Explain the deciding step aloud before checking the changed case; pointing to the final answer alone does not show how it follows.

What this example does not establish

A dataset card documents evidence and gaps; writing one does not fix biased sampling or inaccurate labels.

Read the labelled picture: What two counts should the revised card show? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.

Change one thing: The dataset adds 6 outdoor photos of the same plants. Predict the result before checking the explanation. What remains the same?

7. Find and repair a mistake

ERROR DETECTIVE

Compare → Diagnose → Repair

Do not patch the answer. Find the earliest unsupported assumption, data problem, or experiment error.



COMPARE • DIAGNOSE • REPAIR

Tempting claim: A file name and list of columns are enough documentation for responsible reuse. Inspect the dataset, split, calculation, error type, subgroup, or experiment record; identify one testable cause; change one factor; rerun the same evidence.

Evidence record · Expected · Actual · Likely cause / next test

Evidence record: normal case

Expected: passes

Actual: _____

Likely cause / next test: _____

Evidence record: boundary case

Expected: defined behavior

Actual: _____

Likely cause / next test: _____

Evidence record: missing case

Expected: safe fallback

Actual: _____

Likely cause / next test: _____

Evidence record: shifted case

Expected: measured limit

Actual: _____

Likely cause / next test: _____

CLAIM TO CHALLENGE A file name and list of columns are enough documentation for responsible reuse.

Diagnose one cause and name the cleanest next experiment:

A closer explanation

A complete-looking card does not prove the table is accurate or suitable for a model. It makes claims and unanswered questions reviewable. Report raw counts separately from any later accepted subset, with explicit rules for missing, invalid and repeated records. Do not silently change denominators after cleaning.

For this exercise, zero is allowed and unknown is not zero. That is a stated schema choice, not a universal rule for all duration datasets. In a real task, measurement resolution and event definitions determine which values make sense. Dataset inspection also does not replace protecting held out data from model choices.

A card says the dataset is ready because it has no empty cells. Name two checks that claim misses, then write a more accurate statement about Robot practice v1.

8. Practise together

ML STUDIO

Specify Before You Run

A trustworthy result can be reproduced from the recorded data, code, settings, and evaluation rule.



FREEZE THE EXPERIMENT

Record question, intended use, data version, grouped split, preprocessing, features, model settings, seed, metric, predictions, errors, subgroup results, and human review before claiming improvement.

Experiment record

QUESTION + intended use + success threshold

DATA version + grouped split + leakage check

CODE/preprocessing + features + model + settings + seed

METRICS + positive class + units + denominators

PREDICTIONS + errors + subgroup slices + limitations

HUMAN reviewer + release, revise, limit, or stop decision

Write the minimum record needed to reproduce today's result:

A closer explanation

A learner replaces both None values with 0 and removes the negative sign from -2. The table now passes a simple nonnegative check, but those edits invent durations. Keep the original values, record the reason for each flag, and check the source. If a correction is supported, save a new version with the original value, corrected value and reason. If not, keep the uncertainty visible.

Another learner says all robot trips are covered because every row says indoor. A dataset card should instead state that outdoor trips are absent and that the data is fictional. Extra copies of existing rows do not create new conditions or new independent evidence.

Practise with less help: Interpret the audit counts

Use workbook W2 for the connected example “Dataset inspection and cards”. First fill the input and rule boxes. Try the middle steps before asking for a hint. After a hint, try the next step yourself.

Plan · Do · Check

Plan: What is given? Which rule or procedure applies?

Do: Record each intermediate step.

Check: Does the result follow? What remains unknown?

Explain your trace to a partner. Your partner points to one step to check. Swap roles on the next case. Each person writes their own explanation.

Check your reasoning

Use the worked case for Dataset inspection and cards. Proposed explanation: “A complete card certifies data quality”. Decide whether this explanation is supported. Point to the step or supplied fact that decides; explain your repair if needed.

Answer on your own before discussion. If you are unsure, say which step you need help with.

Explain the picture

What two counts should the revised card show? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.

My explanation using a pictured step and the given facts

This is supported practice. Keep the separately named independent case for an attempt before feedback.

9. Run the investigation

EVIDENCE LAB

Predict → Run → Calculate → Interpret

Keep the test conditions fixed, record every result, and change only one planned factor.

LAB

PREDICT → RUN → CALCULATE → INTERPRET

Dataset Card Review: inspect two incomplete cards, find missing provenance or risks, and write a defensible replacement.

1. Write the hypothesis and expected evidence before running.
2. Dataset Card Review: inspect two incomplete cards, find missing provenance or risks, and write a defensible replacement.
3. Run one normal, one boundary, one missing, and one shifted case.
4. Calculate the metric from raw counts or row-level errors.
5. Interpret what changed, what did not, and what remains unsafe or unknown.

Case · Prediction · Actual · Calculation / evidence · Meaning / next test

Case: Normal

Prediction: _____

Actual: _____

Calculation / evidence: _____

Meaning / next test: _____

Case: Boundary

Prediction: _____

Actual: _____

Calculation / evidence: _____

Meaning / next test: _____

Case: Missing

Prediction: _____

Actual: _____

Calculation / evidence: _____

Meaning / next test: _____

Case: Shifted

Prediction: _____

Actual: _____

Calculation / evidence: _____

Meaning / next test: _____

Plan → try → check


Before trying: what do I expect, and why? While trying: which step could first go wrong? Afterwards: what did I actually observe, and what should change? Keep “not run” if you only traced the procedure.

10. Build your understanding

MASTERY

From Knowing the Word to Owning the Idea

Move upward only when you can produce evidence without copying the example.



ML BUILDER MASTERY

1 Define • 2 Calculate • 3 Trace code • 4 Run experiment • 5 Diagnose errors • 6 Improve one cause • 7 Defend limits

DEFINE Use all four terms precisely: dataset card, provenance, limitation, prohibited use.

CALCULATE Reproduce the worked evidence with units and denominator.

TRACE Explain the data, code or decision path.

DIAGNOSE A file name and list of columns are enough documentation for responsible reuse.

TRANSFER Solve a fresh dataset or model case.

IMPROVE Change one cause and rerun fixed evidence.

DEFEND State intended use, metric, limitation, and human responsibility.

MASTERY RULE A coached calculation becomes mastery only after a different case is completed independently.

11. Connect your project

CUMULATIVE PROJECT

Responsible Plant Classifier

Every week adds one reviewable artifact to the same evidence-based capstone.



RESPONSIBLE PLANT CLASSIFIER

Publish the Plant Classifier dataset card before model experiments begin; update it when data changes.

THIS WEEK'S CHECKPOINT Publish the Plant Classifier dataset card before model experiments begin; update it when data changes.

Which requirement or risk does this artifact address?

Which dataset, split, code, and metric version does it use?

What normal and shifted cases were tested?

What remains uncertain or weak?

What decision still belongs to a teacher or qualified human?

RESPONSIBLE RELEASE Never use this classroom prototype for medical, pesticide, agricultural, or other consequential decisions.

My evidence-linked project decision:

Evidence for your project

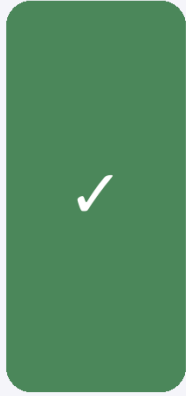
For your Responsible Plant Classifier, save the input, procedure, observed result and limitation of this check. Label this worked example as practice; use a different, previously unreviewed case when checking independent understanding.

12. Show what you know

INDEPENDENT PROOF

Explain • Calculate • Transfer • Defend

Complete this with a fresh case after guided practice and compare it with the mastery criteria.



INDEPENDENT ML BUILDER PROOF

Complete a dataset card and identify two uses the data cannot support.

REQUIRED PROOF Complete a dataset card and identify two uses the data cannot support.

- ☐ Define the task, feature/target roles, and intended-use boundary.
- ☐ Show the calculation, code trace, or experiment record.
- ☐ Test a normal case and one revealing boundary, missing, or shifted case.
- ☐ Diagnose one failure without changing several factors at once.
- ☐ Name a limitation, fairness concern, and the human-review action.

Independent evidence and defense:

My level: ☐ Not yet ☐ With one hint ☐ Independently ☐ I can teach and defend it

NEXT BRIDGE Carry the experiment evidence into: Baselines First.

13. Study a complete case

Week 9 Worked case

Use the supplied details to practise the weekly concept before attempting the independent question. This case extends the illustrated lesson.

The situation

A dataset card lists fields but omits source and intended use.

The question

Add information needed to assess suitability.

Worked explanation

Provenance, permissions, collection method/time/conditions, label definitions, coverage, exclusions, missingness, use and limitations. Fields alone are insufficient.

Explain it in your own words

Which detail supports the answer, and why does it matter?

Write or draw your reasoning and evidence.

Trace the supplied case visually

A dataset card explains suitability

Evidence or stage · Value or result

Evidence or stage: Origin

Value or result: Source, permission and collection

Evidence or stage: Meaning

Value or result: Fields and operational labels

Evidence or stage: Coverage

Value or result: Conditions, exclusions and missingness

Evidence or stage: Use

Value or result: Intended tasks and limitations

A dataset card lists fields but omits source and intended use.

Add information needed to assess suitability.

Provenance, permissions, collection method/time/conditions, label definitions, coverage, exclusions, missingness, use and limitations. Fields alone are insufficient.

14. Try a new case independently

Independent case: Changed independent case

Write key dataset-card entries for 20 photos: five plants, four views each, one phone, one indoor setting. Growth labels have not been supplied.

My independent explanation, trace and limits

Attempt this case before feedback. Record any help. Earlier examples below are further practice; a case already practised with feedback cannot establish fresh transfer.

Week 9 Independent application

Close the worked explanation. Apply the idea to this different question without copying.
A dataset card lists 500 images but omits collection conditions and intended use. Add two essential entries and explain their value.

My answer, with a trace, calculation or supporting evidence

Write or draw your reasoning and evidence.

Create and test

Change one relevant detail. Predict the result before testing. Include a boundary or missing-information case when it helps reveal a limit.

My changed input and expected result

Write or draw your reasoning and evidence.

Connect to your project

Publish the Plant Classifier dataset card before model experiments begin; update it when data changes.

The evidence I will keep

Write or draw your reasoning and evidence.

Your teacher has the independent answer and scoring criteria. Complete the matching workbook week to keep your practice evidence together.

Transfer practice from the original programme

Write five dataset-card fields for a leaf dataset and explain why '1000 images' alone is insufficient.

Explain the mechanism, evidence and one limit.

Use feedback to improve

Attempt the independent case before hints. Record whether you worked alone, used a hint or followed a model. After feedback, fix the first unsupported step and explain why it changed. A later check uses a changed case, not a copied answer.

15. Your lesson route

Start with this question: Build a dataset card that reveals what the photo count hides.

By the end, I can · How I will show it

By the end, I can: Inspect row meaning missing values and repeated records

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Document dataset source schema and limitations

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Separate observed quality flags from justified corrections

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

Remember → watch and explain → try with help → investigate → try independently → keep project evidence. Your teacher may spread this lesson across several classes.

Keep your first prediction. Compare it with what the evidence shows and explain any change.

16. Finish and return

Close your notes. Explain this goal in your own words: Inspect row meaning missing values and repeated records

My deciding step, evidence and one remaining question

If you are unsure, return to the worked picture and explain one arrow. If you are ready, change an input or assumption and justify your prediction.

Save your workbook evidence. At the next review, try a different case before opening your old answer.

17. Supported practice, repair and portfolio

Use this supported practice once, in the browser or on paper. Keep the reserved independent assessment for a separate first attempt before teacher feedback.

Build a dataset card that reveals what the photo count hides.

Practice 1: explain the deciding evidence

Which labelled step matches this detail? Split by plant; state that outdoor use and disease diagnosis have not been evaluated.

1. Record limits
2. Describe units
3. Describe source

My selected conclusion, deciding fact and explanation

Hint if needed: Read the steps in the pictorial case. Match the supplied detail to the stage that performs or records it.

After your attempt, compare with the supported explanation: Record limits Split by plant; state that outdoor use and disease diagnosis have not been evaluated. Read the picture in order. Split by plant; state that outdoor use and disease diagnosis have not been evaluated. The arrows show which recorded input or intermediate result each decision uses. Explain the deciding step aloud before checking the changed case; pointing to the final answer alone does not show how it follows.

Practice 2: explain the deciding evidence

Changed-case practice: The dataset adds 6 outdoor photos of the same plants. Which conclusion is supported by the supplied facts?

1. We can decide the result without examining the changed input or the procedure.
2. Photo count becomes 24, plant count remains 6; lighting coverage changes.
3. This one example guarantees correct results for every future input.

My selected conclusion, deciding fact and explanation

Hint if needed: Use the changed condition and the deciding step in the pictorial trace. The answer must say what follows from those facts.

After your attempt, compare with the supported explanation: Photo count becomes 24, plant count remains 6; lighting coverage changes. Additional views broaden conditions without creating new independent plant identities.

Repair a missing building block

Review Leakage (ai-4/8). Goal: Choose a split unit that matches the intended prediction claim. Short practice: How many independently new plant identities are in the original test? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.. Return to this week and explain what you changed.

Review Label Quality (ai-4/6). Goal: Write an observable label rule with boundaries and a review option. Short practice: Why should two annotators not vote away Z's missing evidence? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.. Return to this week and explain what you changed.

Review Classification or Regression? (ai-4/3). Goal: Identify a task from the meaning of its target. Short practice: Did regression MAE improve when the boundary changed? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.. Return to this week and explain what you changed.

Predict, run and debug

The browser lab lets you edit the supported Python examples already included in this lesson. Predict first, run, inspect the actual output or error, change one input, then run again. Keep code, predictions, actual results and a reasoned revision together. On paper, label your result as a trace rather than an execution.

Keep your completed work

Use Download completed work in the hosted lesson to export your answers, reflections, practice feedback and code-run evidence. Open the HTML copy to read or print it. Drafts stay in this browser when storage is available; clear drafts on a shared device.