

Similarity and normalisation

Project: Image-pattern Network Investigation

1. Your mission

Your learning target

I can calculate vector length and normalise a nonzero vector.

I can compare squared distance with cosine similarity and explain a changed ranking.

I can handle zero vectors explicitly and justify what information a comparison keeps.

Remember before reading

Close your notes first. A rough first answer helps us choose the right starting point; it is not a score for your ability.

Without opening the old lesson, explain one key idea from Embeddings and neighbourhoods. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Augmentation and leakage. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Convolution as a local operation. Give an example and a mistake that the idea helps you avoid.

After trying, check the relevant lesson or ask your teacher. Change the part of your explanation that the evidence shows was wrong.

This week you will practise:

Calculate vector length and normalise a nonzero vector

Compare squared distance with cosine similarity and explain a changed ranking

Handle zero vectors explicitly and justify what information a comparison keeps

Retrieve what you already know

For query [3, 4] and candidate [3, 0], calculate squared Euclidean distance. Does a smaller distance prove equivalent meaning?

My starting explanation

This week's end goal

By the end, I can explain similarity and normalisation using a traced example and a stated limitation.

2. Notice the evidence

Last week, a neighbourhood depended on a chosen distance rule. This week, compare the same fictional vectors in two ways. Squared distance rewards nearby coordinates. Cosine similarity compares direction and ignores positive changes in length. Neither rule is automatically right for every task.

Vector · Coordinates · Length

Vector: Query Q

Coordinates: [3, 4]

Length: 5

Vector: Candidate A

Coordinates: [6, 8]

Length: 10

Vector: Candidate B

Coordinates: [3, 0]

Length: 3

Vector: Candidate C

Coordinates: [-3, -4]

Length: 5

For Q, length = square root of (3 squared + 4 squared) = square root of 25 = 5. Divide both components by 5 to get its unit vector [0.6, 0.8]. Its length is one because 0.6 squared + 0.8 squared = 1. A becomes the same unit vector. C becomes [-0.6, -0.8], pointing in the opposite direction.

These supplied coordinates are for arithmetic practice. They are not learned language embeddings. Normalising each row keeps its direction but removes its original length unless that length is saved separately.

List the given facts and one fact that is not supplied.

3. See how it works

Multiply matching components and add to obtain the dot product. Cosine similarity equals $\text{dot}(Q, \text{candidate})$ divided by $(\text{length}(Q) \times \text{length}(\text{candidate}))$. Higher cosine means a closer direction. For A, the dot product is 3 times 6 + 4 times 8 = 50, so cosine = $50 / (5 \times 10) = 1$. For B it is $9 / 15 = 0.6$; for C it is $-25 / 25 = -1$.

```
from math import sqrt

def unit(v):
    if not v:
        raise ValueError("Need a nonempty vector")
    length = sqrt(sum(x*x for x in v))
    if length == 0:
        raise ValueError("Zero vector has no direction")
    return [x / length for x in v]

def cosine(a, b):
    if len(a) != len(b):
        raise ValueError("Need matching dimensions")
    ua, ub = unit(a), unit(b)
    return sum(x*y for x, y in zip(ua, ub))

q = [3, 4]
for name, v in [("A", [6, 8]), ("B", [3, 0]),
                ("C", [-3, -4])]:
    print(name, round(cosine(q, v), 6))
```

Trace before running: A gives 1, B gives 0.6 and C gives -1. This classroom code assumes finite, moderate numeric inputs. It rejects empty, unequal-length and zero vectors. Keep full precision for

decisions; rounding here is only for display. Cosine lies between -1 and 1 in exact arithmetic and is not a probability.

Connect the picture and the explanation

Point to the input, the deciding step and the result in this week's visual. Explain the same step in ordinary words. A picture helps when you can say what each part represents.

4. Compare the cases

Compare a familiar case with a changed case

For query Q, write the raw squared distances and cosines to A, B and C from the lesson. Name the winner under each rule. Calculate squared distances after normalising both sides and check the relationship to cosine.

Keep the procedure fixed. Explain what changed in the inputs and result.

Change one thing in the pictured case

Change: Replace P with $[0, 0]$.

Prediction: Cosine and unit normalisation for P are undefined.

Reason: The zero vector has norm 0, so dividing by its norm is invalid.

5. Words that unlock the idea

Word · Meaning

Word: Vector length

Meaning: The square root of the sum of squared components, also called the L2 norm

Word: L2 normalisation

Meaning: Dividing every component of a nonzero vector by its length

Word: Cosine similarity

Meaning: The dot product of two nonzero vectors divided by the product of their lengths

Word: Unit vector

Meaning: A vector with length one

Use two terms to explain a step in this week's example.

6. Follow a worked example

Compare the same candidates using two rules

Candidate · Raw squared distance · Cosine similarity

Candidate: A [6, 8]

Raw squared distance: 25

Cosine similarity: 1

Candidate: B [3, 0]

Raw squared distance: 16

Cosine similarity: 0.6

Candidate: C [-3, -4]

Raw squared distance: 100

Cosine similarity: -1

For Q [3, 4], raw squared distance selects B by minimising the score. Cosine selects A by maximising the score. The candidates stayed the same; the comparison rule changed.

After normalising both Q and every candidate, squared distances to A, B and C are 0, 0.8 and 4. For unit vectors, squared distance = $2 - 2 \times \text{cosine}$. Thus the smallest normalised distance and largest cosine produce the same ranking, including ties, in exact arithmetic.

Add candidate D [0, 5]. Calculate its raw squared distance and cosine from Q [3, 4]. Compare D with A under each rule.

Learn from the worked case

Cover the result before checking it. Say what is given, choose the procedure, and follow one step at a time. Then uncover the result and explain your first mismatch. Ask why a step is allowed, not only what number it produces.

A pictorial worked case: Similarity and normalisation


Compare angular alignment with raw distance on the same vectors.

The facts we will use

Query [2, 0]; P=[8, 0]; Q=[2, 1]. Nonzero vectors only; compare squared Euclidean distance and cosine.

All inputs and outcomes in this worked example are invented classroom data; no real model run is claimed.

Raw proximity differs from angular alignment



Horizontal: first vector coordinate; vertical: second vector coordinate

1: P at (8, 0)

2: Q at (2, 1)

Cross: query at (2, 0)

Distance rule: Compare squared Euclidean distance and cosine

Similarity and normalisation

Start with the supplied facts

Follow each arrow. Ask which fact allows the next step.



1. Raw distances

P distance 36; Q distance 1, so Q is closer in raw coordinates.



2. Angular similarity

$\text{Cosine}P=1$; $\text{cosine}Q=2/\sqrt{5}\approx 0.8944$, so P aligns more closely.



3. Normalise

Unit query $[1, 0]$, P $[1, 0]$, Q $[2/\sqrt{5}, 1/\sqrt{5}]$; normalised distance P0.

CHANGE ONE THING

Replace P with $[0, 0]$.

Explain the changed result using the matching lesson.

Read the picture one step at a time

Step · What happens and why

Step: 1. Raw distances

What happens and why: P distance 36; Q distance 1, so Q is closer in raw coordinates.

Step: 2. Angular similarity

What happens and why: $\text{Cosine}P=1$; $\text{cosine}Q=2/\sqrt{5}\approx 0.8944$, so P aligns more closely.

Step: 3. Normalise

What happens and why: Unit query $[1, 0]$, $P [1, 0]$, $Q [2/\sqrt{5}, 1/\sqrt{5}]$; normalised distance $P0$.

Match each pictured step to the supplied facts. This is a worked teaching case, not an independent assessment.

Why this result follows

Raw Euclidean distance includes differences in both direction and magnitude. Cosine compares direction using the vector lengths, so a long vector can align perfectly while being far away in raw coordinates. Declare which question your similarity measure answers. A zero vector has no direction and cannot be divided by its norm for cosine or unit normalisation.

What this example does not establish

Neither metric is universally right; choose it from the task and handle zero vectors explicitly.

Read the labelled picture: Why do the original two metrics select different winners? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.

Change one thing: Replace P with $[0, 0]$. Predict the result before checking the explanation. What remains the same?

7. Find and repair a mistake

Suppose $[3, 4]$ records counts of two event types, while $[6, 8]$ records twice as many of each. Both have the same proportions and direction, but totals are 7 and 14. Normalisation makes their vectors identical. That may suit a question about proportions, but it removes information needed to compare total activity. For learned embeddings, the usefulness of length must be tested for the task.

The zero vector $[0, 0]$ has no direction, so its cosine is undefined. Do not silently assign zero similarity: zero cosine already means perpendicular nonzero vectors. This lesson rejects a zero query and excludes zero candidates with an explicit report. If no valid candidates remain, return no comparison. PAD entries stay excluded by the mask; a real zero input needs its own policy.

Positive scaling preserves cosine: Q and 2 times Q have the same direction. Negative scaling reverses direction; Q and $-Q$ have cosine -1 . Candidate labels, self-exclusion and a stated floating-point tie tolerance are still required when retrieving neighbours.

Locate the earliest unsupported step

A candidate set contains $Z [0, 0]$ and $B [3, 0]$. With nonzero query Q , what happens under this lesson's zero-vector policy? What if Q is zero, or Z is the only candidate? Explain why replacing an undefined cosine with zero is misleading.

State one repair and the checks you would rerun.

8. Practise together

For Q [3, 4], A [6, 8] and C [-3, -4], give each length and unit vector. Explain one distinction that normalisation removes and one that it preserves.

Add candidate D [0, 5]. Calculate its raw squared distance and cosine from Q [3, 4]. Compare D with A under each rule.

Explain your prediction to a partner before comparing results.

Practise with less help: Compare the rankings

Use workbook W2 for the connected example “Similarity and normalisation”. First fill the input and rule boxes. Try the middle steps before asking for a hint. After a hint, try the next step yourself.

Plan · Do · Check

Plan: What is given? Which rule or procedure applies?

Do: Record each intermediate step.

Check: Does the result follow? What remains unknown?

Explain your trace to a partner. Your partner points to one step to check. Swap roles on the next case. Each person writes their own explanation.

Check your reasoning

Use the worked case for Similarity and normalisation. Proposed explanation: "Cosine is a probability". Decide whether this explanation is supported. Point to the step or supplied fact that decides; explain your repair if needed.

Answer on your own before discussion. If you are unsure, say which step you need help with.

Explain the picture

Why do the original two metrics select different winners? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.

My explanation using a pictured step and the given facts

This is supported practice. Keep the separately named independent case for an attempt before feedback.

9. Run the investigation

Predict, investigate and preserve the evidence

For query Q, write the raw squared distances and cosines to A, B and C from the lesson. Name the winner under each rule. Calculate squared distances after normalising both sides and check the relationship to cosine.

```
from math import sqrt

def unit(v):
    if not v:
        raise ValueError("Need a nonempty vector")

    length = sqrt(sum(x*x for x in v))

    if length == 0:
        raise ValueError("Zero vector has no direction")
```

```

    return [x / length for x in v]

def cosine(a, b):
    if len(a) != len(b):
        raise ValueError("Need matching dimensions")

    ua, ub = unit(a), unit(b)

    return sum(x*y for x, y in zip(ua, ub))

q = [3, 4]

for name, v in [("A", [6, 8]), ("B", [3, 0]),
                ("C", [-3, -4])]:
    print(name, round(cosine(q, v), 6))

```

Record · Your evidence

Record: Given input and fixed rule

Your evidence: _____

Record: Prediction before checking

Your evidence: _____

Record: Actual result, or “not run”

Your evidence: _____

Record: First mismatch and one explanation

Your evidence: _____

Record: Boundary or missing input and expected response

Your evidence: _____

On paper, trace the supplied procedure. If code is provided, predict before running it in a local practice file. Never write an expected answer as though it were an observed run.

Plan → try → check

Before trying: what do I expect, and why? While trying: which step could first go wrong? Afterwards: what did I actually observe, and what should change? Keep “not run” if you only traced the procedure.

10. Build your understanding

Explain → apply → challenge

For $Q [3, 4]$, $A [6, 8]$ and $C [-3, -4]$, give each length and unit vector. Explain one distinction that normalisation removes and one that it preserves.

A candidate set contains $Z [0, 0]$ and $B [3, 0]$. With nonzero query Q , what happens under this lesson’s zero-vector policy? What if Q is zero, or Z is the only candidate? Explain why replacing an undefined cosine with zero is misleading.

For event-count vectors $[3, 4]$ and $[6, 8]$, what useful information can normalisation remove? A team chooses cosine because it improves results on the final evaluation set, then calls that same set independent evidence. Explain how to repair the evaluation plan.

Explain which evidence supports your conclusion and which claim remains unproved.

11. Connect your project

Add a labelled similarity and normalisation evidence sheet to Image-pattern Network Investigation. Show how you can calculate vector length and normalise a nonzero vector. Use the supplied fictional case and keep its inputs, intermediate steps, calculation or prediction, and limitation. Connect this record to last week’s record; a matching answer without a trace is insufficient.

Project evidence · Record

Project evidence: Component and version

Record: _____

Project evidence: Inputs permitted at prediction time

Record: _____

Project evidence: Trace and comparison with earlier work

Record: _____

Project evidence: Failure, limit and human review

Record: _____

Evidence for your project

For your Image-pattern Network Investigation, save the input, procedure, observed result and limitation of this check. Label this worked example as practice; use a different, previously unreviewed case when checking independent understanding.

12. Show what you know

L2 normalisation divides each individual vector by its own length; it does not estimate a mean or scale from a dataset. Feature standardisation is a different operation and must fit its statistics on training data only. Choose between raw distance and cosine using the task, training and validation evidence. Do not repeatedly change the rule after seeing final evaluation results.

A teammate says a cosine of 1 proves two records mean the same thing, and normalisation always improves a model. Use A and Q to challenge both claims. What must be checked before adopting the rule?

Record the representation version, comparison rule, normalisation, zero-vector handling, candidate set and tie tolerance. Preserve a baseline and evaluate on the same reserved examples. Next week, transfer learning asks when a representation learned for one task can help another.

For event-count vectors [3, 4] and [6, 8], what useful information can normalisation remove? A team chooses cosine because it improves results on the final evaluation set, then calls that same set independent evidence. Explain how to repair the evaluation plan.

My limited claim and a useful next check

13. Study a complete case

Complete evidence record

Compare the same candidates using two rules

Candidate · Raw squared distance · Cosine similarity

Candidate: A [6, 8]

Raw squared distance: 25

Cosine similarity: 1

Candidate: B [3, 0]

Raw squared distance: 16

Cosine similarity: 0.6

Candidate: C [-3, -4]

Raw squared distance: 100

Cosine similarity: -1

For Q [3, 4], raw squared distance selects B by minimising the score. Cosine selects A by maximising the score. The candidates stayed the same; the comparison rule changed.

Follow the input through the stated procedure to its result. Cover the result first, predict it, then explain every intermediate step. The text/table accompanies the figure so the example can be read without colour.

My complete trace of the worked case

14. Try a new case independently

Independent case: Compare a new case independently

Use query [5, 0] and candidates P [10, 0], R [4, 3] and S [0, 5]. Calculate all raw squared distances and cosines, name each winner, and give squared distances after normalising both query and candidates. Explain the changed ranking.

My independent explanation, trace and limits

Attempt this case before feedback. Record any help. Earlier examples below are further practice; a case already practised with feedback cannot establish fresh transfer.

Attempt this without hints

Use query [5, 0] and candidates P [10, 0], R [4, 3] and S [0, 5]. Calculate all raw squared distances and cosines, name each winner, and give squared distances after normalising both query and candidates. Explain the changed ranking.

My independent procedure, result and limitation

Record support used: none, hint or teacher explanation. After feedback, a different case is needed to demonstrate independence.

Use feedback to improve

Attempt the independent case before hints. Record whether you worked alone, used a hint or followed a model. After feedback, fix the first unsupported step and explain why it changed. A later check uses a changed case, not a copied answer.

15. Your lesson route

Start with this question: Compare angular alignment with raw distance on the same vectors.

By the end, I can · How I will show it

By the end, I can: Calculate vector length and normalise a nonzero vector

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Compare squared distance with cosine similarity and explain a changed ranking

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Handle zero vectors explicitly and justify what information a comparison keeps

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

Remember → watch and explain → try with help → investigate → try independently → keep project evidence. Your teacher may spread this lesson across several classes.

Keep your first prediction. Compare it with what the evidence shows and explain any change.

16. Finish and return

Close your notes. Explain this goal in your own words: Calculate vector length and normalise a nonzero vector

My deciding step, evidence and one remaining question

If you are unsure, return to the worked picture and explain one arrow. If you are ready, change an input or assumption and justify your prediction.

Save your workbook evidence. At the next review, try a different case before opening your old answer.

17. Supported practice, repair and portfolio

Use this supported practice once, in the browser or on paper. Keep the reserved independent assessment for a separate first attempt before teacher feedback.

Compare angular alignment with raw distance on the same vectors.

Practice 1: explain the deciding evidence

Which labelled step matches this detail? Unit query $[1, 0]$, P $[1, 0]$, Q $[2/\sqrt{5}, 1/\sqrt{5}]$; normalised distanceP0.

1. Angular similarity
2. Normalise
3. Raw distances

My selected conclusion, deciding fact and explanation

Hint if needed: Read the steps in the pictorial case. Match the supplied detail to the stage that performs or records it.

After your attempt, compare with the supported explanation: Normalise Unit query $[1, 0]$, P $[1, 0]$, Q $[2/\sqrt{5}, 1/\sqrt{5}]$; normalised distanceP0. Raw Euclidean distance includes differences in both direction and magnitude. Cosine compares direction using the vector lengths, so a long vector can align perfectly while being far away in raw coordinates. Declare which question your similarity measure answers. A zero vector has no direction and cannot be divided by its norm for cosine or unit normalisation.

Practice 2: explain the deciding evidence

Changed-case practice: Replace P with $[0, 0]$. Which conclusion is supported by the supplied facts?

1. Cosine and unit normalisation for P are undefined.
2. This one example guarantees correct results for every future input.
3. We can decide the result without examining the changed input or the procedure.

My selected conclusion, deciding fact and explanation

Hint if needed: Use the changed condition and the deciding step in the pictorial trace. The answer must say what follows from those facts.

After your attempt, compare with the supported explanation: Cosine and unit normalisation for P are undefined. The zero vector has norm 0, so dividing by its norm is invalid.

Repair a missing building block

Review Embeddings and neighbourhoods (ai-5/26). Goal: Look up a vector using a token ID and preserve sequence shape. Short practice: What would justify claiming nearby vectors mean related topics? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.. Return to this week and explain what you changed.

Review Augmentation and leakage (ai-5/24). Goal: Check whether a transformation preserves the task label. Short practice: How many independent source groups are in the final test? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.. Return to this week and explain what you changed.

Review Convolution as a local operation (ai-5/21). Goal: Compute a local weighted sum at each valid image position. Short practice: Why is flipping the kernel an implementation change here? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.. Return to this week and explain what you changed.

Predict, run and debug

The browser lab lets you edit the supported Python examples already included in this lesson. Predict first, run, inspect the actual output or error, change one input, then run again. Keep code, predictions, actual results and a reasoned revision together. On paper, label your result as a trace rather than an execution.

Keep your completed work

Use Download completed work in the hosted lesson to export your answers, reflections, practice feedback and code-run evidence. Open the HTML copy to read or print it. Drafts stay in this browser when storage is available; clear drafts on a shared device.