

Dataset shift in neural models

Project: Image-pattern Network Investigation

1. Your mission

Your learning target

I can explain how a changed input mix can change overall accuracy.

I can calculate group and overall results with the correct denominators.

I can plan a shift investigation while preserving independent evaluation.

Remember before reading

Close your notes first. A rough first answer helps us choose the right starting point; it is not a score for your ability.

Without opening the old lesson, explain one key idea from Uncertainty and failure analysis. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Similarity and normalisation. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Augmentation and leakage. Give an example and a mistake that the idea helps you avoid.

After trying, check the relevant lesson or ask your teacher. Change the part of your explanation that the evidence shows was wrong.

This week you will practise:

Explain how a changed input mix can change overall accuracy

Calculate group and overall results with the correct denominators

Plan a shift investigation while preserving independent evaluation

Retrieve what you already know

Can a high confidence score establish that a new input resembles the development data or that its prediction is correct?

My starting explanation

This week's end goal

By the end, I can explain dataset shift in neural models using a traced example and a stated limitation.

2. Notice the evidence

A neural image classifier developed using mostly clear pictures may later receive mostly dim pictures. Its learned weights can stay unchanged while its overall results change. Dataset shift concerns changes in the data encountered, including inputs, label frequencies or the relationship between inputs and labels. A shift does not always reduce accuracy, and an accuracy change alone does not identify its cause.

Use the fictional audit below. Clear and dim are disjoint groups defined before the audit, and every case belongs to exactly one group. The same model checkpoint, preprocessing and prediction threshold are used throughout. Counts are constructed teaching data, not results from a deployed model. Both cohorts have checked reference labels.

Cohort · Group · Cases · Correct

Cohort: Development audit

Group: Clear

Cases: 80

Correct: 72

Cohort: Development audit

Group: Dim

Cases: 20

Correct: 10

Cohort: New-use audit

Group: Clear

Cases: 20

Correct: 18

Cohort: New-use audit

Group: Dim

Cases: 80

Correct: 40

Clear accuracy is $72 / 80 = 90$ percent in development and $18 / 20 = 90$ percent in new use. Dim accuracy is $10 / 20 = 50$ percent and $40 / 80 = 50$ percent. The observed group rates are unchanged, but the group proportions are different.

List the given facts and one fact that is not supplied.

3. See how it works

Development accuracy is $(72 + 10) / (80 + 20) = 82$ percent. New-use accuracy is $(18 + 40) / (20 + 80) = 58$ percent. The drop is 24 percentage points. Taking the unweighted mean of 90 and 50 gives 70 percent, which is not either cohort's overall accuracy because neither cohort has equal group sizes.

```
def accuracy(groups):
    if not groups:
        raise ValueError("Need at least one group")

    total = correct = 0

    for n, k in groups:
        if n <= 0 or not 0 <= k <= n:
            raise ValueError("Invalid group counts")

        total += n
        correct += k

    return correct / total

development = [(80, 72), (20, 10)]
new_use = [(20, 18), (80, 40)]

old = accuracy(development)
new = accuracy(new_use)
```

```
print(round(100 * old, 1)) # 82.0

print(round(100 * new, 1)) # 58.0

print(round(100 * (old - new), 1)) # 24.0
```

Each pair is (cases, correct), and the function assumes integer counts for non-overlapping groups. It rejects empty inputs, groups with no cases and impossible correct counts. A group with no observed cases has an undefined empirical accuracy, not zero accuracy. Do not count the same example in two groups when forming an overall total.

Equivalently, weight each group accuracy by its share of cases: development is $0.80 \times 0.90 + 0.20 \times 0.50 = 0.82$; new use is $0.20 \times 0.90 + 0.80 \times 0.50 = 0.58$.

Connect the picture and the explanation

Point to the input, the deciding step and the result in this week's visual. Explain the same step in ordinary words. A picture helps when you can say what each part represents.

4. Compare the cases

Compare a familiar case with a changed case

Calculate overall correct counts, accuracy and errors in both cohorts, plus the percentage-point drop. Explain why averaging 90 percent and 50 percent without weights gives the wrong overall result here.

Keep the procedure fixed. Explain what changed in the inputs and result.

Change one thing in the pictured case

Change: New-use mix returns to 40 bright and 10 dim with the same slice rates.

Prediction: Pooled accuracy returns to 74%.

Reason: The mixture weights change, while each slice's conditional performance stays fixed.

5. Words that unlock the idea

Word · Meaning

Word: Dataset shift

Meaning: A change in the distribution of inputs, labels or their relationship between settings

Word: Subgroup

Meaning: A declared subset of examples sharing a relevant condition

Word: Group mix

Meaning: The proportions of examples belonging to different groups

Word: Monitoring

Meaning: Repeatedly checking recorded inputs, outputs and available outcomes against a reference

Use two terms to explain a step in this week's example.

6. Follow a worked example

The same group rates meet a different mix

Group · Development share · New-use share

Group: Clear, 90% correct

Development share: 80 of 100

New-use share: 20 of 100

Group: Dim, 50% correct

Development share: 20 of 100

New-use share: 80 of 100

The weaker group becomes more common. Overall mistakes rise from 18 of 100 to 42 of 100 without any change in these observed within-group accuracy rates.

The arithmetic explains this constructed aggregate difference. It does not prove that dim lighting caused each real error or that these group rates will stay fixed in future samples. Report group counts and uncertainty about generalisation. Review specific errors and reference labels as in Week 29.

If a new cohort has 50 clear and 50 dim cases, and exactly 90 percent and 50 percent are correct respectively, how many correct cases and what overall accuracy result? State the assumption behind this calculation.

Learn from the worked case

Cover the result before checking it. Say what is given, choose the procedure, and follow one step at a time. Then uncover the result and explain your first mismatch. Ask why a step is allowed, not only what number it produces.

A pictorial worked case: Dataset shift in neural models

Separate within-setting performance from changing setting proportions.

The facts we will use

Toy old-use mix: bright 40 cases32 correct; dim 10 cases5 correct. New-use mix: bright 10 cases8 correct; dim 40 cases20 correct.

All inputs and outcomes in this worked example are invented classroom data; no real model run is claimed.

Fixed slice rates, changed pooled accuracy

accuracy percent



Bar length shows the supplied value; the marked baseline is zero.

Dataset shift in neural models

Start with the supplied facts

Follow each arrow. Ask which fact allows the next step.



1. Keep slice rates

Bright is 80% in both mixes; dim is 50% in both.



2. Old pooled rate

Old accuracy $37/50=74\%$.



3. New pooled rate

New accuracy $28/50=56\%$; more use now occurs in the harder dim setting.

CHANGE ONE THING

New-use mix returns to 40 bright and 10 dim with the same slice rates.

Explain the changed result using the matching lesson.

Read the picture one step at a time

Step · What happens and why

Step: 1. Keep slice rates

What happens and why: Bright is 80% in both mixes; dim is 50% in both.

Step: 2. Old pooled rate

What happens and why: Old accuracy $37/50=74\%$.

Step: 3. New pooled rate

What happens and why: New accuracy $28/50=56\%$; more use now occurs in the harder dim setting.

Match each pictured step to the supplied facts. This is a worked teaching case, not an independent assessment.

Why this result follows

A pooled rate depends on how often each condition appears as well as performance within that condition. Bright and dim accuracy stay fixed in this example. The new use contains more dim cases, so its pooled accuracy falls. Report both group counts and group rates before concluding that the network itself deteriorated.

What this example does not establish

Real shift can also change labels, sensors and within-slice behaviour; this toy decomposition isolates only mixture change.

Read the labelled picture: Did the supplied model's dim-light accuracy deteriorate? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.

Change one thing: New-use mix returns to 40 bright and 10 dim with the same slice rates. Predict the result before checking the explanation. What remains the same?

7. Find and repair a mistake

Start with model and input-pipeline version checks. A changed resize operation, channel order or unit conversion is an implementation change, not evidence that the input population itself shifted. Compare image conditions, missing values, source groups and label definitions. A neural network may have learned a shortcut tied to a background or capture condition; that is a hypothesis to test, not a conclusion from a histogram.

Input monitoring without new reference labels can reveal a changed group mix. It cannot calculate current accuracy or prove that the relationship between inputs and labels stayed the same. Conversely, stable input summaries can miss a changed labelling policy. Obtain a representative, correctly labelled audit sample and document unresolved label quality.

For a controlled lighting study, keep scene identity and the label fixed while varying lighting. Keep related images together in one split. If new data guide retraining or threshold choices, they are development data; reserve separate source groups or a later appropriate time period for final evaluation. More data alone do not guarantee that a shortcut is removed.

Locate the earliest unsupported step

New-use monitoring shows more dim images but reference labels have not arrived. What can you report now, what can you not calculate, and what checks should precede blaming the neural network?

State one repair and the checks you would rerun.

8. Practise together

From the table, calculate clear and dim accuracy in each cohort and the clear/dim proportions. Which observed quantities change, and which remain the same?

If a new cohort has 50 clear and 50 dim cases, and exactly 90 percent and 50 percent are correct respectively, how many correct cases and what overall accuracy result? State the assumption behind this calculation.

Explain your prediction to a partner before comparing results.

Practise with less help: Use the right denominator

Use workbook W2 for the connected example “Dataset shift in neural models”. First fill the input and rule boxes. Try the middle steps before asking for a hint. After a hint, try the next step yourself.

Plan · Do · Check

Plan: What is given? Which rule or procedure applies?

Do: Record each intermediate step.

Check: Does the result follow? What remains unknown?

Explain your trace to a partner. Your partner points to one step to check. Swap roles on the next case. Each person writes their own explanation.

Check your reasoning

Use the worked case for Dataset shift in neural models. Proposed explanation: "Average the percentages regardless of group sizes". Decide whether this explanation is supported. Point to the step or supplied fact that decides; explain your repair if needed.

Answer on your own before discussion. If you are unsure, say which step you need help with.

Explain the picture

Did the supplied model's dim-light accuracy deteriorate? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.

My explanation using a pictured step and the given facts

This is supported practice. Keep the separately named independent case for an attempt before feedback.

9. Run the investigation

Predict, investigate and preserve the evidence

Calculate overall correct counts, accuracy and errors in both cohorts, plus the percentage-point drop. Explain why averaging 90 percent and 50 percent without weights gives the wrong overall result here.

```
def accuracy(groups):  
    if not groups:  
        raise ValueError("Need at least one group")  
  
    total = correct = 0  
  
    for n, k in groups:
```

```

    if n <= 0 or not 0 <= k <= n:

        raise ValueError("Invalid group counts")

    total += n

    correct += k

    return correct / total

development = [(80, 72), (20, 10)]
new_use = [(20, 18), (80, 40)]
old = accuracy(development)
new = accuracy(new_use)

print(round(100 * old, 1)) # 82.0
print(round(100 * new, 1)) # 58.0
print(round(100 * (old - new), 1)) # 24.0

```

Record · Your evidence

Record: Given input and fixed rule

Your evidence: _____

Record: Prediction before checking

Your evidence: _____

Record: Actual result, or “not run”

Your evidence: _____

Record: First mismatch and one explanation

Your evidence: _____

Record: Boundary or missing input and expected response

Your evidence: _____

On paper, trace the supplied procedure. If code is provided, predict before running it in a local practice file. Never write an expected answer as though it were an observed run.

Plan → try → check

Before trying: what do I expect, and why? While trying: which step could first go wrong? Afterwards: what did I actually observe, and what should change? Keep “not run” if you only traced the procedure.

10. Build your understanding

Explain → apply → challenge

From the table, calculate clear and dim accuracy in each cohort and the clear/dim proportions. Which observed quantities change, and which remain the same?

New-use monitoring shows more dim images but reference labels have not arrived. What can you report now, what can you not calculate, and what checks should precede blaming the neural network?

A team retrains using new-use errors, then evaluates on those same images and their brightened copies. Explain why the score is not independent evidence and propose a comparison that includes the original model.

Explain which evidence supports your conclusion and which claim remains unproved.

11. Connect your project

Add a labelled dataset shift in neural models evidence sheet to Image-pattern Network Investigation. Show how you can explain how a changed input mix can change overall accuracy. Use the supplied fictional case and keep its inputs, intermediate steps, calculation or prediction, and limitation. Connect this record to last week's record; a matching answer without a trace is insufficient.

Project evidence · Record

Project evidence: Component and version

Record: _____

Project evidence: Inputs permitted at prediction time

Record: _____

Project evidence: Trace and comparison with earlier work

Record: _____

Project evidence: Failure, limit and human review

Record: _____

Evidence for your project

For your Image-pattern Network Investigation, save the input, procedure, observed result and limitation of this check. Label this worked example as practice; use a different, previously unreviewed case when checking independent understanding.

12. Show what you know

Possible responses include fixing a confirmed pipeline bug, gathering missing training coverage, changing a representation or routing some cases for review. Choose a response based on evidence. Compare the original and revised model on identical reserved cases, report both group and overall results, and measure review workload if a review policy is used. Preserve the original checkpoint for comparison.

A team sees more dim images but has no new reference labels. It announces accuracy has fallen, retrains, and claims success using the same cases that guided its changes. Identify what is known, what is missing and how to repair the evaluation.

Record the reference period, group definitions, model and pipeline versions, counts, available labels, decisions and follow-up outcomes. Avoid silently replacing the baseline or the evaluation set. Next week measures efficiency and resource use alongside model quality.

A team retrains using new-use errors, then evaluates on those same images and their brightened copies. Explain why the score is not independent evidence and propose a comparison that includes the original model.

My limited claim and a useful next check

13. Study a complete case

Complete evidence record

The same group rates meet a different mix

Group · Development share · New-use share

Group: Clear, 90% correct

Development share: 80 of 100

New-use share: 20 of 100

Group: Dim, 50% correct

Development share: 20 of 100

New-use share: 80 of 100

The weaker group becomes more common. Overall mistakes rise from 18 of 100 to 42 of 100 without any change in these observed within-group accuracy rates.

Follow the input through the stated procedure to its result. Cover the result first, predict it, then explain every intermediate step. The text/table accompanies the figure so the example can be read without colour.

My complete trace of the worked case

14. Try a new case independently

Independent case: Analyse a new mixture independently

A separate audit has development groups P: 60 cases, 48 correct and Q: 40 cases, 20 correct. New-use groups are P: 20 cases, 16 correct and Q: 80 cases, 40 correct. Compute each group rate, both overall accuracies, and the percentage-point change. Under the same group rates, what accuracy would equal-sized P and Q groups have?

My independent explanation, trace and limits

Attempt this case before feedback. Record any help. Earlier examples below are further practice; a case already practised with feedback cannot establish fresh transfer.

Attempt this without hints

A separate audit has development groups P: 60 cases, 48 correct and Q: 40 cases, 20 correct. New-use groups are P: 20 cases, 16 correct and Q: 80 cases, 40 correct. Compute each group rate, both overall accuracies, and the percentage-point change. Under the same group rates, what accuracy would equal-sized P and Q groups have?

My independent procedure, result and limitation

Record support used: none, hint or teacher explanation. After feedback, a different case is needed to demonstrate independence.

Use feedback to improve

Attempt the independent case before hints. Record whether you worked alone, used a hint or followed a model. After feedback, fix the first unsupported step and explain why it changed. A later check uses a changed case, not a copied answer.

15. Your lesson route

Start with this question: Separate within-setting performance from changing setting proportions.

By the end, I can · How I will show it

By the end, I can: Explain how a changed input mix can change overall accuracy

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Calculate group and overall results with the correct denominators

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Plan a shift investigation while preserving independent evaluation

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

Remember → watch and explain → try with help → investigate → try independently → keep project evidence. Your teacher may spread this lesson across several classes.

Keep your first prediction. Compare it with what the evidence shows and explain any change.

16. Finish and return

Close your notes. Explain this goal in your own words: Explain how a changed input mix can change overall accuracy

My deciding step, evidence and one remaining question

If you are unsure, return to the worked picture and explain one arrow. If you are ready, change an input or assumption and justify your prediction.

Save your workbook evidence. At the next review, try a different case before opening your old answer.