Dataset and benchmark limitations

Learn to understand, design, build, test and explain AI systems responsibly.

Project: Reproducible AI Research Study

Before you start

Recall Sensitivity analysis. Explain why a recorded result needs its input conditions and evidence source before it can support a claim.

Plan time to read, investigate and explain. Keep predictions separate from observations.

Your mission

> Start with the target: Explain the mechanism and terminology of dataset and benchmark limitations. Try the recall before opening notes.

Your learning target

I can explain the **mechanism** and terminology of dataset and **benchmark** limitations.

I can apply the stated procedure to a new case with a traceable result.

I can identify an assumption or limitation and design a revealing follow-up.

Remember before reading

Close your notes first. A rough first answer helps us choose the right starting point; it is not a score for your ability.

Without opening the old lesson, explain one key idea from Sensitivity analysis. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Multiple comparisons. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Uncertainty intervals. Give an example and a mistake that the idea helps you avoid.

After trying, check the relevant lesson or ask your teacher. Change the part of your explanation that the **evidence** shows was wrong.

This week you will practise:

Explain the **mechanism** and terminology of dataset and **benchmark** limitations

Apply the stated procedure to a new case with a traceable result

Identify an assumption or limitation and design a revealing follow-up

Retrieve what you already know

Recall Sensitivity analysis. Explain why a recorded result needs its **input** conditions and **evidence** source before it can support a claim.

> **My starting explanation**

This week's end goal

Explain dataset and **benchmark** limitations using the records and boundaries of this fictional investigation.

Notice the evidence

> Investigate the supplied Dataset and benchmark limitations case. Separate observations, given facts and assumptions before choosing an explanation.

A **benchmark** measures performance on its particular data and labels. Duplicates, obsolete cases or poor coverage can weaken claims.

Begin by writing the question the procedure is intended to answer. Identify the information already supplied, the information it will produce and the information still missing.

These are constructed classroom examples; they do not require live accounts, private records or an external service. Do not treat a plausible explanation as a measured outcome.

> **List the given facts and one fact that is not supplied.**

STUDENT LESSON · 3 OF 18

See how it works

> Follow the mechanism for Dataset and benchmark limitations. Connect each labelled input, deciding step and result.

Inspect eligibility, source dates, related units and label definitions before scoring. Deduplicate by declared unit identity, not merely file name; report excluded and unassessed cases.

Dataset and benchmark limitations — evidence trace

INPUT AND SCOPE	CHECK OR CALCULATION	CONCLUSION AND LIMIT
Ten benchmark rows include three copies of one source and seven other unique sources.	Inspect eligibility, source dates, related units and label definitions before scoring. Deduplicate by declared unit identity, not merely file name; report excluded and unassessed cases.	Eight unique source units, not ten independent sources. Identify duplicate weighting before analysis; do not silently discard results.

The figure connects this chapter's worked input, declared procedure and result. Read the complete worked reasoning for qualifiers and intermediate steps; the abbreviated labels do not replace that reasoning.

Before applying the procedure, write its units, denominator or decision boundary where relevant. A partner should be able to repeat each step using the same **input**s.

If a required **input** is absent or violates the stated contract, record that condition explicitly instead of quietly replacing it with a convenient value.

Connect the picture and the explanation

Point to the **input**, the deciding step and the result in this week's visual. Explain the same step in ordinary words. A picture helps when you can say what each part represents.

STUDENT LESSON · 4 OF 18

Compare the cases

> Compare the supplied cases for Dataset and benchmark limitations. Identify the detail that changes the conclusion.

Compare a familiar case with a changed case

Ten **benchmark** rows include three copies of one source and seven other unique sources. Show the intermediate steps and give the conclusion with its scope.

Keep the procedure fixed. Explain what changed in the inputs and result.

Change one thing in the pictured case

Change one thing

Change: Add four reserved outdoor cases, of which two are correct.

Predicted result

Prediction: Report the outdoor slice as 2/4=50%, separately from indoor results.

Why the result follows

Reason: The new records finally measure the missing condition, with a small denominator.

STUDENT LESSON · 5 OF 18

Words that unlock the idea

Use the vocabulary of Dataset and benchmark limitations to explain a concrete decision, rather than only memorising definitions.

Word	Meaning
Benchmark	A declared task dataset and scoring procedure
Duplicate unit	Repeated representation of the same underlying entity
Coverage gap	A relevant condition absent from evaluated data

Use two terms to explain a step in this week's example.

STUDENT LESSON · 6 OF 18

Follow a worked example

Follow the worked trace for Dataset and benchmark limitations, predict a step, then compare with its explanation.

Ten **benchmark** rows include three copies of one source and seven other unique sources.

Eight unique source units, not ten independent sources. Identify duplicate weighting before analysis; do not silently discard results.

Worked explanation · Dataset and benchmark limitations

Compare the intermediate **evidence** with the final statement. Explain why each number or decision follows from the stated **input**s.

The conclusion belongs to this scenario and procedure; changing the **input**s, evaluation population or authority can change what it establishes. Retain the raw record so a later revision can be compared honestly.

Learn from the worked case

Cover the result before checking it. Say what is given, choose the procedure, and follow one step at a time. Then uncover the result and explain your first mismatch. Ask why a step is allowed, not only what number it produces.

A pictorial worked case: Dataset and benchmark limitations

A **benchmark** can miss the intended setting. A plant model is evaluated on clean indoor photos but proposed for shaded outdoor photos.

The facts we will use

Supplied facts

Indoor **benchmark**: 9/10 correct. No outdoor cases are present. The same model is intended for both.

AI ACADEMY / AI-7/18

Benchmark conditions

	Correct	Total
Indoor benchmark	9	10
Added outdoor slice	2	4

Read a cell using both its row and column labels.

Follow the labelled picture with the matching explanation. Open the vector diagram

AI ACADEMY / AI-7/18

Dataset and benchmark limitations

Start with the supplied facts

Follow each arrow. Ask which fact allows the next step.



1. Name measured setting

All ten benchmark images are clean indoor photos.



2. Report the result

The observed indoor score is 90%.



3. Name missing evidence

Outdoor shade performance remains unmeasured.

CHANGE ONE THING

Add four reserved outdoor cases, of which two are correct.

Explain the changed result using the matching lesson.

Follow the labelled picture with the matching explanation. Open the vector diagram

Read the picture one step at a time

STEP	WHAT HAPPENS AND WHY
1. Name measured setting	All ten benchmark images are clean indoor photos.
2. Report the result	The observed indoor score is 90%.
3. Name missing evidence	Outdoor shade performance remains unmeasured.

Match each pictured step to the supplied facts. This is a worked teaching case, not an independent assessment.

Why this result follows

Why the result follows

A **benchmark** represents particular tasks and collection conditions. High performance on it cannot automatically establish fitness for a different deployment setting.

What this example does not establish

Keep the stated limit

Neither slice estimates all future photos; inspect sampling, labels and grouping.

Keep the stated limit

Read the labelled picture: Report the indoor **benchmark** with its denominator and missing condition. After two correct outdoor cases out of four are added, what slice rate is supported and what remains uncertain?

Keep the stated limit

Change one thing: Add four reserved outdoor cases, of which two are correct. Predict the result before checking the explanation. What remains the same?

STUDENT LESSON · 7 OF 18

Find and repair a mistake

> Find the unsupported step in the Dataset and benchmark limitations claim and explain a repair using the supplied evidence.

Benchmark success does not establish performance on a deployment population missing from its design.

Check the claim against the facts

A useful check asks what would make the reasoning fail. Distinguish a failure to execute the declared procedure from a weakness in the procedure itself.

The first may require a correction; the second requires an explicitly revised method and new **evidence**. Preserve unsuccessful cases so an improved result cannot hide the original problem.

Guided: restate the worked result, then name the assumption or limitation that prevents a broader conclusion.

Locate the earliest unsupported step

A report omits this qualification: **Benchmark** success does not establish performance on a deployment population missing from its design.

Explain how omitting it could change a reader's interpretation, and write an accurate replacement.

State one repair and the checks you would rerun.

STUDENT LESSON · 8 OF 18

Practise together

Practise Dataset and benchmark limitations with support, then explain the deciding step in your own words.

Explain the weekly **mechanism** in your own words. Use the worked scenario to name its **input**, decision rule and **output**.

Guided: restate the worked result, then name the assumption or limitation that prevents a broader conclusion.

Explain your prediction to a partner before comparing results.

Practise with less help: Trace the worked case

Use workbook W2 for the connected example “Dataset and **benchmark** limitations”. First fill the **input** and rule boxes. Try the middle steps before asking for a hint. After a hint, try the next step yourself.

Plan	Do	Check
What is given? Which rule or procedure applies?	Record each intermediate step.	Does the result follow? What remains unknown?

Explain your trace to a partner. Your partner points to one step to check. Swap roles on the next case. Each person writes their own explanation.

Check your reasoning

Use the worked case for Dataset and **benchmark** limitations. Proposed explanation: “The worked result establishes every future case”. Decide whether this explanation is supported.

Point to the step or supplied fact that decides; explain your repair if needed.

Answer on your own before discussion. If you are unsure, say which step you need help with.

Explain the picture

> Report the indoor **benchmark** with its denominator and missing condition. After two correct outdoor cases out of four are added, what slice rate is supported and what remains uncertain?

My explanation using a pictured step and the given facts

This is supported practice. Keep the separately named independent case for an attempt before feedback.

STUDENT LESSON · 9 OF 18

Run the investigation

> Investigate Dataset and benchmark limitations. Record the prediction, actual observation or labelled paper trace, explanation and next test separately.

Predict, investigate and preserve the evidence

Ten **benchmark** rows include three copies of one source and seven other unique sources. Show the intermediate steps and give the conclusion with its scope.

Record	Your evidence
Given input and fixed rule	
Prediction before checking	

Record	Your evidence
Actual result, or “not run”	
First mismatch and one explanation	
Boundary or missing input and expected response	

On paper, trace the supplied procedure. If code is provided, predict before running it in a local practice file. Never write an expected answer as though it were an observed run.

Plan → try → check

Before trying: what do I expect, and why? While trying: which step could first go wrong? Afterwards: what did I actually observe, and what should change? Keep “not run” if you only traced the procedure.

STUDENT LESSON · 10 OF 18

Build your understanding

> Move from recall to tracing, diagnosis and transfer in Dataset and benchmark limitations. Explain why each step is justified.

Explain → apply → challenge

Explain the weekly **mechanism** in your own words. Use the worked scenario to name its **input**, decision rule and **output**.

A report omits this qualification: **Benchmark** success does not establish performance on a deployment population missing from its design.

Explain how omitting it could change a reader's interpretation, and write an accurate replacement.

For the worked scenario, propose one additional case that could reveal a failure, describe its expected outcome before running it, and explain what would need to remain fixed for a fair comparison.

Do not reuse W4 as a new final test after receiving feedback.

> **Explain which evidence supports your conclusion and which claim remains unproved.**

STUDENT LESSON · 11 OF 18

Connect your project

> Apply this week's idea to Reproducible AI Research Study. Keep an annotated decision and the evidence supporting it.

Add a labelled dataset and **benchmark** limitations **evidence** sheet to Reproducible AI Research Study. Show how you can explain the **mechanism** and terminology of dataset and **benchmark** limitations.

Use the supplied fictional case and keep its **input**s, intermediate steps, calculation or **prediction**, and limitation. Connect this record to last week's record; a matching answer without a trace is insufficient.

Project evidence	Record
Component and version	
Inputs permitted at prediction time	
Trace and comparison with earlier work	
Failure, limit and human review	

Evidence for your project

For the Reproducible AI Research Study, retain an annotated worked trace and the changed-case reasoning. Connect the claim to its records, units and assumptions so another investigator can challenge or reproduce it.

STUDENT LESSON · 12 OF 18

Show what you know

> Use a genuinely new case for independent evidence. Previously discussed examples remain practice; use the separately named independent case before feedback.

Complete the supported workbook tasks before attempting W4 independently. Record any hints or help used. After feedback, use the teacher's fresh retry rather than copying the earlier answer.

For a project, select a relevant new case, predict the outcome from the declared rule and preserve the result alongside any revision.

> **Exit: propose one revealing change to the worked scenario. State the expected outcome and the evidence that would make you revise your explanation.**

For the worked scenario, propose one additional case that could reveal a failure, describe its expected outcome before running it, and explain what would need to remain fixed for a fair comparison.

Do not reuse W4 as a new final test after receiving feedback.

My limited claim and a useful next check

STUDENT LESSON · 13 OF 18

Study a complete case

Revisit the complete Dataset and benchmark limitations case for supported practice. Repeated examples do not become fresh assessments.

Complete evidence record

Dataset and benchmark limitations — evidence trace

INPUT AND SCOPE

Ten benchmark rows include three copies of one source and seven other unique sources.

CHECK OR CALCULATION

Inspect eligibility, source dates, related units and label definitions before scoring. Deduplicate by declared unit identity, not merely file name; report excluded and unassessed cases.

CONCLUSION AND LIMIT

Eight unique source units, not ten independent sources. Identify duplicate weighting before analysis; do not silently discard results.

The figure connects this chapter's worked input, declared procedure and result. Read the complete worked reasoning for qualifiers and intermediate steps; the abbreviated labels do not replace that reasoning.

Follow the **input** through the stated procedure to its result. Cover the result first, predict it, then explain every intermediate step. The text/table accompanies the figure so the example can be read without colour.

My complete trace of the worked case

STUDENT LESSON · 14 OF 18

Try a new case independently

> Attempt the reserved new case before feedback. Record support; the earlier examples below remain additional practice.

Independent case: Apply a new case independently

Twelve rows include four copies of one source and eight others. Count unique units and name one coverage concern.

My independent explanation, trace and limits

Attempt this case before feedback. Record any help. Earlier examples below are further practice; a case already practised with feedback cannot establish fresh transfer.

Attempt this without hints

Twelve rows include four copies of one source and eight others. Count unique units and name one coverage concern.

My independent procedure, result and limitation

Record support used: none, hint or teacher explanation. After feedback, a different case is needed to demonstrate independence.

Use feedback to improve

Attempt the independent case before hints. Record whether you worked alone, used a hint or followed a model. After feedback, fix the first unsupported step and explain why it changed.

A later check uses a changed case, not a copied answer.

STUDENT LESSON · 15 OF 18

Your lesson route

> Use this map to navigate Dataset and benchmark limitations, connect the goals to tasks and keep one evidence record for each attempt.

Start with this question: A **benchmark** can miss the intended setting. A plant model is evaluated on clean indoor photos but proposed for shaded outdoor photos.

By the end, I can	How I will show it
Explain the mechanism and terminology of dataset and benchmark limitations	Explain the deciding step, show a trace or test, and state one limit.
Apply the stated procedure to a new case with a traceable result	Explain the deciding step, show a trace or test, and state one limit.
Identify an assumption or limitation and design a revealing follow-up	Explain the deciding step, show a trace or test, and state one limit.

Remember → watch and explain → try with help → investigate → try independently → keep project **evidence**. Your teacher may spread this lesson across several classes.

Keep your first **prediction**. Compare it with what the **evidence** shows and explain any change.

STUDENT LESSON · 16 OF 18

Finish and return

> Explain Dataset and benchmark limitations without notes, identify a remaining question and plan a later check with a changed case.

Close your notes. Explain this goal in your own words: Explain the **mechanism** and terminology of dataset and **benchmark** limitations

> **My deciding step, evidence and one remaining question**

If you are unsure, return to the worked picture and explain one arrow. If you are ready, change an **input** or assumption and justify your **prediction**.

Save your workbook **evidence**. At the next review, try a different case before opening your old answer.

STUDENT LESSON · 17 OF 18

Supported practice, repair and portfolio

> Attempt supported questions before using hints or feedback. Review a missing building

> block, then keep the completed work.

Use this supported practice once, in the browser or on paper. Keep the reserved independent assessment for a separate first attempt before teacher feedback.

A **benchmark** can miss the intended setting. A plant model is evaluated on clean indoor photos but proposed for shaded outdoor photos.

Practice 1: explain the deciding evidence

Which labelled step matches this detail? Outdoor shade performance remains unmeasured.

1. Report the result

2. Name measured setting

3. Name missing **evidence**

My selected conclusion, deciding fact and explanation

Hint · try first

Hint if needed: Read the steps in the pictorial case. Match the supplied detail to the stage that performs or records it.

Feedback · compare after trying

After your attempt, compare with the supported explanation: Name missing **evidence** Outdoor shade performance remains unmeasured. A **benchmark** represents particular tasks and collection conditions.

High performance on it cannot automatically establish fitness for a different deployment setting.

Practice 2: explain the deciding evidence

Changed case · predict first

Changed-case practice: Add four reserved outdoor cases, of which two are correct. Which conclusion is supported by the supplied facts?

1. Report the outdoor slice as 2/4=50%, separately from indoor results.

2. We can decide the result without examining the changed **input** or the procedure.

3. This one example guarantees correct results for every future **input**.

My selected conclusion, deciding fact and explanation

Hint · try first

Hint if needed: Use the changed condition and the deciding step in the pictorial trace. The answer must say what follows from those facts.

Feedback · compare after trying

After your attempt, compare with the supported explanation: Report the outdoor slice as 2/4=50%, separately from indoor results. The new records finally measure the missing condition, with a small denominator.

Repair a missing building block

Review Sensitivity analysis (ai-7/17). Goal: Explain the **mechanism** and terminology of sensitivity analysis. Short practice: Classify 0.4, 0.6 and 0.8 at cutoffs 0.5, 0.7 and exactly 0.6.

Which equality rule decides the boundary, and can these unlabelled scores identify the best cutoff?. Return to this week and explain what you changed.

Review Multiple comparisons (ai-7/15). Goal: Explain the **mechanism** and terminology of multiple comparisons. Short practice: Calculate the six-test cutoff and identify which supplied p-values reject.

Why must the team count every planned comparison rather than report only its smallest p-value?. Return to this week and explain what you changed.

Review Uncertainty intervals (ai-7/12). Goal: Explain the **mechanism** and terminology of uncertainty intervals. Short practice: Locate -0.8 and zero relative to the supplied interval [-1.4,-0.2].

Interpret the sign and confidence-method meaning, then describe what changes with interval [-1.4,0.3].. Return to this week and explain what you changed.

Keep your completed work

Use Download completed work in the hosted lesson to export your answers, reflections, practice feedback and code-run **evidence**. Open the HTML copy to read or print it.

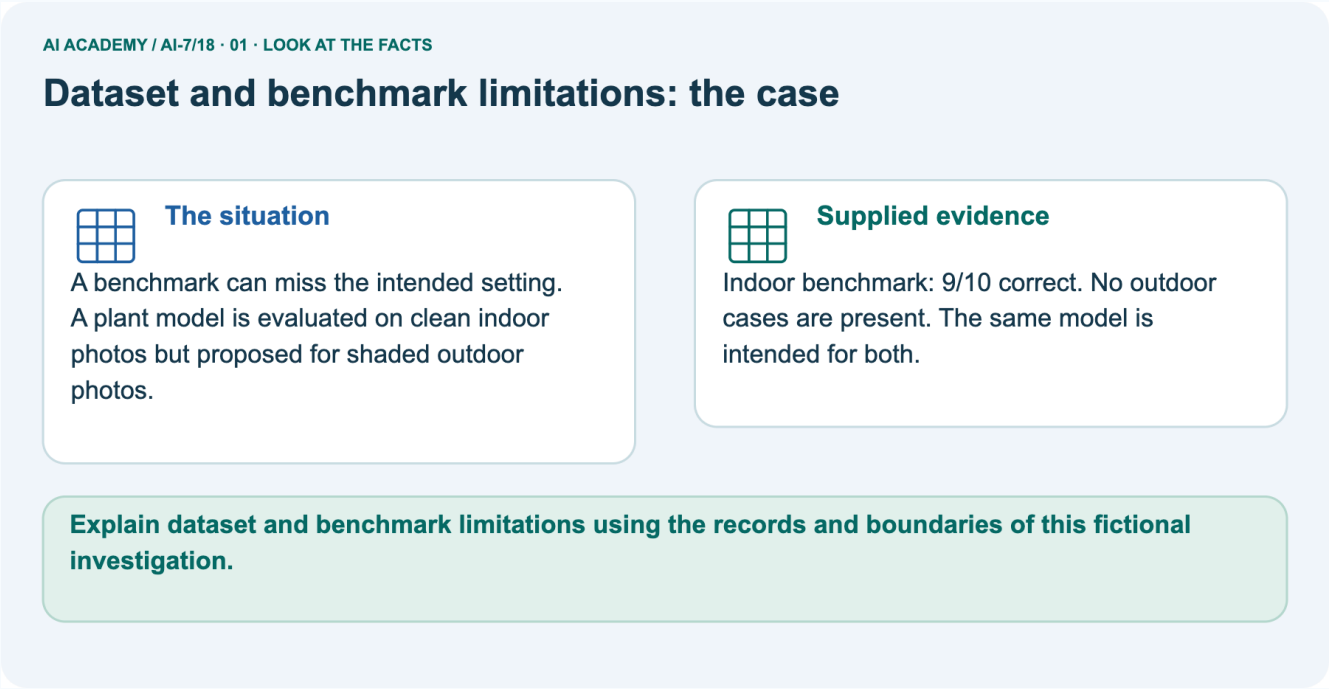
Drafts stay in this browser when storage is available; clear drafts on a shared device.

STUDENT LESSON · 18 OF 18

Visual concept journey

> Explain each Dataset and benchmark limitations graphic using its labels, deciding evidence and relevant limitation.

01. Dataset and benchmark limitations: the case



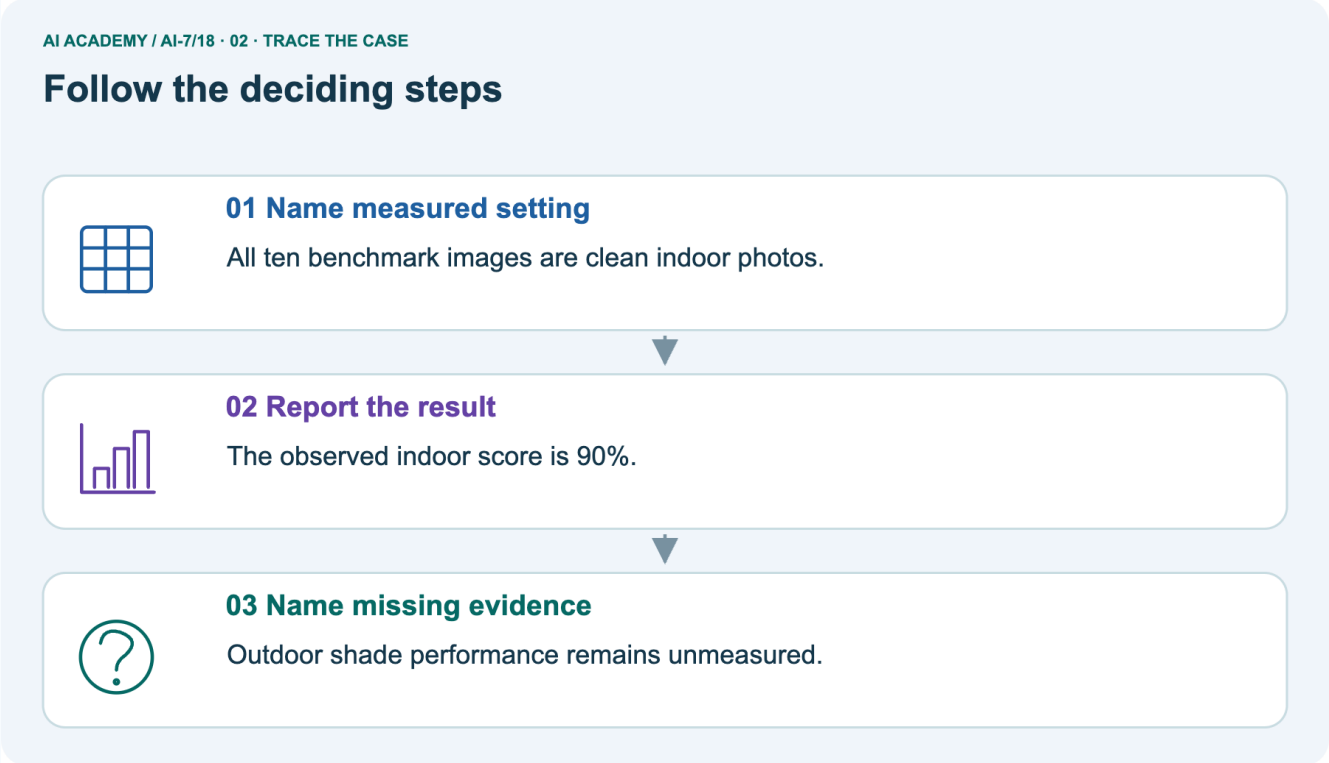
Follow the labelled picture with the matching explanation. Open the vector diagram

A **benchmark** can miss the intended setting. A plant model is evaluated on clean indoor photos but proposed for shaded outdoor photos. Supplied facts: Indoor **benchmark**: 9/10 correct. No outdoor cases are present.

The same model is intended for both.

Point to the given fact that matters most. Explain why it matters.

02. Follow the deciding steps

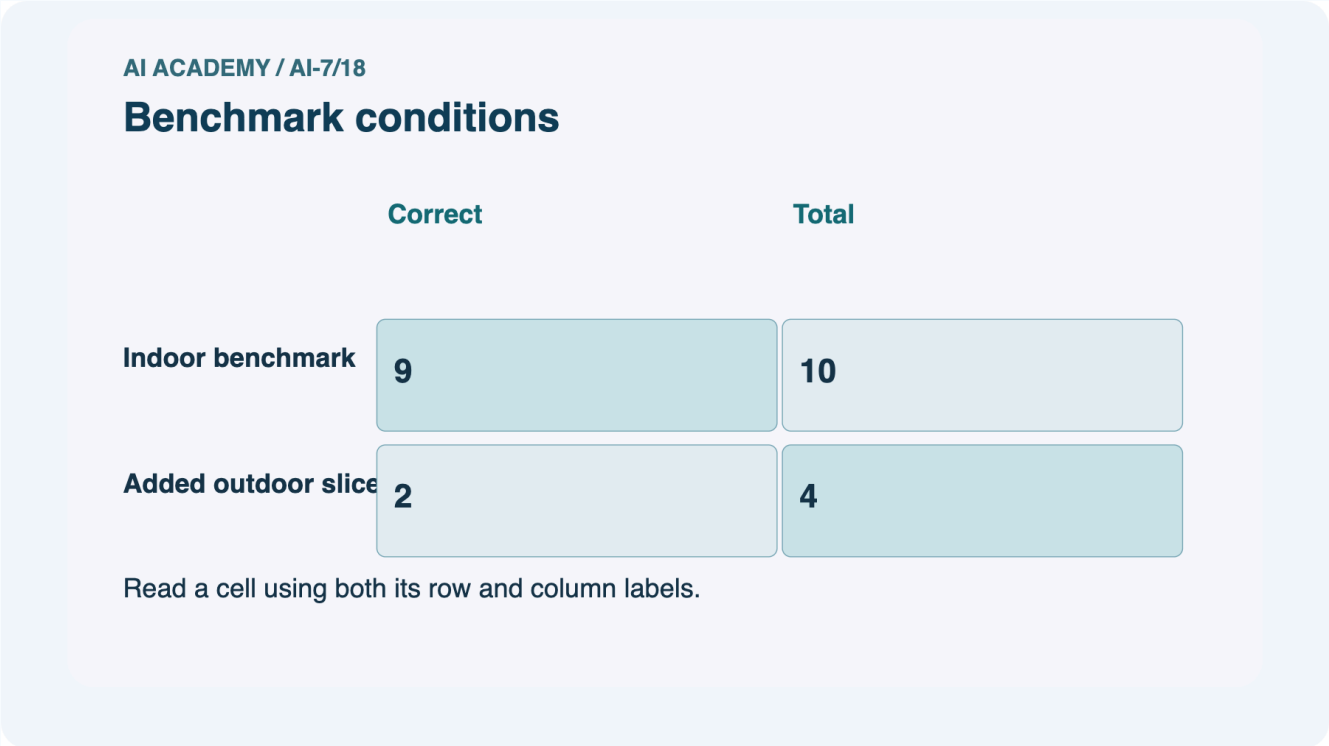


Follow the labelled picture with the matching explanation. Open the vector diagram

Name measured setting: All ten **benchmark** images are clean indoor photos. → Report the result: The observed indoor score is 90%. → Name missing **evidence**: Outdoor shade performance remains unmeasured.

Point to each numbered stage and explain which recorded **input** it uses.

03. Benchmark conditions

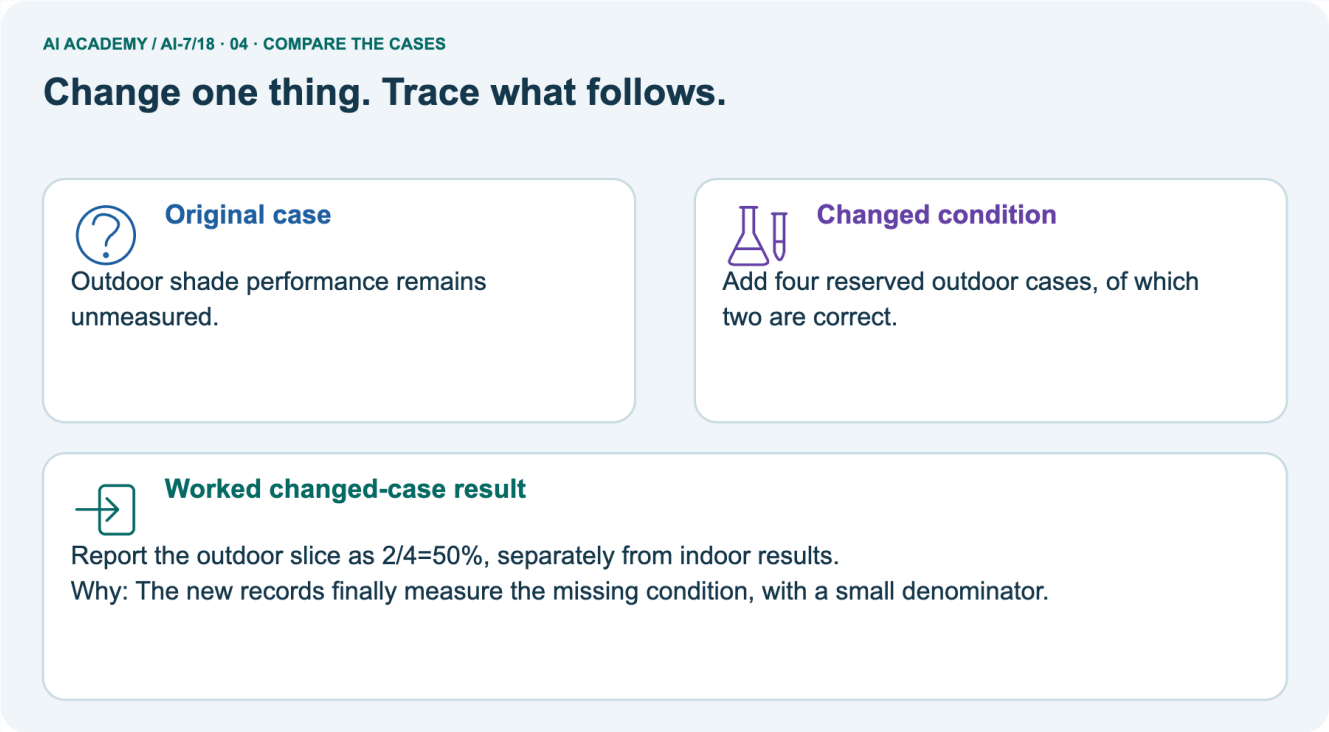


Follow the labelled picture with the matching explanation. Open the vector diagram

Benchmark conditions. Indoor **benchmark**: Correct = 9, Total = 10. Added outdoor slice: Correct = 2, Total = 4.

> Read the labels and values. Explain the relationship the model shows.

04. Change one thing. Trace what follows.



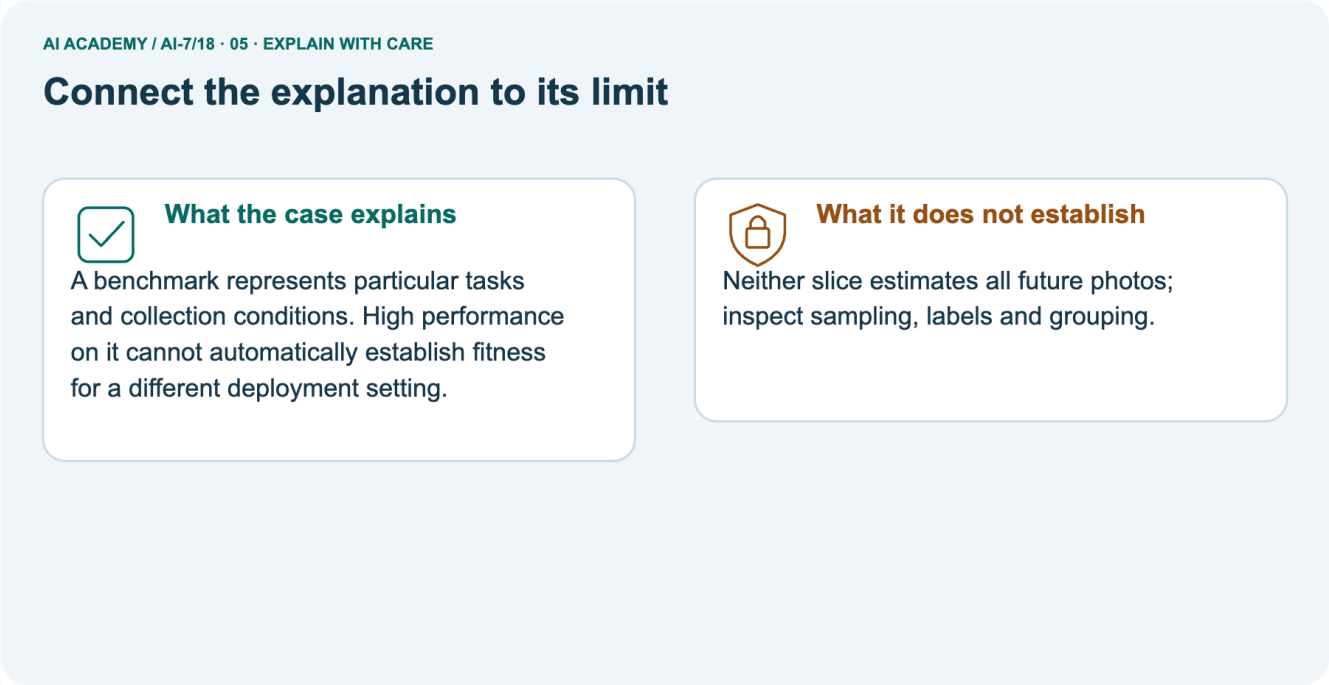
Follow the labelled picture with the matching explanation. Open the vector diagram

Original case: Outdoor shade performance remains unmeasured. Changed condition: Add four reserved outdoor cases, of which two are correct. Expected result: Report the outdoor slice as 2/4=50%, separately from indoor results.

Reason: The new records finally measure the missing condition, with a small denominator.

> Predict before checking the worked changed-case result. Name what changed and what the supplied facts allow you to conclude.

05. Connect the explanation to its limit



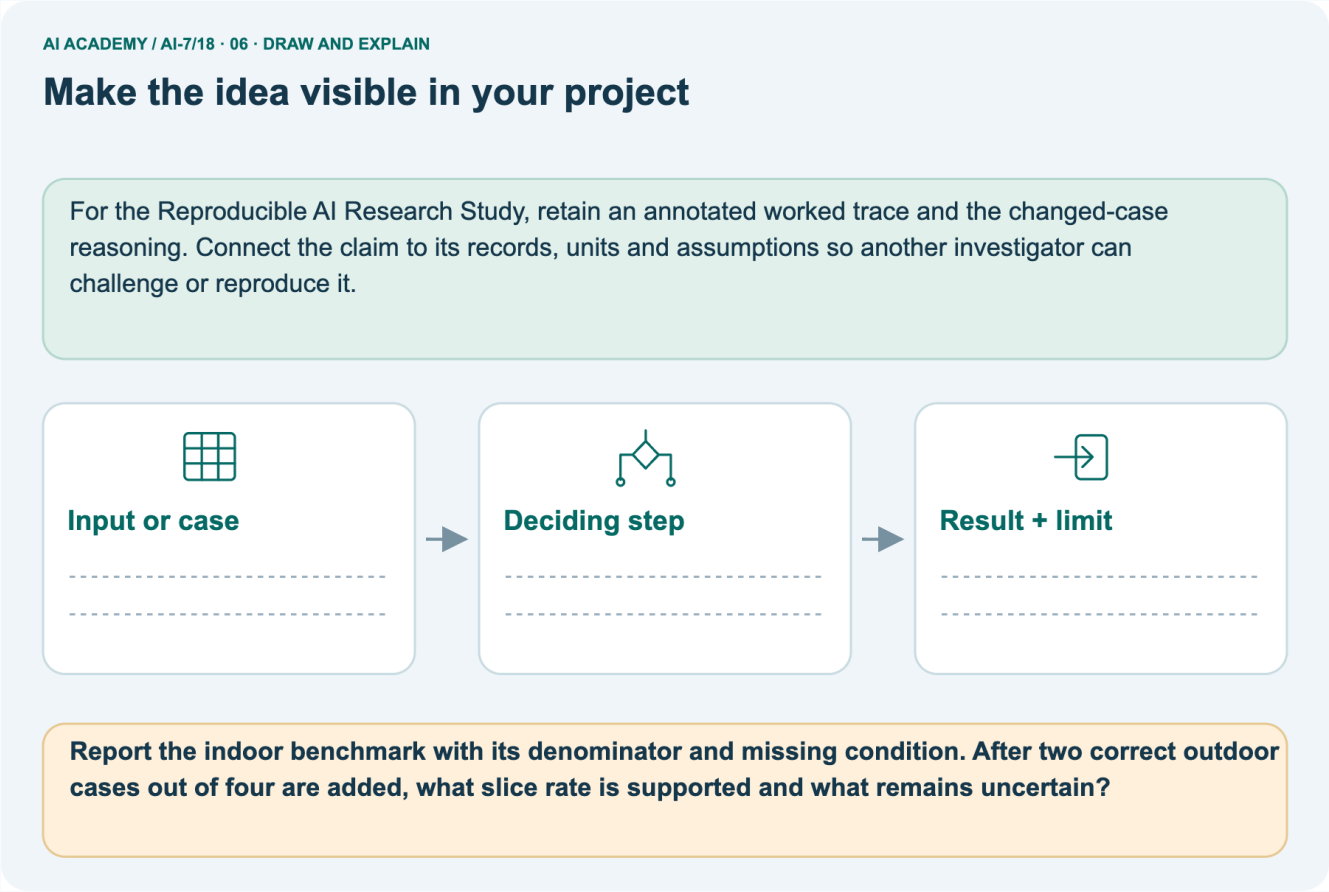
Follow the labelled picture with the matching explanation. Open the vector diagram

A **benchmark** represents particular tasks and collection conditions. High performance on it cannot automatically establish fitness for a different deployment setting.

Limit: Neither slice estimates all future photos; inspect sampling, labels and grouping.

> Explain the case, then point to a limit on the claim.

06. Make the idea visible in your project



Follow the labelled picture with the matching explanation. Open the vector diagram

For the Reproducible AI Research Study, retain an annotated worked trace and the changed-case reasoning. Connect the claim to its records, units and assumptions so another investigator can challenge or reproduce it.

Draw your **input**, deciding step, result and limit. Supported question: Report the indoor **benchmark** with its denominator and missing condition.

After two correct outdoor cases out of four are added, what slice rate is supported and what remains uncertain?

> Draw your project trace and justify the deciding step using **evidence**.

Save or print your responses for your portfolio. Progress and drafts on the hosted website stay in the current browser. Use only fictional or teacher-approved data.