

Dataset and benchmark limitations

Project: Reproducible AI Research Study

1. Your mission

Your learning target

I can explain the mechanism and terminology of dataset and benchmark limitations.

I can apply the stated procedure to a new case with a traceable result.

I can identify an assumption or limitation and design a revealing follow-up.

Remember before reading

Close your notes first. A rough first answer helps us choose the right starting point; it is not a score for your ability.

Without opening the old lesson, explain one key idea from Sensitivity analysis. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Multiple comparisons. Give an example and a mistake that the idea helps you avoid.

Without opening the old lesson, explain one key idea from Uncertainty intervals. Give an example and a mistake that the idea helps you avoid.

After trying, check the relevant lesson or ask your teacher. Change the part of your explanation that the evidence shows was wrong.

This week you will practise:

Explain the mechanism and terminology of dataset and benchmark limitations

Apply the stated procedure to a new case with a traceable result

Identify an assumption or limitation and design a revealing follow-up

Retrieve what you already know

Recall Sensitivity analysis. Explain why a recorded result needs its input conditions and evidence source before it can support a claim.

My starting explanation

This week's end goal

Explain dataset and benchmark limitations using the records and boundaries of this fictional investigation.

2. Notice the evidence

A benchmark measures performance on its particular data and labels. Duplicates, obsolete cases or poor coverage can weaken claims.

Begin by writing the question the procedure is intended to answer. Identify the information already supplied, the information it will produce and the information still missing. These are constructed classroom examples; they do not require live accounts, private records or an external service. Do not treat a plausible explanation as a measured outcome.

List the given facts and one fact that is not supplied.

3. See how it works

Inspect eligibility, source dates, related units and label definitions before scoring. Deduplicate by declared unit identity, not merely file name; report excluded and unassessed cases.

Dataset and benchmark limitations — evidence trace

Input and scope · Check or calculation · Conclusion and limit

Input and scope: Ten benchmark rows include three copies of one source and seven other unique sources.

Check or calculation: Inspect eligibility, source dates, related units and label definitions before scoring. Deduplicate by declared unit identity, not merely file name; report excluded and unassessed cases.

Conclusion and limit: Eight unique source units, not ten independent sources. Identify duplicate weighting before analysis; do not silently discard results.

The figure connects this chapter's worked input, declared procedure and result. Read the complete worked reasoning for qualifiers and intermediate steps; the abbreviated labels do not replace that reasoning.

Before applying the procedure, write its units, denominator or decision boundary where relevant. A partner should be able to repeat each step using the same inputs. If a required input is absent or violates the stated contract, record that condition explicitly instead of quietly replacing it with a convenient value.

Connect the picture and the explanation

Point to the input, the deciding step and the result in this week's visual. Explain the same step in ordinary words. A picture helps when you can say what each part represents.

4. Compare the cases

Compare a familiar case with a changed case

Ten benchmark rows include three copies of one source and seven other unique sources. Show the intermediate steps and give the conclusion with its scope.

Keep the procedure fixed. Explain what changed in the inputs and result.

Change one thing in the pictured case

Change: Add four reserved outdoor cases, of which two are correct.

Prediction: Report the outdoor slice as $2/4=50\%$, separately from indoor results.

Reason: The new records finally measure the missing condition, with a small denominator.

5. Words that unlock the idea

Word · Meaning

Word: Benchmark

Meaning: A declared task dataset and scoring procedure

Word: Duplicate unit

Meaning: Repeated representation of the same underlying entity

Word: Coverage gap

Meaning: A relevant condition absent from evaluated data

Use two terms to explain a step in this week's example.

6. Follow a worked example

Ten benchmark rows include three copies of one source and seven other unique sources.

Eight unique source units, not ten independent sources. Identify duplicate weighting before analysis; do not silently discard results.

Compare the intermediate evidence with the final statement. Explain why each number or decision follows from the stated inputs. The conclusion belongs to this scenario and procedure; changing the inputs, evaluation population or authority can change what it establishes. Retain the raw record so a later revision can be compared honestly.

Learn from the worked case

Cover the result before checking it. Say what is given, choose the procedure, and follow one step at a time. Then uncover the result and explain your first mismatch. Ask why a step is allowed, not only what number it produces.

A pictorial worked case: Dataset and benchmark limitations

A benchmark can miss the intended setting. A plant model is evaluated on clean indoor photos but proposed for shaded outdoor photos.

The facts we will use

Indoor benchmark: 9/10 correct. No outdoor cases are present. The same model is intended for both.

Benchmark conditions

	Correct	Total
Indoor benchmark	9	10
Added outdoor slice	2	4

Read a cell using both its row and column labels.

Dataset and benchmark limitations

Start with the supplied facts

Follow each arrow. Ask which fact allows the next step.



1. Name measured setting

All ten benchmark images are clean indoor photos.



2. Report the result

The observed indoor score is 90%.



3. Name missing evidence

Outdoor shade performance remains unmeasured.

CHANGE ONE THING

Add four reserved outdoor cases, of which two are correct.

Explain the changed result using the matching lesson.

Read the picture one step at a time

Step · What happens and why

Step: 1. Name measured setting

What happens and why: All ten benchmark images are clean indoor photos.

Step: 2. Report the result

What happens and why: The observed indoor score is 90%.

Step: 3. Name missing evidence

What happens and why: Outdoor shade performance remains unmeasured.

Match each pictured step to the supplied facts. This is a worked teaching case, not an independent assessment.

Why this result follows

A benchmark represents particular tasks and collection conditions. High performance on it cannot automatically establish fitness for a different deployment setting.

What this example does not establish

Neither slice estimates all future photos; inspect sampling, labels and grouping.

Read the labelled picture: Report the indoor benchmark with its denominator and missing condition. After two correct outdoor cases out of four are added, what slice rate is supported and what remains uncertain?

Change one thing: Add four reserved outdoor cases, of which two are correct. Predict the result before checking the explanation. What remains the same?

7. Find and repair a mistake

Benchmark success does not establish performance on a deployment population missing from its design.

A useful check asks what would make the reasoning fail. Distinguish a failure to execute the declared procedure from a weakness in the procedure itself. The first may require a correction; the second requires an explicitly revised method and new evidence. Preserve unsuccessful cases so an improved result cannot hide the original problem.

Guided: restate the worked result, then name the assumption or limitation that prevents a broader conclusion.

Locate the earliest unsupported step

A report omits this qualification: Benchmark success does not establish performance on a deployment population missing from its design. Explain how omitting it could change a reader's interpretation, and write an accurate replacement.

State one repair and the checks you would rerun.

8. Practise together

Explain the weekly mechanism in your own words. Use the worked scenario to name its input, decision rule and output.

Guided: restate the worked result, then name the assumption or limitation that prevents a broader conclusion.

Explain your prediction to a partner before comparing results.

Practise with less help: Trace the worked case

Use workbook W2 for the connected example “Dataset and benchmark limitations”. First fill the input and rule boxes. Try the middle steps before asking for a hint. After a hint, try the next step yourself.

Plan · Do · Check

Plan: What is given? Which rule or procedure applies?

Do: Record each intermediate step.

Check: Does the result follow? What remains unknown?

Explain your trace to a partner. Your partner points to one step to check. Swap roles on the next case. Each person writes their own explanation.

Check your reasoning

Use the worked case for Dataset and benchmark limitations. Proposed explanation: “The worked result establishes every future case”. Decide whether this explanation is supported. Point to the step or supplied fact that decides; explain your repair if needed.

Answer on your own before discussion. If you are unsure, say which step you need help with.

Explain the picture

Report the indoor benchmark with its denominator and missing condition. After two correct outdoor cases out of four are added, what slice rate is supported and what remains uncertain?

My explanation using a pictured step and the given facts

This is supported practice. Keep the separately named independent case for an attempt before feedback.

9. Run the investigation

Predict, investigate and preserve the evidence

Ten benchmark rows include three copies of one source and seven other unique sources. Show the intermediate steps and give the conclusion with its scope.

Record · Your evidence

Record: Given input and fixed rule

Your evidence: _____

Record: Prediction before checking

Your evidence: _____

Record: Actual result, or “not run”

Your evidence: _____

Record: First mismatch and one explanation

Your evidence: _____

Record: Boundary or missing input and expected response

Your evidence: _____

On paper, trace the supplied procedure. If code is provided, predict before running it in a local practice file. Never write an expected answer as though it were an observed run.

Plan → try → check

Before trying: what do I expect, and why? While trying: which step could first go wrong? Afterwards: what did I actually observe, and what should change? Keep “not run” if you only traced the procedure.

10. Build your understanding

Explain → apply → challenge

Explain the weekly mechanism in your own words. Use the worked scenario to name its input, decision rule and output.

A report omits this qualification: Benchmark success does not establish performance on a deployment population missing from its design. Explain how omitting it could change a reader's interpretation, and write an accurate replacement.

For the worked scenario, propose one additional case that could reveal a failure, describe its expected outcome before running it, and explain what would need to remain fixed for a fair comparison. Do not reuse W4 as a new final test after receiving feedback.

Explain which evidence supports your conclusion and which claim remains unproved.

11. Connect your project

Add a labelled dataset and benchmark limitations evidence sheet to Reproducible AI Research Study. Show how you can explain the mechanism and terminology of dataset and benchmark limitations. Use the supplied fictional case and keep its inputs, intermediate steps, calculation or prediction, and limitation. Connect this record to last week's record; a matching answer without a trace is insufficient.

Project evidence · Record

Project evidence: Component and version

Record: _____

Project evidence: Inputs permitted at prediction time

Record: _____

Project evidence: Trace and comparison with earlier work

Record: _____

Project evidence: Failure, limit and human review

Record: _____

Evidence for your project

For the Reproducible AI Research Study, retain an annotated worked trace and the changed-case reasoning. Connect the claim to its records, units and assumptions so another investigator can challenge or reproduce it.

12. Show what you know

Complete the supported workbook tasks before attempting W4 independently. Record any hints or help used. After feedback, use the teacher's fresh retry rather than copying the earlier answer. For a project, select a relevant new case, predict the outcome from the declared rule and preserve the result alongside

any revision.

Exit: propose one revealing change to the worked scenario. State the expected outcome and the evidence that would make you revise your explanation.

For the worked scenario, propose one additional case that could reveal a failure, describe its expected outcome before running it, and explain what would need to remain fixed for a fair comparison. Do not reuse W4 as a new final test after receiving feedback.

My limited claim and a useful next check

13. Study a complete case

Complete evidence record

Dataset and benchmark limitations — evidence trace

Input and scope · Check or calculation · Conclusion and limit

Input and scope: Ten benchmark rows include three copies of one source and seven other unique sources.

Check or calculation: Inspect eligibility, source dates, related units and label definitions before scoring. Deduplicate by declared unit identity, not merely file name; report excluded and unassessed cases.

Conclusion and limit: Eight unique source units, not ten independent sources. Identify duplicate weighting before analysis; do not silently discard results.

The figure connects this chapter's worked input, declared procedure and result. Read the complete worked reasoning for

qualifiers and intermediate steps; the abbreviated labels do not replace that reasoning.

Follow the input through the stated procedure to its result. Cover the result first, predict it, then explain every intermediate step. The text/table accompanies the figure so the example can be read without colour.

My complete trace of the worked case

14. Try a new case independently

Independent case: Apply a new case independently

Twelve rows include four copies of one source and eight others. Count unique units and name one coverage concern.

My independent explanation, trace and limits

Attempt this case before feedback. Record any help. Earlier examples below are further practice; a case already practised with feedback cannot establish fresh transfer.

Attempt this without hints

Twelve rows include four copies of one source and eight others. Count unique units and name one coverage concern.

My independent procedure, result and limitation

Record support used: none, hint or teacher explanation. After feedback, a different case is needed to demonstrate independence.

Use feedback to improve

Attempt the independent case before hints. Record whether you worked alone, used a hint or followed a model. After feedback, fix the first unsupported step and explain why it changed. A later check uses a changed case, not a copied answer.

15. Your lesson route

Start with this question: A benchmark can miss the intended setting. A plant model is evaluated on clean indoor photos but proposed for shaded outdoor photos.

By the end, I can · How I will show it

By the end, I can: Explain the mechanism and terminology of dataset and benchmark limitations

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Apply the stated procedure to a new case with a traceable result

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

By the end, I can: Identify an assumption or limitation and design a revealing follow-up

How I will show it: Explain the deciding step, show a trace or test, and state one limit.

Remember → watch and explain → try with help → investigate → try independently → keep project evidence. Your teacher may spread this lesson across several classes.

Keep your first prediction. Compare it with what the evidence shows and explain any change.

16. Finish and return

Close your notes. Explain this goal in your own words: Explain the mechanism and terminology of dataset and benchmark limitations

My deciding step, evidence and one remaining question

If you are unsure, return to the worked picture and explain one arrow. If you are ready, change an input or assumption and justify your prediction.

Save your workbook evidence. At the next review, try a different case before opening your old answer.

17. Supported practice, repair and portfolio

Use this supported practice once, in the browser or on paper. Keep the reserved independent assessment for a separate first attempt before teacher feedback.

A benchmark can miss the intended setting. A plant model is evaluated on clean indoor photos but proposed for shaded outdoor photos.

Practice 1: explain the deciding evidence

Which labelled step matches this detail? Outdoor shade performance remains unmeasured.

1. Report the result
2. Name measured setting
3. Name missing evidence

My selected conclusion, deciding fact and explanation

Hint if needed: Read the steps in the pictorial case. Match the supplied detail to the stage that performs or records it.

After your attempt, compare with the supported explanation: Name missing evidence Outdoor shade performance remains unmeasured. A benchmark represents particular tasks and collection conditions. High performance on it cannot automatically establish fitness for a different deployment setting.

Practice 2: explain the deciding evidence

Changed-case practice: Add four reserved outdoor cases, of which two are correct. Which conclusion is supported by the supplied facts?

1. Report the outdoor slice as $2/4=50\%$, separately from indoor results.
2. We can decide the result without examining the changed input or the procedure.
3. This one example guarantees correct results for every future input.

My selected conclusion, deciding fact and explanation

Hint if needed: Use the changed condition and the deciding step in the pictorial trace. The answer must say what follows from those facts.

After your attempt, compare with the supported explanation: Report the outdoor slice as $2/4=50\%$, separately from indoor results. The new records finally measure the missing condition, with a small denominator.

Repair a missing building block

Review Sensitivity analysis (ai-7/17). Goal: Explain the mechanism and terminology of sensitivity analysis. Short practice: Classify 0.4, 0.6 and 0.8 at cutoffs 0.5, 0.7 and exactly 0.6. Which equality rule decides the boundary, and can these unlabelled scores identify the best cutoff?. Return to this week and explain what you changed.

Review Multiple comparisons (ai-7/15). Goal: Explain the mechanism and terminology of multiple comparisons. Short practice: Calculate the six-test cutoff and identify which supplied p-values reject. Why must the team count every planned comparison rather than report only its smallest p-value?. Return to this week and explain what you changed.

Review Uncertainty intervals (ai-7/12). Goal: Explain the mechanism and terminology of uncertainty intervals. Short practice: Locate -0.8 and zero relative to the supplied interval $[-1.4, -0.2]$. Interpret the sign and confidence-method meaning, then describe what changes with interval $[-1.4, 0.3]$.. Return to this week and explain what you changed.

Keep your completed work

Use Download completed work in the hosted lesson to export your answers, reflections, practice feedback and code-run evidence. Open the HTML copy to read or print it. Drafts stay in this browser when storage is available; clear drafts on a shared device.