

Evaluation Sets

Project: Verified Study Buddy

1. Evaluation Sets

Teach the learning cycle

Primary target: Explain the mechanism and terminology of system evaluation sets

Session · Sequence and minutes · Resources

Session: Session 1 · Understand · 45 min

Sequence and minutes: 0–7: recall and prerequisite; 7–15: inspect the concrete case and predict; 15–30: teacher models the worked trace; 30–40: learners explain a step using the text and visual; 40–45: check the core idea.

Resources: Sections 1–6; record the brief recall in the workbook.

Session: Session 2 · Practise · 45 min

Sequence and minutes: 0–5: retrieve the procedure; 5–15: W2 completion practice; 15–23: explain and check with a partner; 23–30: each learner answers the misconception check; 30–40: targeted reteaching or a harder variation; 40–45: prepare the investigation.

Resources: Sections 7–10; use the matching workbook W2 and misconception check.

Session: Session 3 · Investigate and transfer · 45 min

Sequence and minutes: 0–5: retrieve the decision rule; 5–20: investigate or trace; 20–30: start the independent case; 30–38: connect the project; 38–43: act on feedback; 43–45: plan the return check.

Resources: Sections 9 and 11–14. Give extra time to finish the full task set; use Section 15 when the week includes a coding lab.

This is a starting allocation of 135 minutes, not a validated duration. Preserve all named topics. Schedule extra time for connected examples, the full investigation, independent cases and project building. Repeated copies of a case share one recorded attempt; they are not new evidence.

Case exposure check: this source example appears in ai-3/25, ai-3/26. Reuse a familiar case for recall or

modelling; verify that the case used to assess transfer has not already received feedback.

Retrieval answer guidance

Quality Rubrics

Without opening the old lesson, explain one key idea from Quality Rubrics. Give an example and a mistake that the idea helps you avoid.

Look for: Explain the mechanism and terminology of system evaluation sets

Reference example, not the only acceptable wording: Teacher explanation: A system evaluation set contains cases with expected behaviour and scoring rules. It should cover retrieval, support, fallback and permissions rather than only easy answerable questions. Record case ID, evidence version, expected claims or fallback, permitted actions and scoring rubric. Separate development from final cases. Aggregate metrics only after checking required conditions individually.

Block to Python Bridge

Without opening the old lesson, explain one key idea from Block to Python Bridge. Give an example and a mistake that the idea helps you avoid.

Look for: Translate a short paper calculation into a Python script

Reference example, not the only acceptable wording: `packs = 4; per_pack = 6; total = packs * per_pack; print(total)`, each on its own line. The first two bind quantities, the third computes 24 and the fourth displays 24. Assignment alone produces no displayed output in this script.

Conversation State and Memory

Without opening the old lesson, explain one key idea from Conversation State and Memory. Give an example and a mistake that the idea helps you avoid.

Look for: Explain the mechanism and terminology of conversation state and retention

Reference example, not the only acceptable wording: Teacher explanation: Conversation state stores facts or prior interactions for later use. It needs scope, provenance, correction and retention rules rather than indefinite accumulation. This classroom assistant keeps an approved room ID only for the current session. A later correction replaces the active value but retains an audit event until session end. Do not persist private text merely because it appeared in conversation.

Use the student module and workbook section with this exact week title. The worked case, independent question and answers below form one aligned lesson sequence.

What the learner must demonstrate

A test set contains ten easy questions with direct answers in the notes. Add three different kinds of cases that probe fallback behaviour and define expected outcomes.

Concept to explain

A good eval set represents real use, edge cases, and known risks.

Term · Meaning

Term: evaluation set

Meaning: Cases used to measure a system under declared criteria; specify whether they are validation or final test cases.

Term: coverage

Meaning: The range of relevant cases, conditions or groups included in data or tests.

Term: expected behavior

Meaning: The result or action required by the specification for a given case, stated before observing the system.

Term: regression test

Meaning: A repeatable check that previously correct behaviour still works after a change.

Prepare the class

Use the supplied scenario, student illustrations and blank paper. For numeric work, offer counters or a calculator after learners show the setup. Fictional records are supplied; personal data and paid accounts are unnecessary.

Allow 80 minutes: recall 8; explanation 12; worked demonstration 15; guided practice 20; independent assessment 15; project and reflection 10. Extend the lab into a second session when needed.

Preparation for the connected explanation

Two 45-minute sessions. A system evaluation set contains cases with expected behaviour and scoring rules. It should cover retrieval, support, fallback and permissions rather than only easy answerable questions. Use the supplied constructed scenarios to make the mechanism inspectable. Collect W4 before sharing solutions. Students should name intermediate evidence and limitations; classroom pacing and learning effectiveness remain unvalidated.

Readiness answer

A result depends on the inputs, procedure and conditions that produced it. Preserve data roles and distinguish observed outcomes from predictions or unmeasured conditions. If the prior concept cannot be explained, use a short retrieval activity before this chapter.

Lesson and workbook; paper and pencil

Calculator and optional local Python 3 for arithmetic; no live API or private data

Prepare the fresh retry separately and withhold its solution until after collection

Apply the connected idea: Ranking and retrieval evaluation

Two 45-minute sessions. A ranking needs task-specific relevance judgments and a declared cutoff. Precision and recall have different denominators. Use the supplied constructed scenarios to make the mechanism inspectable. Collect W4 before sharing solutions. Students should name intermediate evidence and limitations; classroom pacing and learning effectiveness remain unvalidated.

Readiness answer

A result depends on the inputs, procedure and conditions that produced it. Preserve data roles and distinguish observed outcomes from predictions or unmeasured conditions. If the prior concept cannot be explained, use a short retrieval activity before this chapter.

Lesson and workbook; paper and pencil

Calculator and optional local Python 3 for arithmetic; no live API or private data

Prepare the fresh retry separately and withhold its solution until after collection

End goal for the pictorial case

By the end, I can explain evaluation sets using a traced example and a stated limitation.

This added worked case illustrates the week's main idea. Retain all original and connected topics; schedule its practice within the teaching cycle.

2. Explain the mechanism

Explain the mechanism

Input → Process → Output

Trace the information. At each stage, ask what changed and how you could test it.

Which stage could fail even when the other two are correct? Explain:

Predict first. Then check each step. A correct answer without visible reasoning is incomplete evidence.

1. Ten cases include factual, comparison, calculation, ambiguous, and missing-source questions.
2. Add conflicting-source, prompt-injection, sensitive-data, and accessibility cases.
3. Write expected behavior before running the product.
4. After a repair, keep the failed case as a regression test.

Recalculate, retrace, or restate the reasoning in your own representation:

CHECK Did the evidence answer the exact question? Did the total, path, or source record stay consistent?

Concrete teaching case

Ten evaluation questions are easy definitions; none lacks a source.

What is untested?

Exact explanation to model

Absent evidence, ambiguity and boundary handling. Add predefined fallback cases and keep final cases separate from tuning.

Explain the added worked example

A system evaluation set contains cases with expected behaviour and scoring rules. It should cover retrieval, support, fallback and permissions rather than only easy answerable questions.

Begin by writing the question the procedure is intended to answer. Identify the information already supplied, the information it will produce and the information still missing. These are constructed classroom examples; they do not require live accounts, private records or an external service. Do not treat a plausible explanation as a measured outcome.

Record case ID, evidence version, expected claims or fallback, permitted actions and scoring rubric. Separate development from final cases. Aggregate metrics only after checking required conditions individually.

Before applying the procedure, write its units, denominator or decision boundary where relevant. A partner should be able to repeat each step using the same inputs. If a required input is absent or violates the stated contract, record that condition explicitly instead of quietly replacing it with a convenient value.

Key terms

Term · Meaning

Term: Evaluation case

Meaning: An input with evidence and prespecified expected behaviour

Term: Coverage

Meaning: Which relevant conditions are represented

Term: Contract pass

Meaning: Satisfying the declared requirements for one case

Apply the connected idea: Ranking and retrieval evaluation

A ranking needs task-specific relevance judgments and a declared cutoff. Precision and recall have different denominators.

Begin by writing the question the procedure is intended to answer. Identify the information already supplied, the information it will produce and the information still missing. These are constructed classroom examples; they do not require live accounts, private records or an external service. Do not treat a plausible explanation as a measured outcome.

Precision at k is relevant returned items divided by k when exactly k items are returned. Recall at k is relevant returned items divided by all known relevant eligible items. If the relevant set is unknown, do not invent recall.

Before applying the procedure, write its units, denominator or decision boundary where relevant. A partner should be able to repeat each step using the same inputs. If a required input is absent or violates the stated contract, record that condition explicitly instead of quietly replacing it with a convenient value.

Guided: Precision@3=2/3; recall@3=2/3. Duplicates should not count as distinct retrieved documents.

Qualification: Known judgments may be incomplete; a score over one query does not establish coverage over all queries. Exit: accept a concrete variation with its predicted outcome and a stated condition that would challenge the mechanism; reject a claim resting on unmeasured evidence.

Model the reasoning aloud

Use the complete worked example below before asking learners to complete W2. Reveal a small part, invite a prediction, then explain the deciding step. These teacher notes contain solutions; do not reveal the independent case answer during modelling.

Ten cases include six answerable and four missing-evidence cases. System answers five answerable correctly and abstains appropriately on three missing cases.

Answer correctness 5/6; appropriate fallback 3/4; combined contract passes 8/10 if each case has one required outcome. Different denominators answer different questions.

Compare the intermediate evidence with the final statement. Explain why each number or decision follows from the stated inputs. The conclusion belongs to this scenario and procedure; changing the inputs, evaluation population or authority can change what it establishes. Retain the raw record so a later revision can be compared honestly.

Thinking move · Teacher prompt

Thinking move: Plan

Teacher prompt: What is given? What is the question? Which procedure fits these facts?

Thinking move: Do

Teacher prompt: Point to each input and intermediate step in the worked trace. Name the rule that allows the next step.

Thinking move: Check

Teacher prompt: Does the output answer the question? Which assumption or limitation prevents a larger claim?

Use the diagram and verbal explanation together. Ask learners to match a sentence to a specific visual step; do not assign learners fixed visual or auditory learning styles.

A pictorial worked case: Evaluation Sets

Design an evaluation set that checks answering and appropriate uncertainty separately.

The facts we will use

Six toy questions: three have source answers and three do not. Results: 2 supported answers out of 3; 3 correct “not in source” responses out of 3.

All inputs and outcomes in this worked example are invented classroom data; no real model run is claimed.

Evaluation Sets

Start with the supplied facts

Follow each arrow. Ask which fact allows the next step.



1. Separate cases

Keep answerable and missing-evidence cases as different groups.



2. Calculate rates

Answering: 2/3. Appropriate uncertainty: 3/3.



3. Report total

Five of six pass overall; the perfect missing-evidence group hides one answering failure if only totals are discussed.

CHANGE ONE THING

Add three missing-evidence cases that all pass.

Explain the changed result using the matching lesson.

Read the picture one step at a time

Step · What happens and why

Step: 1. Separate cases

What happens and why: Keep answerable and missing-evidence cases as different groups.

Step: 2. Calculate rates

What happens and why: Answering: 2/3. Appropriate uncertainty: 3/3.

Step: 3. Report total

What happens and why: Five of six pass overall; the perfect missing-evidence group hides one answering failure if only totals are discussed.

Match each pictured step to the supplied facts. This is a worked teaching case, not an independent assessment.

Why this result follows

Read the picture in order. Five of six pass overall; the perfect missing-evidence group hides one answering failure if only totals are discussed. The arrows show which recorded input or intermediate result each decision uses. Explain the deciding step aloud before checking the changed case; pointing to the final answer alone does not show how it follows.

What this example does not establish

A small constructed set probes specific behaviours; its pooled rate is not an estimate of every future classroom question.

Change one thing in the pictured case

Change: Add three missing-evidence cases that all pass.

Prediction: Overall becomes 8/9, but supported answering remains 2/3.

Reason: Changing the case mixture raises the pooled rate without improving any existing answerable result.

Ask while showing the picture

Cover the result. Ask learners to name the inputs and predict the next step. Uncover one step, connect the icon to the actual procedure, and explain its evidence. Then change the supplied condition and compare. The picture is a representation; it does not document execution of a real AI model.

3. Demonstrate and guide practice

Demonstrate and guide practice

Board sequence

Write the given facts: Ten evaluation questions are easy definitions; none lacks a source.

Write the question to explain: What is untested?

Reveal after predictions: Absent evidence, ambiguity and boundary handling. Add predefined fallback cases and keep final cases separate from tuning.

Keep for an independent attempt: You improve a prompt after reading every test question. Can those

questions still provide an untouched final evaluation? Propose a repair.

Classroom dialogue

Ask: Which facts are given, and which would we have to assume? Learners mark relevant details before discussing answers.

Say: Predict first. Show the rule, information path or calculation that would produce the result. A correct guess is not yet an explanation.

Reveal reasoning one step at a time. Explain why a tempting alternative fails. Write units and denominators for arithmetic, trace conditions for decisions, and identify supporting sources for claims.

In pairs, one learner solves while the other checks evidence; exchange roles. Remove the worked answer before independent assessment.

Model these guided questions

Guided: Answer correctness 5/6; appropriate fallback 3/4; combined contract passes 8/10 if each case has one required outcome. Different denominators answer different questions. **Qualification:** Evaluation coverage limits the claim; prompts tuned on these cases need new reserved final cases. **Exit:** accept a concrete variation with its predicted outcome and a stated condition that would challenge the mechanism; reject a claim resting on unmeasured evidence.

Case-specific teaching sequence

Stage · Teacher action

Stage: Session 1: 0–7 minutes

Teacher action: Retrieve the prerequisite; identify the input, procedure and evidence source.

Stage: 7–18 minutes

Teacher action: Explain section 1 and vocabulary. Complete W1 with a mechanism-based explanation.

Stage: 18–32 minutes

Teacher action: Trace the declared procedure and worked case. Complete W2, recording intermediate steps.

Stage: 32–45 minutes

Teacher action: Discuss the diagram with the worked input; use W3 to distinguish measured claims and limitations.

Stage: Session 2: 0–8 minutes

Teacher action: Retrieve the worked reasoning and collect the guided response before discussing it.

Stage: 8–20 minutes

Teacher action: Compare a proposed boundary variant with the original case; identify what remains fixed.

Stage: 20–35 minutes

Teacher action: Collect W4 independently. Require the procedure trace and the bounded conclusion.

Stage: 35–45 minutes

Teacher action: Use W5 and exit responses for feedback. Administer the fresh retry after feedback and preserve both attempts.

Teach the connected idea

Stage · Teacher action

Stage: Session 1: 0–7 minutes

Teacher action: Retrieve the prerequisite; identify the input, procedure and evidence source.

Stage: 7–18 minutes

Teacher action: Explain section 1 and vocabulary. Complete W1 with a mechanism-based explanation.

Stage: 18–32 minutes

Teacher action: Trace the declared procedure and worked case. Complete W2, recording intermediate steps.

Stage: 32–45 minutes

Teacher action: Discuss the diagram with the worked input; use W3 to distinguish measured claims and limitations.

Stage: Session 2: 0–8 minutes

Teacher action: Retrieve the worked reasoning and collect the guided response before discussing it.

Stage: 8–20 minutes

Teacher action: Compare a proposed boundary variant with the original case; identify what remains fixed.

Stage: 20–35 minutes

Teacher action: Collect W4 independently. Require the procedure trace and the bounded conclusion.

Stage: 35–45 minutes

Teacher action: Use W5 and exit responses for feedback. Administer the fresh retry after feedback and preserve both attempts.

Reduce help deliberately

Completion task: Ten cases include six answerable and four missing-evidence cases. System answers five answerable correctly and abstains appropriately on three missing cases. Show the intermediate steps and give the conclusion with its scope.

Teacher check: Teacher explanation: Answer correctness 5/6; appropriate fallback 3/4; combined contract passes 8/10 if each case has one required outcome. Different denominators answer different questions.

Start by identifying the given inputs together. Leave the intermediate steps blank. If needed, model one step, then withdraw that help on the next step. Move to a full trace only when the partial trace is understood. For partners, alternate explainer and checker; collect a response from every learner before discussion.

Individual reasoning check

Use the worked case for System evaluation sets. Proposed explanation: “The worked result establishes every future case”. Decide whether this explanation is supported. Point to the step or supplied fact that decides; explain your repair if needed.

Expected correction: Evaluation coverage limits the claim; prompts tuned on these cases need new reserved final cases.

Collect explanations, not only yes/no votes. If a misunderstanding is common, re-model the relevant step for the class. If it affects a few learners, give a short targeted group explanation while others solve a more demanding variation.

Connected teaching cycle: Ranking and retrieval evaluation

Target: Explain the mechanism and terminology of ranking and retrieval evaluation

Reserve further sessions for this topic; the 135-minute primary plan does not absorb all added content.

Completion practice: Relevant documents are A, C, E. Top three results A, B, C contain two relevant items. Show the intermediate steps and give the conclusion with its scope.

Teacher check: Teacher explanation: $\text{Precision}@3=2/3$; $\text{recall}@3=2/3$. Duplicates should not count as distinct retrieved documents.

Individual check: For Ranking and retrieval evaluation, use its worked case to evaluate this proposed explanation: “The worked result establishes every future case”. Explain the deciding step.

Expected correction: Known judgments may be incomplete; a score over one query does not establish coverage over all queries.

Apply the same support, independent transfer and delayed-review decisions to the named connected-topic workbook tasks.

Pictorial practice question

Did the added cases demonstrate better answering? Point to the relevant step in the picture, then write the changed result and the rule or evidence that supports it.

Reasoned teacher answer

No. Answerable cases still pass 2/3. Report both group rates and the changed evaluation mixture.

Changed-case check: Overall becomes 8/9, but supported answering remains 2/3. Changing the case mixture raises the pooled rate without improving any existing answerable result.

Changed-case reasoning: Overall becomes 8/9, but supported answering remains 2/3. Changing the case mixture raises the pooled rate without improving any existing answerable result.

Collect each learner’s explanation before discussion. Use the first missing connection between input, deciding step and result to choose support. This question is worked-case practice, not fresh mastery evidence.

4. Run the weekly evidence lab

Run the weekly evidence lab

Words You Can Use to Reason

A definition matters when you can recognize the idea, use it, and contrast it with a near-miss.

Predict → Test → Record → Explain

Keep expected and actual results separate. Surprises are useful data when recorded honestly.

Map user tasks and risks, author ten labeled cases, peer-audit missing coverage, then freeze expected behaviors.

What did the hardest test teach you that the normal test did not?

Make the lab runnable

Use the concrete case as the shared normal case and the independent question as a second case. Each pair creates a boundary or missing-information case and predicts its result before testing. Keep the expected result separate from an observed mistake.

Record · What to capture

Record: Input or evidence

What to capture: Exact data, instruction or claim

Record: Expected result

What to capture: Answer or safe fallback declared before the run

Record: Observed result

What to capture: Actual result, including errors

Record: Explanation

What to capture: Rule, calculation or evidence explaining the outcome

Record: Next test

What to capture: One justified change and a new prediction

Diagnose and reteach

For guessing, return to given facts and request a trace. For vocabulary without application, use one example and one non-example. For arithmetic errors, check the setup before the calculation. Finish with a different uncoached case before recording mastery.

Address these additional misconceptions

The worked result establishes every future case: Evaluation coverage limits the claim; prompts tuned on these cases need new reserved final cases.

A missing input may be replaced without recording the change: A replacement changes the evidence or contract. Record the change and its justification; do not silently turn unknown into a measured value.

Repeating an assessed case is new independent evidence: A replay checks consistency. An example used for feedback or revision cannot later be described as untouched final evaluation.

Check these misconceptions

The worked result establishes every future case: Known judgments may be incomplete; a score over one query does not establish coverage over all queries.

A missing input may be replaced without recording the change: A replacement changes the evidence or contract. Record the change and its justification; do not silently turn unknown into a measured value.

Repeating an assessed case is new independent evidence: A replay checks consistency. An example used for feedback or revision cannot later be described as untouched final evaluation.

Match the support to the first error

What the response shows · Next teaching action

What the response shows: Missing or misread facts

Next teaching action: Read the case aloud, mark supplied versus unknown facts, and let the learner restate it.

What the response shows: Wrong procedure

Next teaching action: Return to the deciding step in the worked case. Contrast it with the proposed incorrect explanation.

What the response shows: Arithmetic or implementation error

Next teaching action: Trace intermediate values or lines; preserve the first mismatch and retry that step.

What the response shows: Correct result without a reason

Next teaching action: Ask for a trace and a changed case before marking independent understanding.

What the response shows: An unsupported wider claim

Next teaching action: Ask which supplied evidence supports the claim and which further observation would be needed.

Offer an oral explanation, drawing, enlarged visual or fewer examples when that removes a reading/writing barrier. Keep the same reasoning goal. Stretch secure learners by changing a boundary or assumption

and asking them to defend the effect, not simply by assigning more of the same questions.

5. Assess and explain the answers

Assess and explain the answers

Worked case answer

Absent evidence, ambiguity and boundary handling. Add predefined fallback cases and keep final cases separate from tuning.

Independent question

You improve a prompt after reading every test question. Can those questions still provide an untouched final evaluation? Propose a repair.

Expected independent answer

They have become development examples. Keep a separate unseen evaluation set and report the development results honestly.

What counts as evidence

Criterion · 0 points · 1 point · 2 points

Criterion: Result

0 points: Incorrect or absent

1 point: Partly correct

2 points: Correct with necessary conditions

Criterion: Reasoning

0 points: No supporting trace

1 point: Some steps

2 points: Clear mechanism, calculation or source support

Criterion: Limit

0 points: Overclaims certainty

1 point: Vague limitation

2 points: Relevant uncertainty or missing fact

Criterion: Transfer

0 points: Copies without adapting

1 point: Adapts with prompts

2 points: Solves a changed case independently

Use 6/8 as a classroom checkpoint, with 2 required for both Result and Reasoning. This is a proposed teaching rule, not a validated certification threshold. Do not count a coached answer as independent mastery; correct unsafe or unsupported claims regardless of total.

Feedback

Ask the learner to find the evidence that challenges this claim: “A test set stays independent after its answers guide revisions.” Keep the first attempt and an uncoached retry separate.

Original transfer answer

Add an absent fact requiring 'not enough evidence'; an ambiguous referent requiring clarification; and conflicting notes requiring an attributed conflict rather than an invented reconciliation. Keep final cases separate from tuning and specify acceptable answers before runs.

Explanations for the additional workbook examples

W1 Explain the mechanism

Explain the weekly mechanism in your own words. Use the worked scenario to name its input, decision rule and output.

Teacher explanation: A system evaluation set contains cases with expected behaviour and scoring rules. It should cover retrieval, support, fallback and permissions rather than only easy answerable questions. Record case ID, evidence version, expected claims or fallback, permitted actions and scoring rubric. Separate development from final cases. Aggregate metrics only after checking required conditions individually.

W2 Trace the worked case

Ten cases include six answerable and four missing-evidence cases. System answers five answerable correctly and abstains appropriately on three missing cases. Show the intermediate steps and give the conclusion with its scope.

Teacher explanation: Answer correctness 5/6; appropriate fallback 3/4; combined contract passes 8/10 if each case has one required outcome. Different denominators answer different questions.

W3 Audit the limit

A report omits this qualification: Evaluation coverage limits the claim; prompts tuned on these cases need new reserved final cases. Explain how omitting it could change a reader's interpretation, and write an accurate replacement.

Teacher explanation: Evaluation coverage limits the claim; prompts tuned on these cases need new reserved final cases. State the observed result separately from the unverified or out-of-scope claim; preserve the original evidence rather than inventing a missing fact.

W4 Apply a new case independently

Twelve cases: eight answerable, four unanswerable. Correct answers six; appropriate abstentions two. Compute separate and combined pass rates.

Teacher explanation: $6/8=75\%$; $2/4=50\%$; combined $8/12=2/3$. Do not call all twelve answer-accuracy cases.

W5 Design a revealing follow-up

For the worked scenario, propose one additional case that could reveal a failure, describe its expected outcome before running it, and explain what would need to remain fixed for a fair comparison. Do not reuse W4 as a new final test after receiving feedback.

Teacher explanation: Accept a concrete variant of the worked input or a missing/conflicting/boundary condition, with an expectation derived from the declared procedure. Require preserved settings, original outputs and data roles. Evaluation coverage limits the claim; prompts tuned on these cases need new reserved final cases. The new case should test this particular mechanism rather than merely restate the original result.

Ranking and retrieval evaluation — W1 Explain the mechanism

Explain the weekly mechanism in your own words. Use the worked scenario to name its input, decision rule and output.

Teacher explanation: A ranking needs task-specific relevance judgments and a declared cutoff. Precision and recall have different denominators. Precision at k is relevant returned items divided by k when exactly k items are returned. Recall at k is relevant returned items divided by all known relevant eligible items. If the relevant set is unknown, do not invent recall.

Ranking and retrieval evaluation — W2 Trace the worked case

Relevant documents are A, C, E. Top three results A, B, C contain two relevant items. Show the intermediate steps and give the conclusion with its scope.

Teacher explanation: $\text{Precision}@3=2/3$; $\text{recall}@3=2/3$. Duplicates should not count as distinct retrieved documents.

Ranking and retrieval evaluation — W3 Audit the limit

A report omits this qualification: Known judgments may be incomplete; a score over one query does not establish coverage over all queries. Explain how omitting it could change a reader's interpretation, and write an accurate replacement.

Teacher explanation: Known judgments may be incomplete; a score over one query does not establish coverage over all queries. State the observed result separately from the unverified or out-of-scope claim; preserve the original evidence rather than inventing a missing fact.

Ranking and retrieval evaluation — W4 Apply a new case independently

Relevant set A, D, F, G; top three A, B, D. Compute precision and recall. Explain which omitted records affect recall.

Teacher explanation: $\text{Precision@3} = 2/3$; $\text{recall@3} = 2/4 = 1/2$. F and G are relevant omissions. The metric uses this judged set only.

Ranking and retrieval evaluation — W5 Design a revealing follow-up

For the worked scenario, propose one additional case that could reveal a failure, describe its expected outcome before running it, and explain what would need to remain fixed for a fair comparison. Do not reuse W4 as a new final test after receiving feedback.

Teacher explanation: Accept a concrete variant of the worked input or a missing/conflicting/boundary condition, with an expectation derived from the declared procedure. Require preserved settings, original outputs and data roles. Known judgments may be incomplete; a score over one query does not establish coverage over all queries. The new case should test this particular mechanism rather than merely restate the original result.

Assess independent transfer

Selected case: Approved notes: Robotics Tuesday Room2; Art Wednesday Room5. Requests: "Where does the club meet?" and "Who leads drama?" An unversioned added note says Robotics Room3. Design checks and safe responses.

Selected case answer

Clarify which club for the ambiguous request; the leader is absent, so abstain or request a source. Flag the conflicting Robotics room and resolve source authority/version before answering it. Include ambiguity, absent facts and conflicts in evaluation; easy answered cases alone do not establish robustness.

Dimension · Evidence required

Dimension: Explanation

Evidence required: Names the applicable idea or procedure and ties it to the supplied facts.

Dimension: Trace and result

Evidence required: Preserves the given inputs and shows the relevant intermediate steps; distinguishes a prediction from an observed execution.

Dimension: Limits

Evidence required: States the assumptions, missing evidence or conditions under which the result would change.

Record each dimension as not yet, with support, or independent. Accept a valid alternative procedure with a defensible trace. Do not penalise an honestly recorded model failure if the learner correctly explains its evidence and limits.

Feedback that the learner uses

Name the earliest unsupported step and give one action to repair it. Have the learner make that repair now and explain why. Then assess the relevant skill on a fresh, fully specified case. Confirm it was not previously practised with feedback; solve and check the new case before assigning it. A follow-up design prompt is a starting brief, not automatically an equivalent test.

6. Connect and extend learning

Connect and extend learning

Verified Study Buddy

Maintain a balanced eval set covering utility, grounding, format, privacy, and safety.

Keep the decision, input, result, explanation and remaining limitation in the project evidence record. Retrieve this record next week rather than restarting the project.

Support and stretch

Provide one short approved passage; separate instruction, input, output and evidence. Stretch: introduce an unsupported claim or conflicting source and validate claim by claim with a fallback.

Exit reflection

Ask for one decision the learner can now justify and one situation requiring more evidence or human review. Confidence ratings are separate from correct independent performance.

Bridge to the next week

Next: Week 27: Red-Team the Product. Retrieve today's evidence and identify what carries forward and what is new.

Support and fresh reassessment

Provide an input/procedure/result/limit frame and a calculator when needed. Ask students to identify the first step before completing the whole trace. Extension: compare two concrete variants of this mechanism while holding unrelated conditions fixed; report both successful and unsuccessful outcomes.

W4: $6/8=75\%$; $2/4=50\%$; combined $8/12=2/3$. Do not call all twelve answer-accuracy cases. Require intermediate reasoning and the stated scope. Fresh retry: Five answerable and five unanswerable; four correct answers and four appropriate fallbacks. Expected reasoning: $4/5=80\%$ in each category; combined $8/10$. Preserve individual failures and permission checks. If students need support, record it and reassess with another new case rather than awarding independence to a copied answer.

Next connection: Capstone requirements and corpus. Preserve the current case record and retrieve its evidence boundary before moving on.

Readiness and reassessment for Ranking and retrieval evaluation

Provide an input/procedure/result/limit frame and a calculator when needed. Ask students to identify the first step before completing the whole trace. Extension: compare two concrete variants of this mechanism while holding unrelated conditions fixed; report both successful and unsuccessful outcomes.

W4: Precision@3= $2/3$; recall@3= $2/4=1/2$. F and G are relevant omissions. The metric uses this judged set only. Require intermediate reasoning and the stated scope. Fresh retry: Relevant set B, E; top four A, B, C, D. Expected reasoning: Precision@4= $1/4$; recall@4= $1/2$. If no relevant items exist, report recall denominator zero under an explicit policy. If students need support, record it and reassess with another new case rather than awarding independence to a copied answer.

Plan the delayed return

Suggested return points in this level: ai-3/27, ai-3/29, ai-3/32. Use actual classroom dates, adjusting for holidays and gaps. One, three and six teaching weeks are scheduling choices, not a claimed optimal spacing.

At a delayed check, ask for the idea from memory and a changed application. Compare with the earlier independent response and record support used. For a new level, revisit the previous capstone before checking the new prerequisites.

Record · Teacher evidence

Record: Learner code and date

Teacher evidence: _____

Record: Concept and case used; prior exposure

Teacher evidence: _____

Record: Explanation / trace / limits

Teacher evidence: _____

Record: Support and first repair

Teacher evidence: _____

Record: Changed case and delayed outcome

Teacher evidence: _____

Record: Next teaching action

Teacher evidence: _____

Use the evidence to choose reteaching or the next step. A completed page, confidence rating or attractive project does not by itself demonstrate understanding.

Project evidence from this case

For your Verified Study Buddy, save the input, procedure, observed result and limitation of this check. Label this worked example as practice; use a different, previously unreviewed case when checking independent understanding.

7. Prepare, teach and assess

Teaching target: Explain the mechanism and terminology of system evaluation sets

Prepare the labelled worked-case picture, the matching workbook, fictional input cards and a place for individual responses. No paid AI account is required for the paper reasoning activities.

Session · Plan · Resources

Session: Session 1 · Understand · 45 min

Plan: 0–7: recall and prerequisite; 7–15: inspect the concrete case and predict; 15–30: teacher models the worked trace; 30–40: learners explain a step using the text and visual; 40–45: check the core idea.

Resources: Sections 1–6; record the brief recall in the workbook.

Session: Session 2 · Practise · 45 min

Plan: 0–5: retrieve the procedure; 5–15: W2 completion practice; 15–23: explain and check with a partner; 23–30: each learner answers the misconception check; 30–40: targeted reteaching or a harder variation; 40–45: prepare the investigation.

Resources: Sections 7–10; use the matching workbook W2 and misconception check.

Session: Session 3 · Investigate and transfer · 45 min

Plan: 0–5: retrieve the decision rule; 5–20: investigate or trace; 20–30: start the independent case; 30–38: connect the project; 38–43: act on feedback; 43–45: plan the return check.

Resources: Sections 9 and 11–14. Give extra time to finish the full task set; use Section 15 when the week includes a coding lab.

Before class, solve the supported practice: Ten cases include six answerable and four missing-evidence cases. System answers five answerable correctly and abstains appropriately on three missing cases. Show the intermediate steps and give the conclusion with its scope.

Reasoned practice answer: Teacher explanation: Answer correctness 5/6; appropriate fallback 3/4; combined contract passes 8/10 if each case has one required outcome. Different denominators answer different questions.

Check every learner before discussion: Use the worked case for System evaluation sets. Proposed explanation: “The worked result establishes every future case”. Decide whether this explanation is supported. Point to the step or supplied fact that decides; explain your repair if needed.

Misconception repair: Evaluation coverage limits the claim; prompts tuned on these cases need new reserved final cases.

Support: read the exact case aloud, enlarge the labelled picture, supply input cards and model only the first deciding step. Accept a spoken explanation or drawing; retain the reasoning goal.

Extension: Add three missing-evidence cases that all pass. Require a prediction and a reason before comparing. Expected result: Overall becomes 8/9, but supported answering remains 2/3. Reason: Changing the case mixture raises the pooled rate without improving any existing answerable result.

Independent assessment: withhold the reserved answer until the learner has attempted the new case. Record support separately from correctness.

Criterion · Not yet · With support · Independent

Criterion: Explanation

Not yet: Names an answer without a deciding fact

With support: Uses a hint to link fact and decision

Independent: Explains the deciding fact in own words

Criterion: Trace or test

Not yet: Cannot connect input and result

With support: Completes steps with prompts

Independent: Traces or tests a changed case accurately

Criterion: Limits and revision

Not yet: Claims more than evidence supports

With support: Names a limit after a prompt

Independent: States a defensible limit and useful next test

Project checkpoint: Prepare the demonstration. Gather reproducible instructions, evidence, access checks and remaining limitations.

Use the weekly teacher answer for topic-specific reasoning. A saved progress tick is an activity record; judge mastery from the learner's explanation and evidence.

8. Use feedback, labs and repair paths

Use the two formative questions as supported practice, not as independent mastery scores. Discuss the deciding evidence rather than accepting a guessed selection.

The feedback comes from the reviewed public worked case. Do not reveal the separately reserved independent teacher answer until the learner has attempted that case.

Ask students to predict before running code. Preserve errors and actual outputs; collect a changed test and an explanation. The branch-boundary and indexing lessons deliberately include failure cases.

Use the earlier worked-explanation links to address a specific missing building block, then return to the current question. Adapt this suggestion after observing the learner; it is not an automatic diagnosis.

Collect the completed-work portfolio as evidence. Record support separately from correctness, and assess explanation, trace/test and limitations. Progress ticks and quiz selections alone do not establish mastery.