Baselines First

Learn to understand, design, build, test and explain AI systems responsibly.

Project: Responsible Plant Classifier

TEACHER GUIDE · 1 OF 9

Baselines First

> Prepare Baselines First: check the target, prerequisites, resources and adaptable teaching schedule.

Teach the learning cycle

Primary target: Fit a **majority label baseline** using **training** labels only

Session · Sequence and minutes · Resources

> **Session:** Session 1 · Understand · 45 min
>
> **Sequence and minutes:** 0–7: recall and prerequisite; 7–15: inspect the concrete case and predict; 15–30: teacher models the worked trace; 30–40: learners explain a step using the text and visual; 40–45: check the core idea.
>
> **Resources:** Sections 1–6; record the brief recall in the workbook.

> **Session:** Session 2 · Practise · 45 min
>
> **Sequence and minutes:** 0–5: retrieve the procedure; 5–15: W2 completion practice; 15–23: explain and check with a partner; 23–30: each learner answers the misconception check; 30–40: targeted reteaching or a harder variation; 40–45: prepare the investigation.
>
> **Resources:** Sections 7–10; use the matching workbook W2 and misconception check.

> **Session:** Session 3 · Investigate and transfer · 45 min

> **Sequence and minutes:** 0–5: retrieve the decision rule; 5–20: investigate or trace; 20–30: start the independent case; 30–38: connect the project; 38–43: act on feedback; 43–45: plan the return check.
>
> **Resources:** Sections 9 and 11–14. Give extra time to finish the full task set; use Section 15 when the week includes a coding lab.

This is a starting allocation of 135 minutes, not a validated duration. Preserve all named topics. Schedule extra time for connected examples, the full investigation, independent cases and project building.

Repeated copies of a case share one recorded attempt; they are not new **evidence**.

Retrieval answer guidance

Dataset Cards

Without opening the old lesson, explain one key idea from Dataset Cards. Give an example and a mistake that the idea helps you avoid.

Look for: Inspect row meaning missing values and repeated records

Reference example, not the only acceptable wording: One row is intended to represent one indoor practice trip.

The source is a teacher-authored fictional export for this lesson, duration uses minutes, and None means unknown rather than zero.

Sampling and Intended Use

Without opening the old lesson, explain one key idea from Sampling and Intended Use. Give an example and a mistake that the idea helps you avoid.

Look for: Inspect row meaning missing values and repeated records

Reference example, not the only acceptable wording: One row is intended to represent one indoor practice trip.

The source is a teacher-authored fictional export for this lesson, duration uses minutes, and None means unknown rather than zero.

Python and Table Readiness

Without opening the old lesson, explain one key idea from Python and Table Readiness. Give an example and a mistake that the idea helps you avoid.

Look for: Specify a table schema

Reference example, not the only acceptable wording: activity: nonempty text; minutes: integer from 0 to 60 inclusive.

Use the student module and workbook section with this exact week title. The worked case, independent question and answers below form one aligned lesson sequence.

What the learner must demonstrate

> A dataset contains 27 healthy and 3 spotted leaves. An always-healthy **baseline** gets what **accuracy** and spotted recall? What must a new model improve upon?

Concept to explain

A model must be compared with a simple transparent strategy on the same untouched examples.

Term	Meaning
baseline	simple comparison method
majority class	most frequent training label
rule baseline	transparent fixed decision rule
incremental value	measured gain beyond baseline

Prepare the class

Use the supplied scenario, student illustrations and blank paper. For numeric work, offer counters or a calculator after learners show the setup. Fictional records are supplied; personal data and paid accounts are unnecessary.

Alternative compact allocation

> This compact option is not the full-week schedule. Use the main teaching cycle and allow extra sessions for the investigation, independent work and project.

Allow 80 minutes: recall 8; explanation 12; worked demonstration 15; guided practice 20; independent assessment 15; project and reflection 10. Extend the lab into a second session when needed.

Preparation for the connected explanation

Teach two 45-minute sessions. Fit a binary majority-label **baseline** from **training** labels and freeze its **prediction** before scoring. The invented set uses one frame per distinct trip; it is separate from Week 6 data.

Supplied candidate **prediction**s are comparison examples, not an implemented fitted model. Treat this as a first benchmark and **evidence** lesson, not a readiness test.

Readiness answer

Related frames can share conditions and objects. Keeping the trip together prevents that trip from appearing in both development and evaluation sets.

Printed lesson and workbook

Six **training** label cards and four validation cards kept in order

Optional Python 3 editor with no extra packages

Blank comparison table with row reference **baseline** candidate and matches

End goal for the pictorial case

By the end, I can explain **baseline**s first using a traced example and a stated limitation.

This added worked case illustrates the week’s main idea. Retain all original and connected topics; schedule its practice within the teaching cycle.

Explain the mechanism

> Model Baselines First using the supplied case, labelled mechanism and reasoning. Make the scope of each conclusion explicit.

Trace the information, fitted behavior, and evaluation proof separately.

Question to explain

QUESTION Build majority-class and one-**rule baseline**s, then calculate their results.

MODEL IDEA A model must be compared with a simple transparent strategy on the same untouched examples.

PROOF NEEDED Use an unseen case, visible calculation, error analysis, and a stated limitation.

How the mechanism works

UNDER THE HOOD Test set has 12 healthy and 8 spotted leaves. Always-healthy **baseline** gets 12 ÷ 20 = **60% accuracy** but 0% spotted recall.

Draw a fresh data → model → **prediction** → evaluation trace:

Predict first. Then check each step. A correct answer without visible **evidence** cannot be audited.

Step 1

STEP 1 Test set has 12 healthy and 8 spotted leaves.

Step 2

STEP 2 Always-healthy **baseline** gets 12 ÷ 20 = **60% accuracy** but 0% spotted recall.

Step 3

STEP 3 On the same 20 leaves, the visible-spot rule correctly labels 5 spotted leaves, misses 3 spotted leaves, correctly labels 10 healthy leaves and falsely marks 2 healthy leaves as spotted.

Accuracy=(5+10)/20=75%; spotted recall=5/8=62.5%.

Step 4

STEP 4 A **learned model** must beat relevant **baseline** metrics—not only one convenient total.

CHECK Verify the denominator, unit, split, class definition, and data version before interpreting a score.

Recalculate, retrace, or restate the **evidence** independently:

Concrete teaching case

The teaching case

A test set has eighteen healthy and two spotted leaves. **Baseline** always predicts healthy.

The teaching case

Calculate and critique **accuracy**.

Exact explanation to model

> ### Why the result follows
>
> 18/20=90%, but spotted recall is 0/2=0%. Compare with task-relevant metrics and the trained model.

Explain the added worked example

Explain the added worked example · 1

A learner reports that a classifier gets 80 percent correct. Is that useful? We need to know the task, the examples and what a simple method achieves on those same examples. A **baseline** is a reference method.

This week our **baseline** ignores each scene and always predicts the label that is most common in the **training** set.

Explain the added worked example · 2

Use a new teacher-authored fictional dataset for predicting clear or blocked lanes. Each listed row describes one frame from a different trip on a familiar route; **training** and validation trips do not overlap.

This is a new practice set, not the twelve-frame table from Week 6. All supplied labels have been resolved under the classroom rule. There are no review cases in this exercise, and no real navigation claim is made.

Explain the added worked example · 3

The **training** counts are four clear and two blocked. Our stated method therefore predicts clear for every new example. A **baseline** can still learn a value from **training** data; simple does not mean it uses no data.

If the two **training** counts tie, our rule chooses clear, declared before scoring.

The **output** is clear 2 4. The fixed **prediction** matches two of the four validation labels, so **accuracy** is 2/4 or 50 percent. The **training** majority share, 4/6, is not the validation **accuracy**.

The **baseline** does not inspect the validation labels to decide which constant to predict.

Explain the added worked example · 4

The code assumes nonempty **training** data with only clear and blocked labels. It counts matches on the supplied validation list; a percentage also requires a nonempty validation set.

Do not run it on empty **training** data and describe the default clear value as learned **evidence**. Other labels and unresolved review statuses need an explicit policy.

Key terms

Term	Meaning
Baseline	A simple reference method used to judge a more complex approach
Majority label	The label with more examples than the other label in this binary task
Accuracy	Correct predictions divided by the number of scored examples
Percentage point	A unit for the difference between two percentages

Model the reasoning aloud

Use the complete worked example below before asking learners to complete W2. Reveal a small part, invite a **prediction**, then explain the deciding step.

These teacher notes contain solutions; do not reveal the independent case answer during modelling.

Model the reasoning aloud · 5

Suppose a candidate method supplies blocked, clear, clear, clear for the four validation rows in the table. These are supplied **prediction**s for a comparison exercise; we have not fitted that candidate.

It gets three of four correct, or 75 percent. The **baseline** gets two of four, or 50 percent. The candidate is ahead by one correct example and 25 **percentage point**s on this set.

Model the reasoning aloud · 6

Use identical row IDs, reference labels, exclusions and metric for both methods. Match **prediction**s to their rows even if files are reordered. A score on easier examples is not a fair comparison.

Neither this four-row result nor beating one **baseline** proves that the candidate will perform well on future trips.

The candidate misses the third validation row. Which rows does the baseline miss? Explain how the candidate changes the number correct. Why is the improvement 25 percentage points rather than a claim of 25 percent relative improvement?

Model the reasoning aloud · 7

Record **training** counts, the fixed **prediction** and tie rule, validation IDs and labels, **prediction** lists and correct/total counts. Keep the final test separate from development decisions.

Next week we will search candidate rules using **training evidence**, then compare choices carefully.

Thinking move	Teacher prompt
Plan	What is given? What is the question? Which procedure fits these facts?
Do	Point to each input and intermediate step in the worked trace. Name the rule that allows the next step.
Check	Does the output answer the question? Which assumption or limitation prevents a larger claim?

Use the diagram and verbal explanation together. Ask learners to match a sentence to a specific visual step; do not assign learners fixed visual or auditory learning styles.

A pictorial worked case: Baselines First

Make a simple **baseline** earn its place before claiming a better classifier.

The facts we will use

Supplied facts

Training labels: healthy, healthy, spotted. Majority **baseline** predicts healthy. Validation references: healthy, spotted, healthy, spotted.

Supplied facts

All **input**s and outcomes in this worked example are invented classroom data; no real model run is claimed.

AI ACADEMY / AI-4/10

Baselines First

Start with the supplied facts
Follow each arrow. Ask which fact allows the next step.



1. Fit baseline

Training majority is healthy: 2 of 3 labels.

↓



2. Score baseline

Always healthy matches 2/4 validation cases: 50%.

↓



3. Compare candidate

Supplied candidate predictions healthy, spotted, spotted, spotted match 3/4: 75%.

CHANGE ONE THING

Change the third candidate prediction from spotted to healthy.

Explain the changed result using the matching lesson.

Follow the labelled picture with the matching explanation. Open the vector diagram

Read the picture one step at a time

STEP	WHAT HAPPENS AND WHY
1. Fit baseline	Training majority is healthy: 2 of 3 labels.
2. Score baseline	Always healthy matches 2/4 validation cases: 50%.
3. Compare candidate	Supplied candidate predictions healthy, spotted, spotted, spotted match 3/4: 75%.

Match each pictured step to the supplied facts. This is a worked teaching case, not an independent assessment.

Why this result follows

Why the result follows

Read the picture in order. Supplied candidate **prediction**s healthy, spotted, spotted, spotted match 3/4: 75%. The arrows show which recorded **input** or intermediate result each decision uses.

Explain the deciding step aloud before checking the changed case; pointing to the final answer alone does not show how it follows.

What this example does not establish

Keep the stated limit

These supplied **prediction**s illustrate comparison. No actual model was trained, and four

cases are limited **evidence**.

Change one thing in the pictured case

Change one thing

Change: Change the third candidate **prediction** from spotted to healthy.

Predicted result

Prediction: Candidate becomes correct on all 4 validation cases.

Why the result follows

Reason: Only the previously wrong third decision changes; **baseline** stays fixed at 2/4.

Ask while showing the picture

Ask while showing the picture · 8

Cover the result. Ask learners to name the **input**s and predict the next step. Uncover one step, connect the icon to the actual procedure, and explain its **evidence**. Then change the supplied condition and compare.

The picture is a representation; it does not document execution of a real AI model.

Demonstrate and guide practice

> Demonstrate Baselines First, reduce support deliberately and check individual reasoning before discussing answers.

Board sequence

1. Put the facts on the board

Write the given facts: A test set has eighteen healthy and two spotted leaves. **Baseline** always predicts healthy.

2. State the decision to explain

Write the question to explain: Calculate and critique **accuracy**.

3. Reveal the worked reasoning

Reveal after **prediction**s: 18/20=90%, but spotted recall is 0/2=0%. Compare with task-relevant metrics and the trained model.

Hold back the independent task

Keep for an independent attempt: A dataset has 18 grass and 2 fern examples. A **baseline** always predicts grass. Compute **accuracy** and name a serious limitation.

Classroom dialogue

Ask the class

Ask: Which facts are given, and which would we have to assume? Learners mark relevant details before discussing answers.

Language to model

Say: Predict first. Show the rule, information path or calculation that would produce the

> result. A correct guess is not yet an explanation.

Reveal reasoning one step at a time. Explain why a tempting alternative fails. Write units and denominators for arithmetic, trace conditions for decisions, and identify supporting sources for claims.

In pairs, one learner solves while the other checks **evidence**; exchange roles. Remove the worked answer before independent assessment.

Model these guided questions

Guided: the **baseline** misses rows 1 and 3; the candidate misses only row 3, adding one correct example. The difference is 75 minus 50 = 25 **percentage point**s.

Relative improvement in **accuracy** would be (75 minus 50)/50 = 50 percent; keep the measure explicit. Exit: equal **accuracy** does not establish identical mistakes or suitability.

Inspect **prediction**s by row, missed blocked cases, false blocked **prediction**s, sample coverage and consequences, then evaluate using an appropriate held out **process**.

Case-specific teaching sequence

Stage	Teacher action
Session 1 start to 7 minutes	Retrieve the split roles and related-trip issue. Ask what a score means without a reference method.
7 to 18 minutes	Count training labels and write the constant before revealing validation labels. State the tie rule before any comparison.
18 to 32 minutes	Trace the code. Complete W1 and W2, separating training share from validation accuracy.
32 to 45 minutes	Build a row-by-row comparison for the supplied candidate. Complete the guided percentage-point question.
Session 2 start to 8 minutes	Retrieve what changes if the baseline constant is chosen using validation labels.
8 to 20 minutes	Discuss the two unfair comparisons and complete W3. Require common rows, labels and scoring policy.

Stage	Teacher action
20 to 35 minutes	Collect W4 independently before hints. Require the training counts, both prediction lists, scores and restrained claim.
35 to 45 minutes	Use W5 and the exit question to compare error patterns. Give the fresh retry after feedback and record support.

Reduce help deliberately

Completion practice

Completion task: List the four **baseline prediction**s and mark which validation rows are correct. Give correct/total and **accuracy**. Explain why 4/6 is not the answer.

Reasoning to listen for

Teacher check: **Prediction**s are clear, clear, clear, clear. Rows 2 and 4 match, giving 2/4 or 50 percent. The fraction 4/6 describes clear labels in **training**, not correct **prediction**s on the four validation rows.

Start by identifying the given **input**s together. Leave the intermediate steps blank. If needed, model one step, then withdraw that help on the next step. Move to a full trace only when the partial trace is understood.

For partners, alternate explainer and checker; collect a response from every learner before discussion.

Individual reasoning check

Use the worked case for A **baseline** before a model. Proposed explanation: “A simple **baseline** does not need **training** data”. Decide whether this explanation is supported.

Point to the step or supplied fact that decides; explain your repair if needed.

Expected repair

Expected correction: This method learns its constant from **training** label counts, although it ignores **input** features.

Collect explanations, not only yes/no votes. If a misunderstanding is common, re-model the relevant step for the class.

If it affects a few learners, give a short targeted group explanation while others solve a more demanding variation.

Pictorial practice question

Can the candidate's **training accuracy** replace this validation comparison? Point to the relevant step in the picture, then write the changed result and the rule or **evidence** that supports it.

Reasoned teacher answer

No. Compare candidate and **baseline** on the same held-out cases and references; **training** fit answers a different question. Changed-case check: Candidate becomes correct on all 4 validation cases.

Only the previously wrong third decision changes; **baseline** stays fixed at 2/4.

Changed-case reasoning: Candidate becomes correct on all 4 validation cases. Only the previously wrong third decision changes; **baseline** stays fixed at 2/4.

Collect each learner's explanation before discussion. Use the first missing connection between **input**, deciding step and result to choose support. This question is worked-case practice, not fresh mastery **evidence**.

Run the weekly evidence lab

> Run the Baselines First investigation and use the first unsupported step to choose a teaching

> response.

Keep the test conditions fixed, record every result, and change only one planned factor.

1. Write the hypothesis and expected **evidence** before running.

2. **Baseline** Challenge: calculate two **baseline**s on one fixed set, compare class-specific errors, and choose the honest benchmark.

3. Run one normal, one boundary, one missing, and one shifted case.

4. Calculate the metric from raw counts or row-level errors.

5. Interpret what changed, what did not, and what remains unsafe or unknown.

Make the lab runnable

Use the concrete case as the shared normal case and the independent question as a second case. Each pair creates a boundary or missing-information case and predicts its result before testing.

Keep the expected result separate from an observed mistake.

Record	What to capture
Input or evidence	Exact data, instruction or claim
Expected result	Answer or safe fallback declared before the run

Record	What to capture
Observed result	Actual result, including errors
Explanation	Rule, calculation or evidence explaining the outcome
Next test	One justified change and a new prediction

Diagnose and reteach

MISCONCEPTION TO REPAIR Any trained model is more intelligent and therefore better than a simple rule.

For guessing, return to given facts and request a trace. For vocabulary without application, use one example and one non-example. For arithmetic errors, check the setup before the calculation.

Finish with a different uncoached case before recording mastery.

Address these additional misconceptions

A simple **baseline** does not need **training** data: This method learns its constant from **training** label counts, although it ignores **input** features.

Training majority share is validation **accuracy**: Fit using **training** labels, then count matches separately on validation rows. Their label proportions can differ.

Beating a **baseline** proves the system is ready: It is one comparison on specified examples. Error types, coverage, uncertainty and final evaluation remain relevant.

Match the support to the first error

What the response shows	Next teaching action
Missing or misread facts	Read the case aloud, mark supplied versus unknown facts, and let the learner restate it.
Wrong procedure	Return to the deciding step in the worked case. Contrast it with the proposed incorrect explanation.
Arithmetic or implementation error	Trace intermediate values or lines; preserve the first mismatch and retry that step.

What the response shows	Next teaching action
Correct result without a reason	Ask for a trace and a changed case before marking independent understanding.
An unsupported wider claim	Ask which supplied evidence supports the claim and which further observation would be needed.

Offer an oral explanation, drawing, enlarged visual or fewer examples when that removes a reading/writing barrier. Keep the same reasoning goal.

Stretch secure learners by changing a boundary or assumption and asking them to defend the effect, not simply by assigning more of the same questions.

Assess and explain the answers

> Use the Baselines First questions, teacher keys and criteria after the relevant attempts. Previously discussed cases are practice; protect the separately selected independent task.

Worked case answer

Teacher key · Worked case answer

18/20=90%, but spotted recall is 0/2=0%. Compare with task-relevant metrics and the trained model.

Independent question

Independent question · withhold the key

A dataset has 18 grass and 2 fern examples. A **baseline** always predicts grass. Compute **accuracy** and name a serious limitation.

Expected independent answer

> **Teacher key · Expected independent answer**
>
> **Accuracy** is 18/20 = 90%, but fern recall is 0/2. The high total hides complete failure on the minority class.

What counts as evidence

Criterion	0 points	1 point	2 points
Result	Incorrect or absent	Partly correct	Correct with necessary conditions
Reasoning	No supporting trace	Some steps	Clear mechanism, calculation or source support
Limit	Overclaims certainty	Vague limitation	Relevant uncertainty or missing fact
Transfer	Copies without adapting	Adapts with prompts	Solves a changed case independently

Use 6/8 as a classroom checkpoint, with 2 required for both Result and Reasoning. This is a proposed teaching rule, not a validated certification threshold.

Do not count a coached answer as independent mastery; correct unsafe or unsupported claims regardless of total.

Feedback

Ask the learner to find the **evidence** that challenges this claim: “A 90% **baseline** necessarily serves both classes well.” Keep the first attempt and an uncoached retry separate.

Original transfer answer

> **Teacher key · Original transfer answer**
>
> **Accuracy**=27/30=90%; spotted recall=0/3=0%. A new model should be compared on task-relevant metrics and error consequences, not only **accuracy**.

Fit a majority **baseline** using **training** labels, without using final labels to choose the class.

Explanations for the additional workbook examples

W1 Choose from training evidence

Workbook question · W1

Count clear and blocked in the worked **training** list. State the constant **prediction** and the declared tie rule. Explain why the validation majority does not choose this **baseline**.

Teacher reasoning · W1

Training has four clear and two blocked, so the constant is clear. A tie chooses clear under the declared classroom rule.

Validation labels score the stored **prediction**; using their majority to choose it would change the **fitting** procedure.

W2 Calculate the validation score

Workbook question · W2

List the four **baseline prediction**s and mark which validation rows are correct. Give correct/total and **accuracy**. Explain why 4/6 is not the answer.

Teacher reasoning · W2

Predictions are clear, clear, clear, clear. Rows 2 and 4 match, giving 2/4 or 50 percent. The fraction 4/6 describes clear labels in **training**, not correct **prediction**s on the four validation rows.

W3 Make the comparison fair

Workbook question · W3

A candidate is scored on ten easy rows and the **baseline** on four difficult rows. Describe a fair comparison and two details you must record.

Teacher reasoning · W3

Score both methods on the same intended validation examples with matched IDs, reference labels, metric and exclusion policy. Record **prediction**s, correct/total counts and the **training**-derived **baseline** rule.

Different example sets do not isolate a method difference.

W4 Fit and compare independently

Workbook question · W4

New **training** labels are blocked, clear, blocked, blocked, clear. New validation labels are clear, blocked, clear, blocked. Fit the same **baseline** rule and list its four **prediction**s.

A supplied candidate predicts clear, blocked, blocked, blocked. Score both on these rows and state the difference in correct examples and **percentage point**s. What claim remains unsupported?

Teacher reasoning · W4

Training counts are three blocked and two clear, so predict blocked four times. The **baseline** matches rows 2 and 4: 2/4 or 50 percent. The candidate matches rows 1,2,4: 3/4 or 75 percent.

It gains one correct example and 25 **percentage point**s. This tiny validation comparison does not prove future performance or readiness for use.

W5 Explain the hidden error

Workbook question · W5

A **training**-derived clear **baseline** gets nine of ten validation cases correct; the labels are nine clear and one blocked. State its **accuracy** and missed case. If a candidate also gets nine correct, must it miss the same case?

Explain.

Teacher reasoning · W5

Accuracy is 90 percent, and the **baseline** misses the sole blocked case. A candidate could also score 90 percent by correctly identifying blocked but missing one clear case.

Equal totals do not imply identical errors; inspect **prediction**s by row and error type.

Assess independent transfer

Independent question · withhold the key

Selected case: New **training** labels are blocked, clear, blocked, blocked, clear. New validation labels are clear, blocked, clear, blocked. Fit the same **baseline** rule and list its four **prediction**s.

A supplied candidate predicts clear, blocked, blocked, blocked. Score both on these rows and state the difference in correct examples and **percentage point**s. What claim remains unsupported?

Selected case answer

Teacher key · Selected case answer

Training counts are three blocked and two clear, so predict blocked four times. The **baseline** matches rows 2 and 4: 2/4 or 50 percent. The candidate matches rows 1,2,4: 3/4 or 75 percent.

It gains one correct example and 25 **percentage point**s. This tiny validation comparison does not prove future performance or readiness for use.

Dimension	Evidence required
Explanation	Names the applicable idea or procedure and ties it to the supplied facts.
Trace and result	Preserves the given inputs and shows the relevant intermediate steps; distinguishes a prediction from an observed execution.
Limits	States the assumptions, missing evidence or conditions under which the result would change.

Teacher key · Selected case answer

Record each dimension as not yet, with support, or independent. Accept a valid alternative procedure with a defensible trace.

Do not penalise an honestly recorded model failure if the learner correctly explains its **evidence** and limits.

Feedback that the learner uses

Name the earliest unsupported step and give one action to repair it. Have the learner make that repair now and explain why. Then assess the relevant skill on a fresh, fully specified case.

Confirm it was not previously practised with feedback; solve and check the new case before assigning it. A follow-up design prompt is a starting brief, not automatically an equivalent test.

Connect and extend learning

Connect Baselines First to Responsible Plant Classifier, adapt support and plan a later changed-case check.

Responsible Plant Classifier

Record majority and transparent-**rule baseline**s for the Plant Classifier using the frozen test plan.

Keep the decision, **input**, result, explanation and remaining limitation in the project **evidence** record. Retrieve this record next week rather than restarting the project.

Support and stretch

Use a tiny labelled dataset, explicit units and a named positive class. Stretch: require a **baseline**, held-out evaluation and subgroup or shifted-condition analysis with denominators.

Exit reflection

Ask for one decision the learner can now justify and one situation requiring more **evidence** or human review. Confidence ratings are separate from correct independent performance.

Bridge to the next week

Next: Week 11: Decision Trees. Retrieve today's **evidence** and identify what carries forward and what is new.

Support and fresh reassessment

Use labelled cards and matching ticks before fractions. Supply a blank row-aligned comparison table without **prediction**s.

Extension: construct two different **prediction** lists with identical **accuracy** but different missed rows, and explain how to report both.

Require W4 blocked from counts 3 versus 2, **baseline** 2/4, candidate 3/4, one extra correct and 25 **percentage point**s. Fresh retry: **training** clear, blocked, blocked, clear is tied, so choose clear by the declared rule.

Validation blocked, clear, blocked, blocked gives **baseline** 1/4 or 25 percent. Candidate blocked, clear, clear, blocked gives 3/4 or 75 percent: two extra correct and 50 **percentage point**s.

Require the tie rule and explain why validation majority blocked cannot replace the stored constant.

Week 8 searches candidate rules. Carry forward a frozen **training**-derived **baseline** and a fair comparison record.

Teacher reference: scikit-learn DummyClassifier documentation, most_frequent strategy, https://scikit-learn.org/dev/modules/generated/sklearn.dummy.DummyClassifier.html (reviewed 2026-09-27).

Plan the delayed return

Suggested return points in this level: ai-4/11, ai-4/13, ai-4/16. Use actual classroom dates, adjusting for holidays and gaps. One, three and six teaching weeks are scheduling choices, not a claimed optimal spacing.

At a delayed check, ask for the idea from memory and a changed application. Compare with the earlier independent response and record support used.

For a new level, revisit the previous capstone before checking the new prerequisites.

Record	Teacher evidence
Learner code and date	
Concept and case used; prior exposure	
Explanation / trace / limits	
Support and first repair	
Changed case and delayed outcome	
Next teaching action	

Use the **evidence** to choose reteaching or the next step. A completed page, confidence rating or attractive project does not by itself demonstrate understanding.

Project evidence from this case

For your Responsible Plant Classifier, save the **input**, procedure, observed result and limitation of this check.

Label this worked example as practice; use a different, previously unreviewed case when checking independent understanding.

Prepare, teach and assess

> Use this preparation summary alongside the detailed Baselines First lesson; it is a checklist rather than a replacement for the full sequence.

Teaching target

Teaching target: Fit a **majority label baseline** using **training** labels only

Prepare the labelled worked-case picture, the matching workbook, fictional **input** cards and a place for individual responses. Check a local Python 3 environment before class; have a paper-trace version ready.

Session · Plan · Resources

> **Session:** Session 1 · Understand · 45 min
>
> **Plan:** 0–7: recall and prerequisite; 7–15: inspect the concrete case and predict; 15–30: teacher models the worked trace; 30–40: learners explain a step using the text and visual; 40–45: check the core idea.
>
> **Resources:** Sections 1–6; record the brief recall in the workbook.

> **Session:** Session 2 · Practise · 45 min
>
> **Plan:** 0–5: retrieve the procedure; 5–15: W2 completion practice; 15–23: explain and check with a partner; 23–30: each learner answers the misconception check; 30–40: targeted reteaching or a harder variation; 40–45: prepare the investigation.
>
> **Resources:** Sections 7–10; use the matching workbook W2 and misconception check.

> **Session:** Session 3 · Investigate and transfer · 45 min
>
> **Plan:** 0–5: retrieve the decision rule; 5–20: investigate or trace; 20–30: start the independent case; 30–38: connect the project; 38–43: act on feedback; 43–45: plan the return check.

> **Resources:** Sections 9 and 11–14. Give extra time to finish the full task set; use Section 15 when the week includes a coding lab.

Before class, solve the supported practice: List the four **baseline prediction**s and mark which validation rows are correct. Give correct/total and **accuracy**. Explain why 4/6 is not the answer.

Supported practice answer

Reasoned practice answer: **Prediction**s are clear, clear, clear, clear. Rows 2 and 4 match, giving 2/4 or 50 percent. The fraction 4/6 describes clear labels in **training**, not correct **prediction**s on the four validation rows.

Check every learner before discussion: Use the worked case for A **baseline** before a model. Proposed explanation: "A simple **baseline** does not need **training** data". Decide whether this explanation is supported.

Point to the step or supplied fact that decides; explain your repair if needed.

Repair the misconception

Misconception repair: This method learns its constant from **training** label counts, although it ignores **input** features.

Support without giving away the answer

Support: read the exact case aloud, enlarge the labelled picture, supply **input** cards and model only the first deciding step. Accept a spoken explanation or drawing; retain the reasoning goal.

Extend the reasoning

Extension: Change the third candidate **prediction** from spotted to healthy. Require a **prediction** and a reason before comparing. Expected result: Candidate becomes correct on all 4 validation cases.

Reason: Only the previously wrong third decision changes; **baseline** stays fixed at 2/4.

Protect the independent attempt

Independent assessment: withhold the reserved answer until the learner has attempted the new case. Record support separately from correctness.

Criterion	Not yet	With support	Independent
Explanation	Names an answer without a deciding fact	Uses a hint to link fact and decision	Explains the deciding fact in own words
Trace or test	Cannot connect input and result	Completes steps with prompts	Traces or tests a changed case accurately
Limits and revision	Claims more than evidence supports	Names a limit after a prompt	States a defensible limit and useful next test

Project checkpoint

Project checkpoint: Make a transparent first version. Keep a trace of how one **input** becomes an **output**; explain the rule or learned component.

Use the weekly teacher answer for topic-specific reasoning. A saved progress tick is an activity record; judge mastery from the learner's explanation and **evidence**.

TEACHER GUIDE · 8 OF 9

Use feedback, labs and repair paths

Use feedback, earlier review and portfolios purposefully. Apply the coding routine only when this lesson has an executable lab.

Use the two formative questions as supported practice, not as independent mastery scores. Discuss the deciding **evidence** rather than accepting a guessed selection.

The feedback comes from the reviewed public worked case. Do not reveal the separately reserved independent teacher answer until the learner has attempted that case.

Ask students to predict before running code. Preserve errors and actual **output**s; collect a changed test and an explanation. The branch-boundary and indexing lessons deliberately include failure cases.

Use the earlier worked-explanation links to address a specific missing building block, then return to the current question. Adapt this suggestion after observing the learner; it is not an automatic diagnosis.

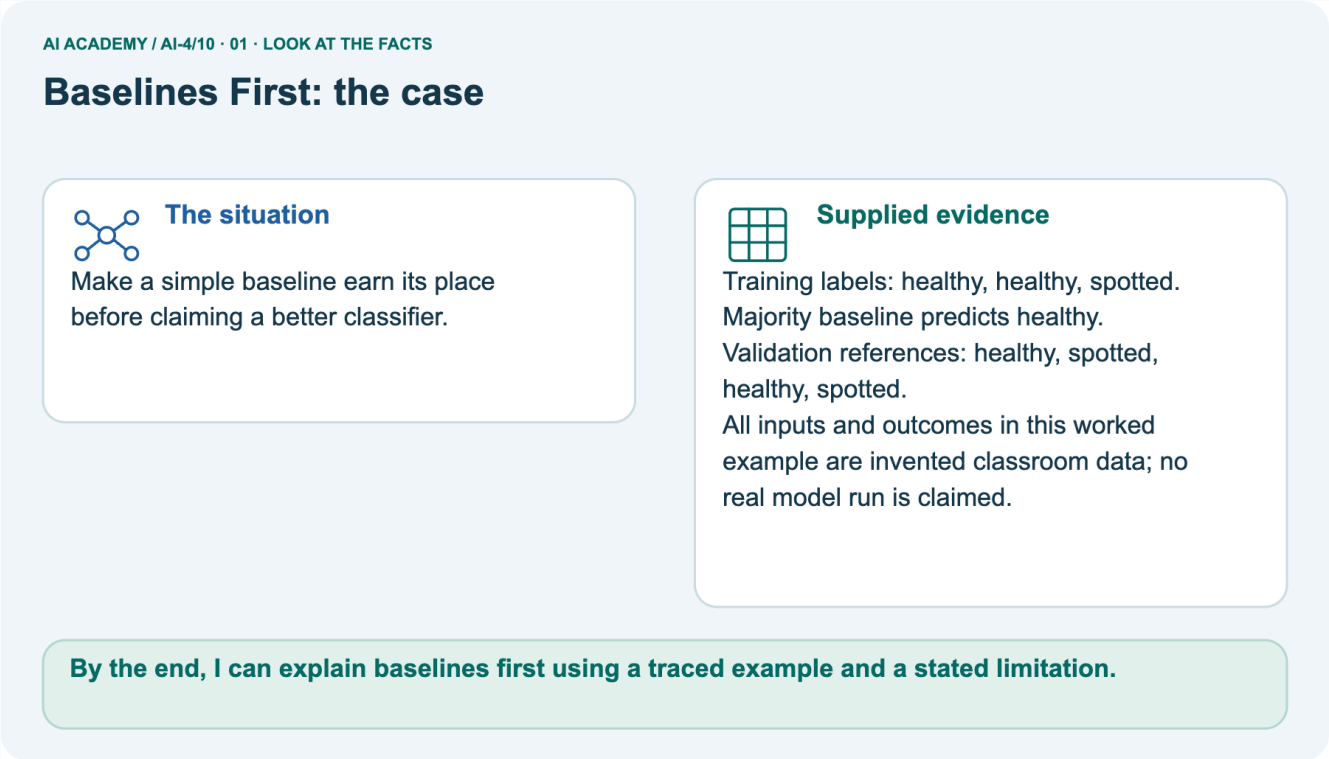
Collect the completed-work portfolio as **evidence**. Record support separately from correctness, and assess explanation, trace/test and limitations. Progress ticks and quiz selections alone do not establish mastery.

Visual concept journey

> Show one Baselines First graphic at a time, request a prediction and discuss its trace before the independent assessment.

Use one picture at a time. Hide a worked result before asking for a **prediction**, then reveal it for feedback. These are supported teaching cases; retain the separate independent assessment for a fresh first attempt.

01. Baselines First: the case



Follow the labelled picture with the matching explanation. Open the vector diagram

Make a simple **baseline** earn its place before claiming a better classifier. Supplied facts: **Training** labels: healthy, healthy, spotted. Majority **baseline** predicts healthy.

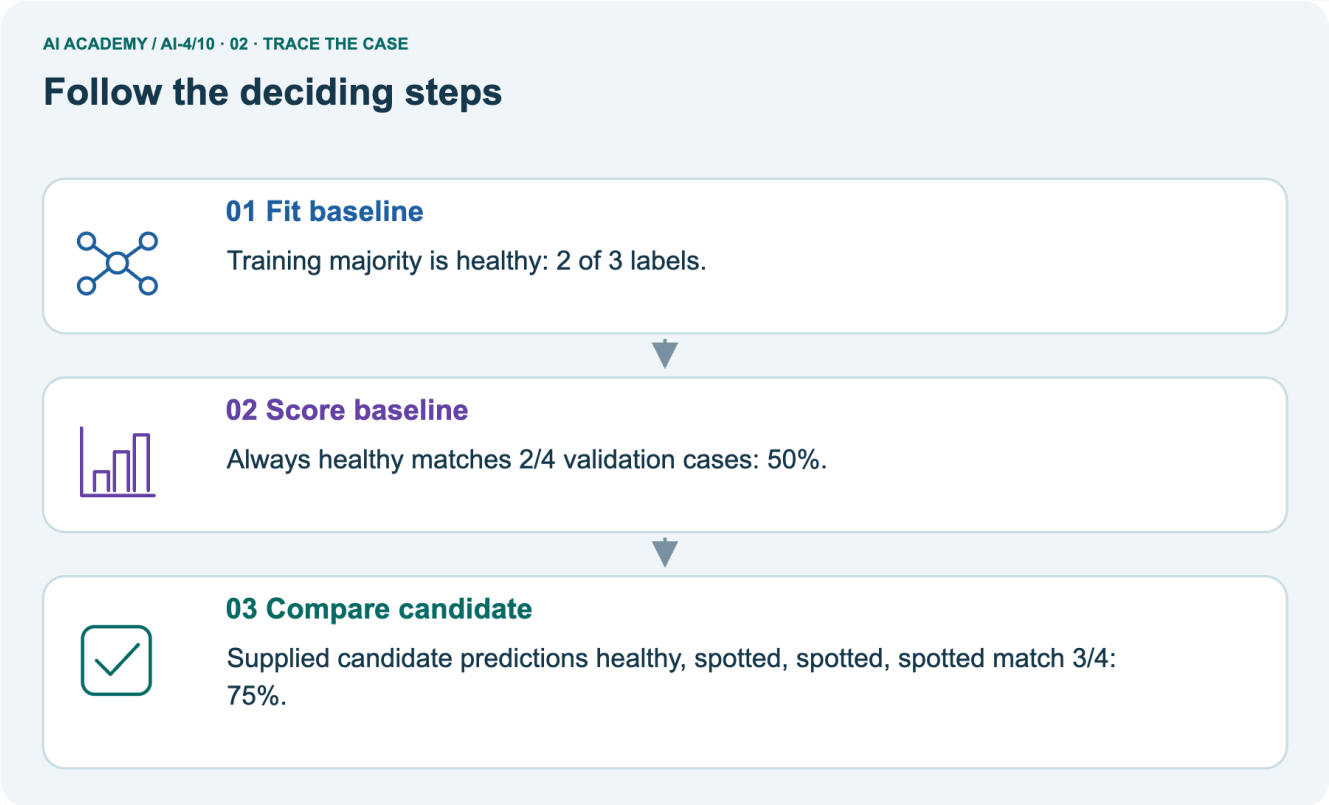
Validation references: healthy, spotted, healthy, spotted. All **input**s and outcomes in this worked example are invented classroom data; no real model run is claimed.

> Point to the given fact that matters most. Explain why it matters.

Discuss while showing the picture

Teacher prompt: Point to the given fact that matters most. Explain why it matters. Ask learners to describe one label or arrow in ordinary words. Use the adjacent transcript if the image is difficult to see.

02. Follow the deciding steps



Follow the labelled picture with the matching explanation. Open the vector diagram

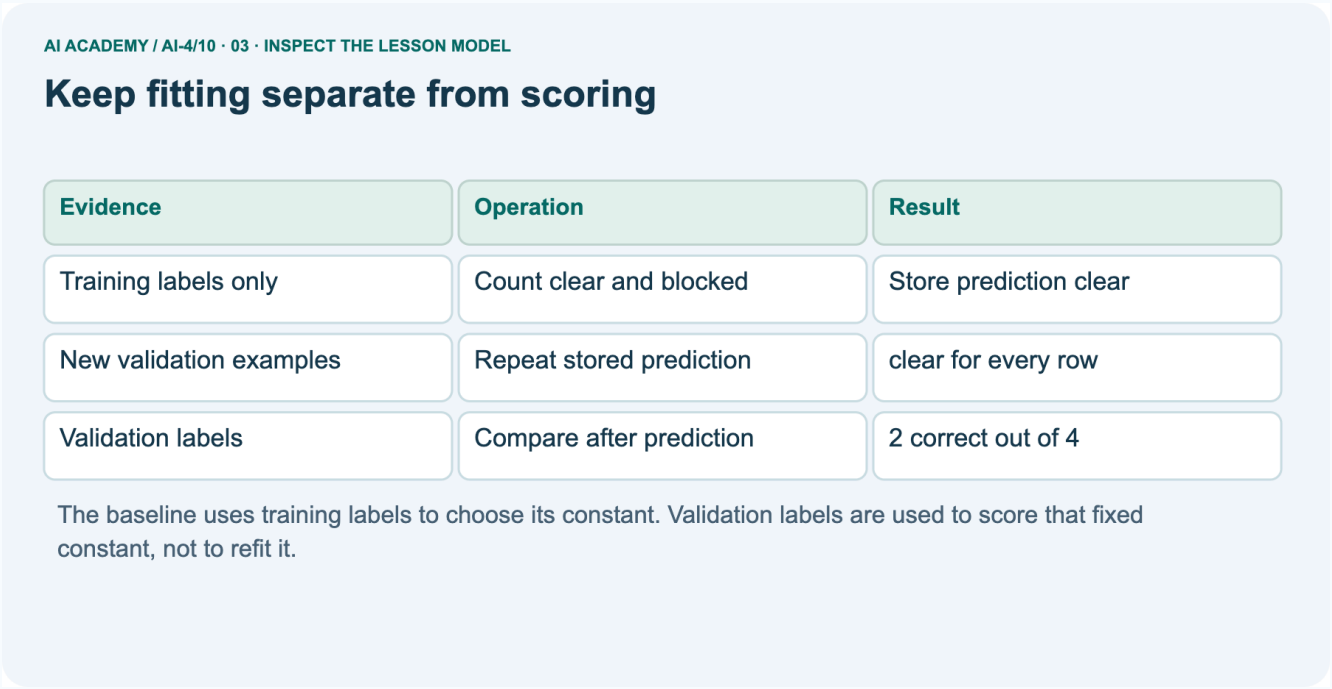
Fit **baseline**: **Training** majority is healthy: 2 of 3 labels. → Score **baseline**: Always healthy matches 2/4 validation cases: 50%. → Compare candidate: Supplied candidate **prediction**s healthy, spotted, spotted, spotted match 3/4: 75%.

Point to each numbered stage and explain which recorded **input** it uses.

Discuss while showing the picture

Teacher prompt: Point to each numbered stage and explain which recorded **input** it uses. Ask learners to describe one label or arrow in ordinary words. Use the adjacent transcript if the image is difficult to see.

03. Keep fitting separate from scoring



Follow the labelled picture with the matching explanation. Open the vector diagram

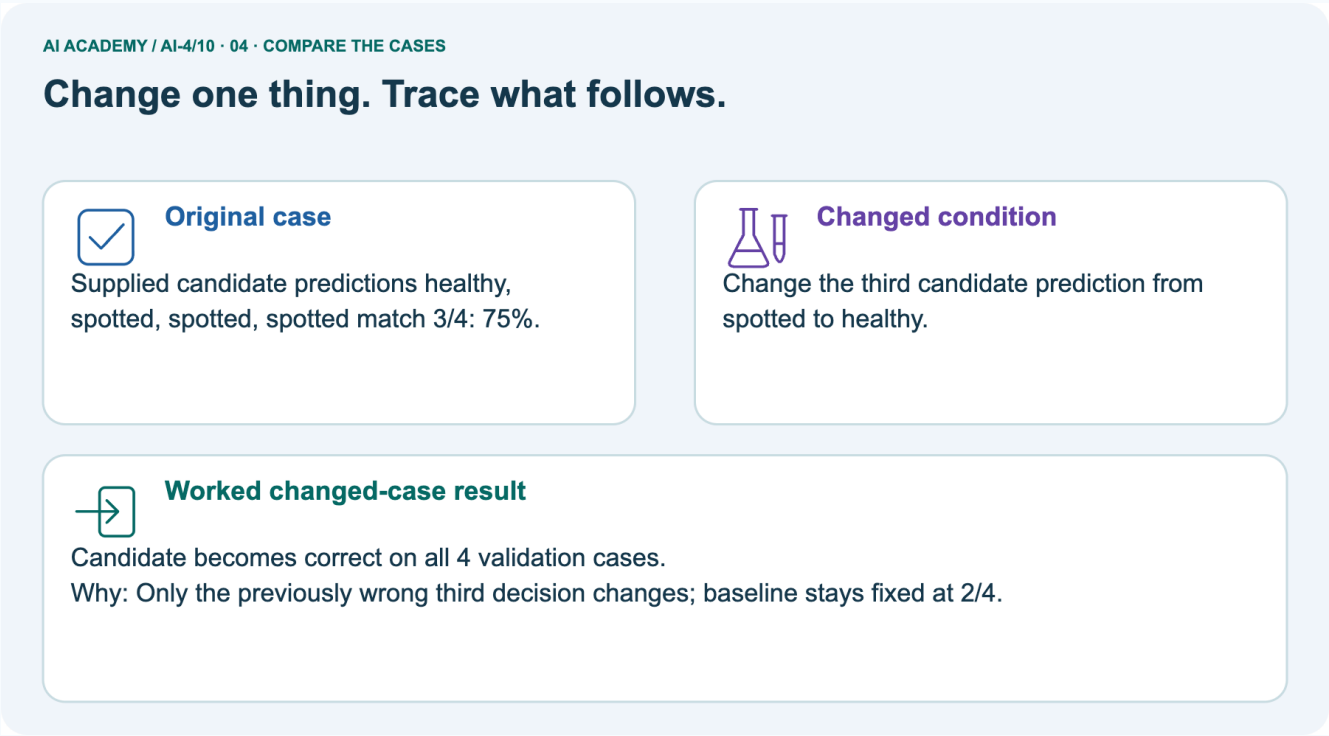
Keep **fitting** separate from scoring. **Training** labels only | Count clear and blocked | Store **prediction** clear. New validation examples | Repeat stored **prediction** | clear for every row.

Validation labels | Compare after **prediction** | 2 correct out of 4

> Read the labels and values. Explain the relationship the model shows.

> **Discuss while showing the picture**
>
> Teacher prompt: Read the labels and values. Explain the relationship the model shows. Ask learners to describe one label or arrow in ordinary words. Use the adjacent transcript if the image is difficult to see.

04. Change one thing. Trace what follows.



Follow the labelled picture with the matching explanation. Open the vector diagram

Original case: Supplied candidate **prediction**s healthy, spotted, spotted, spotted match 3/4: 75%.
Changed condition: Change the third candidate **prediction** from spotted to healthy.

Expected result: Candidate becomes correct on all 4 validation cases. Reason: Only the previously wrong third decision changes; **baseline** stays fixed at 2/4.

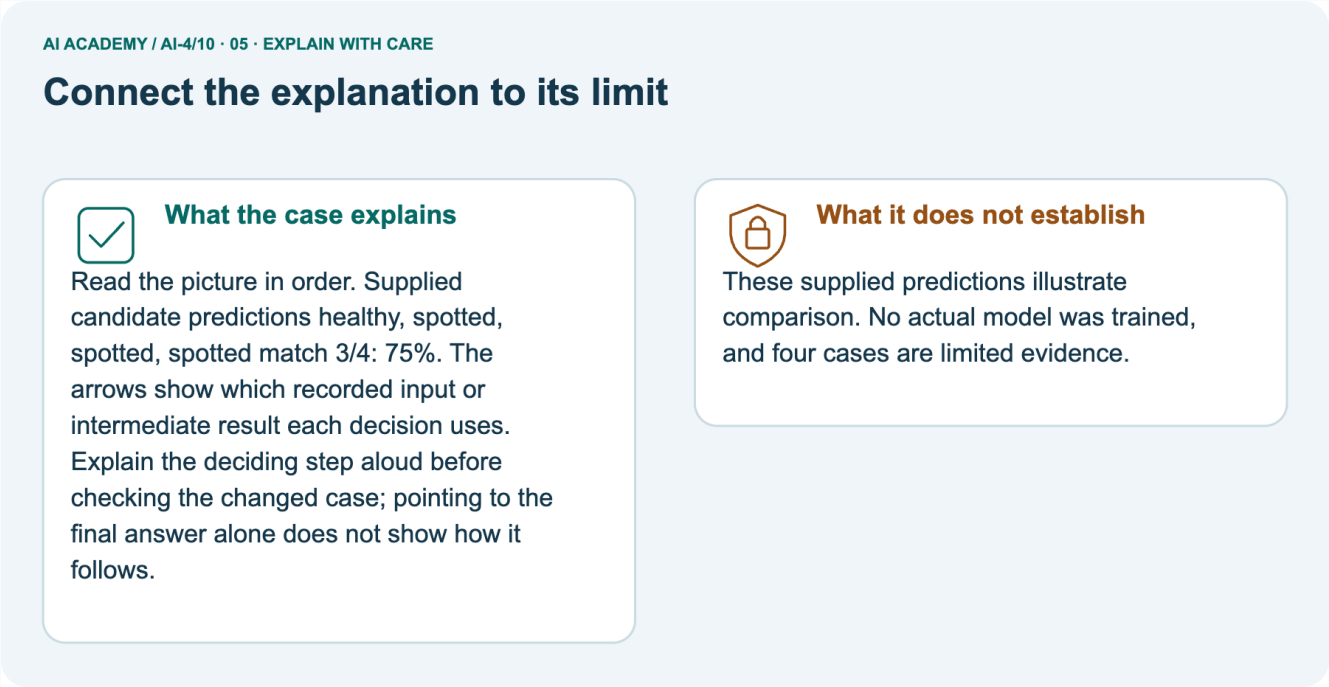
> Predict before checking the worked changed-case result. Name what changed and what the supplied facts allow you to conclude.

Discuss while showing the picture

Teacher prompt: Predict before checking the worked changed-case result. Name what changed and what the supplied facts allow you to conclude. Ask learners to describe one label or arrow in ordinary words.

Use the adjacent transcript if the image is difficult to see. Changed-case reasoning: Only the previously wrong third decision changes; **baseline** stays fixed at 2/4.

05. Connect the explanation to its limit



Follow the labelled picture with the matching explanation. Open the vector diagram

Read the picture in order. Supplied candidate **prediction**s healthy, spotted, spotted, spotted match 3/4: 75%. The arrows show which recorded **input** or intermediate result each decision uses.

Explain the deciding step aloud before checking the changed case; pointing to the final answer alone does not show how it follows. Limit: These supplied **prediction**s illustrate comparison.

No actual model was trained, and four cases are limited **evidence**.

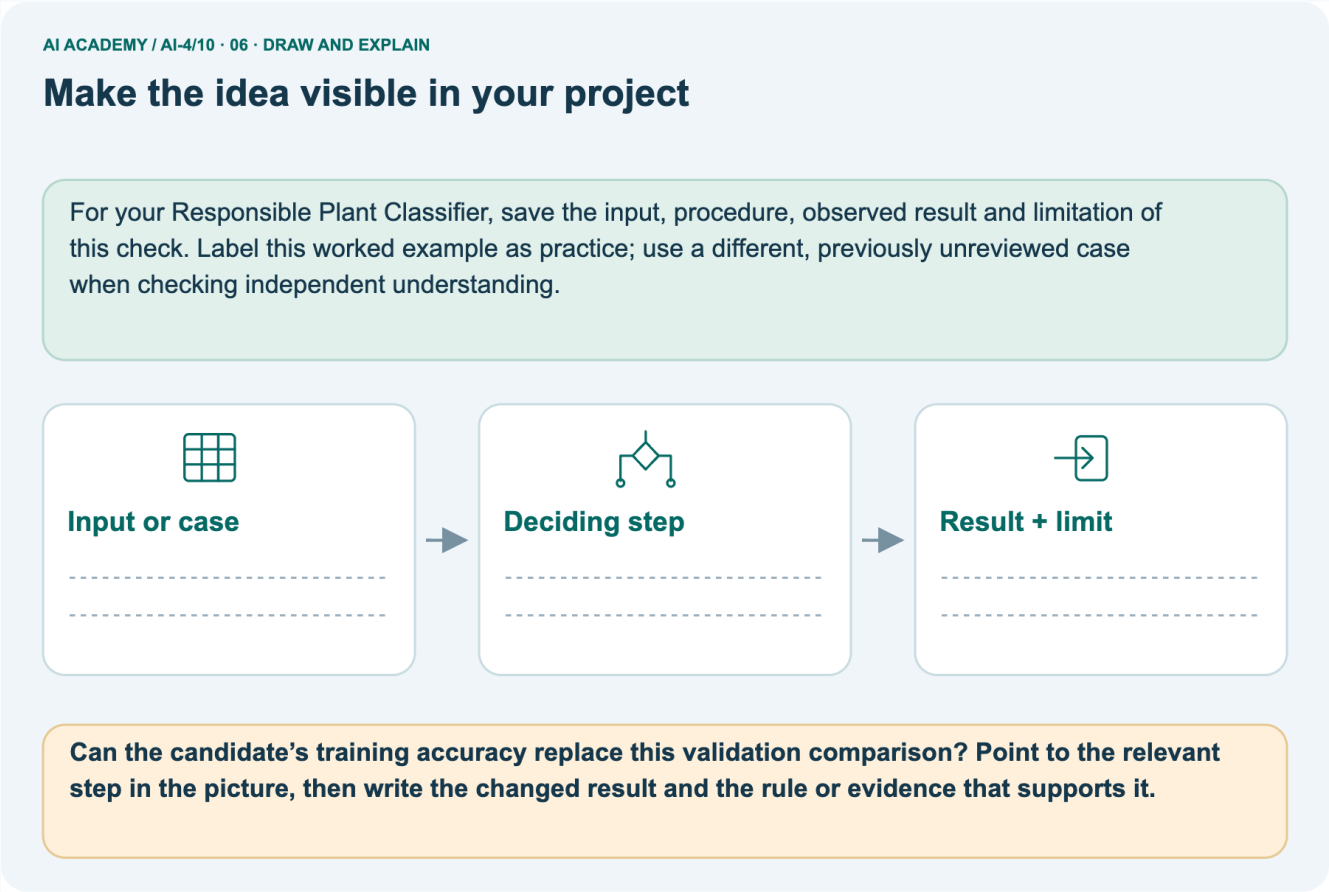
Explain the case, then point to a limit on the claim.

Discuss while showing the picture

Teacher prompt: Explain the case, then point to a limit on the claim. Ask learners to describe one label or arrow in ordinary words. Use the adjacent transcript if the image is difficult to see.

Keep the stated scope: These supplied **prediction**s illustrate comparison. No actual model was trained, and four cases are limited **evidence**.

06. Make the idea visible in your project



Follow the labelled picture with the matching explanation. Open the vector diagram

For your Responsible Plant Classifier, save the **input**, procedure, observed result and limitation of this check.

Label this worked example as practice; use a different, previously unreviewed case when checking independent understanding. Draw your **input**, deciding step, result and limit.

Supported question: Can the candidate's **training accuracy** replace this validation comparison? Point to the relevant step in the picture, then write the changed result and the rule or **evidence** that supports it.

Draw your project trace and justify the deciding step using **evidence**.

Discuss while showing the picture

Teacher prompt: Draw your project trace and justify the deciding step using **evidence**. Ask learners to describe one label or arrow in ordinary words. Use the adjacent transcript if the image is difficult to see.

Supported practice key: No. Compare candidate and **baseline** on the same held-out cases and references; **training** fit answers a different question. Changed-case check: Candidate becomes correct on all 4 validation cases.

Only the previously wrong third decision changes; **baseline** stays fixed at 2/4.

Save or print your responses for your portfolio. Progress and drafts on the hosted website stay in the current browser. Use only fictional or teacher-approved data.